

Big Data Computing

Master's Degree in Computer Science

2025-2026



SAPIENZA
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science
Sapienza Università di Roma
tolomei@di.uniroma1.it

Recap from Last Lecture

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
 - Minimizing “some” aggregated internal cluster distance
- Computing exact solution is **NP-hard** due to exponential search space
- Use an iterative (approximate) solution → e.g., **K-means**

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs
- Tries to minimize the SSD between each point and its assigned centroid

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs
- Tries to minimize the SSD between each point and its assigned centroid
- It iteratively repeats **2 steps:**

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs
- Tries to minimize the SSD between each point and its assigned centroid
- It iteratively repeats **2 steps:**
 - Assignment step → assign each point to the closest centroid

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs
- Tries to minimize the SSD between each point and its assigned centroid
- It iteratively repeats **2 steps:**
 - Assignment step → assign each point to the closest centroid
 - Update step → recompute centroids

One-Slide K-means

- A greedy algorithm to find K hard partitions of N inputs
- Tries to minimize the SSD between each point and its assigned centroid
- It iteratively repeats **2 steps**:
 - **Assignment step** → assign each point to the closest centroid
 - **Update step** → recompute centroids
- The final output depends on key choices (**local optimum**)

Alternative Seed Choice: K-means++

- A method to carefully select initial centroids

Alternative Seed Choice: K-means++

- A method to carefully select initial centroids
- Proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]

Alternative Seed Choice: K-means++

- A method to carefully select initial centroids
- Proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: Spreading out the K initial cluster centers is a good thing

Alternative Seed Choice: K-means++

- A method to carefully select initial centroids
- Proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: Spreading out the K initial cluster centers is a good thing
- Select the i-th centroid as the farthest data point to any other already selected centroids

Alternative Seed Choice: K-means++

1. Choose **one** centroid uniformly at random from among initial data points

Alternative Seed Choice: K-means++

1. Choose **one** centroid uniformly at random from among initial data points
2. For each data point x , compute $D(x)$ as the distance between x and the **nearest centroid** already chosen

Alternative Seed Choice: K-means++

1. Choose **one** centroid uniformly at random from among initial data points
2. For each data point x , compute $D(x)$ as the distance between x and the **nearest centroid** already chosen
3. Choose one new data point at random as a new centroid with probability proportional to $D(x)^2$

Alternative Seed Choice: K-means++

1. Choose **one** centroid uniformly at random from among initial data points
2. For each data point x , compute $D(x)$ as the distance between x and the **nearest centroid** already chosen
3. Choose one new data point at random as a new centroid with probability proportional to $D(x)^2$
4. Repeat steps 2. and 3. until K centroids are chosen, then run Lloyd-Forgy

"Vanilla" K-means vs. K-means++

- Random initialization of "vanilla" K-means may give clusters that are **arbitrarily worse** than optimum

"Vanilla" K-means vs. K-means++

- Random initialization of "vanilla" K-means may give clusters that are **arbitrarily worse** than optimum
- K-means++ provides an upper-bound to the approximation obtained w.r.t. the optimal solution

"Vanilla" K-means vs. K-means++

- Random initialization of "vanilla" K-means may give clusters that are **arbitrarily worse** than optimum
- K-means++ provides an upper-bound to the approximation obtained w.r.t. the optimal solution
- At most, clusters obtained with K-means++ initialization are $O(\log K)$ worse than the optimal partitioning

K-means: How Many Clusters?

- Number of clusters K is given
 - Great! Partition N data points into a predetermined number K of clusters
 - Unfortunately, it is very uncommon to know K in advance

K-means: How Many Clusters?

- Number of clusters K is given
 - Great! Partition N data points into a predetermined number K of clusters
 - Unfortunately, it is very uncommon to know K in advance
- Finding the “right” number K of clusters is part of the problem!
 - Trade-off between having too few and too many clusters
 - Total benefit vs. Total cost

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point x_i to be the similarity to its assigned centroid

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point x_i to be the similarity to its assigned centroid
- Define the total benefit B to be the sum of the individual benefits

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point x_i to be the similarity to its assigned centroid
- Define the total benefit B to be the sum of the individual benefits

NOTE

There is always a clustering whose total benefit $B=N$
(where N is the number of data points)

Why?

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

Goal:

Find the clustering which maximizes V , over all choices of K

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

Goal:

Find the clustering which maximizes V , over all choices of K

B increases with larger values of K , but P allows to stop that

K-means: "Elbow" Method

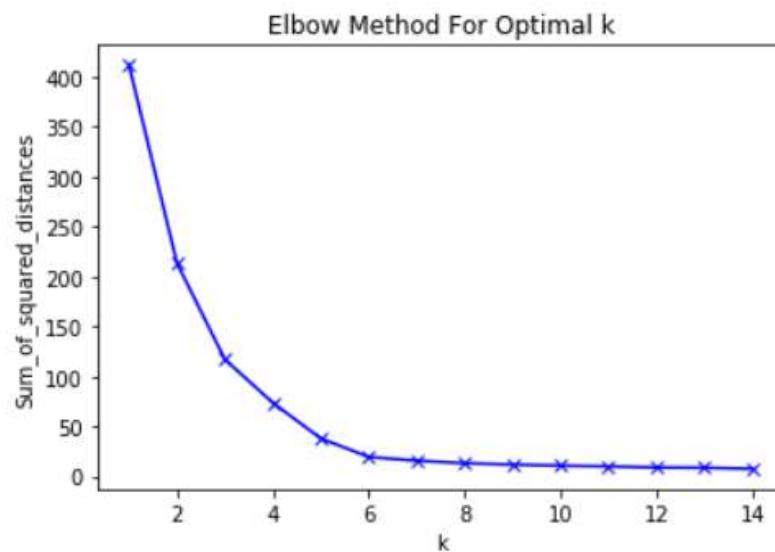
- Empirical method to figure out the right number K of clusters

K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost

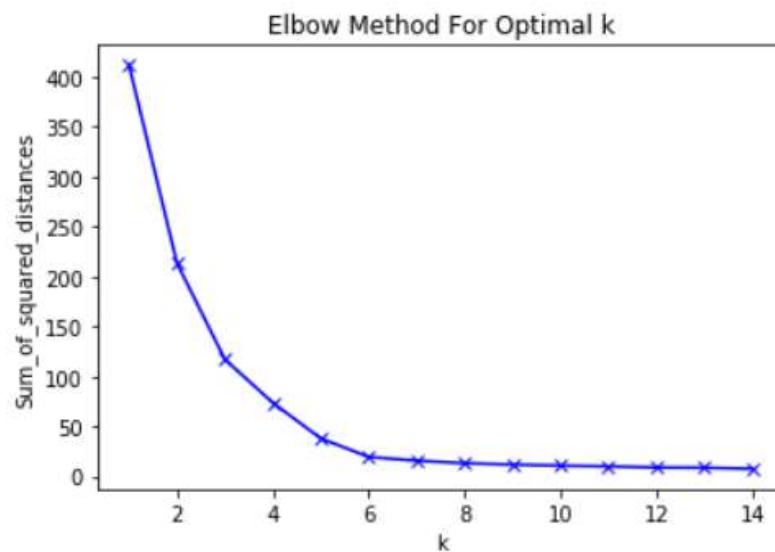
K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



K-means: "Elbow" Method

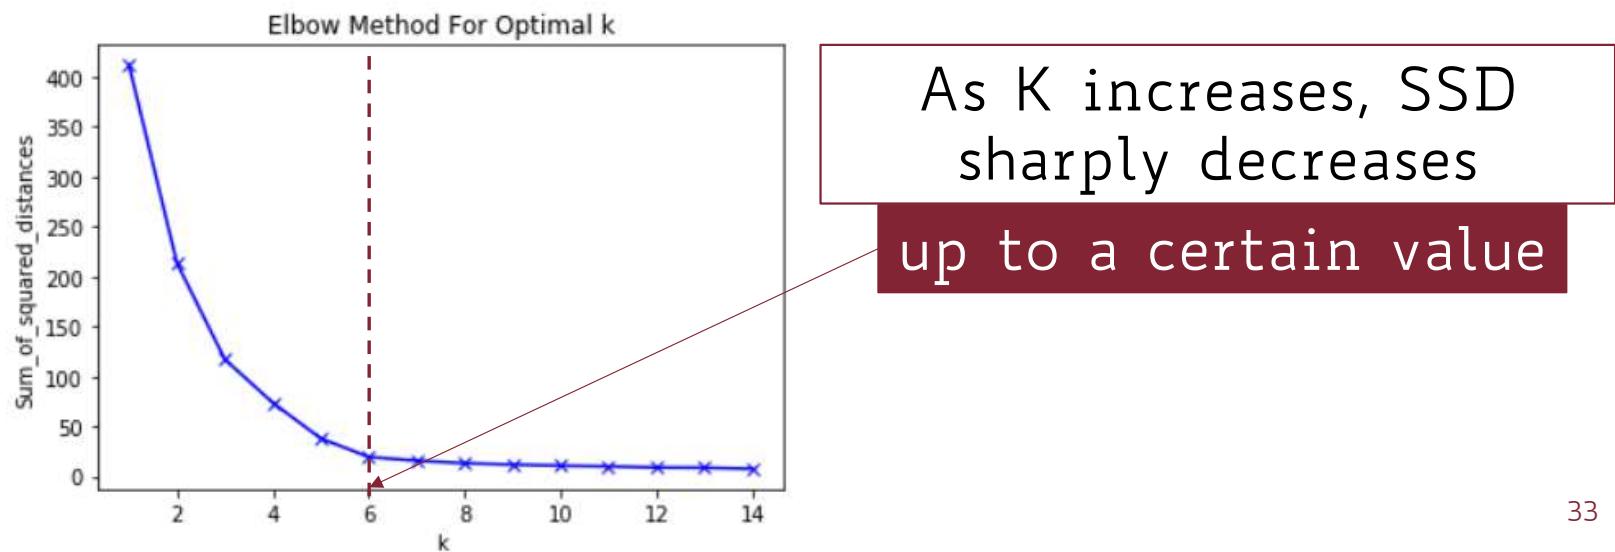
- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



As K increases, SSD sharply decreases

K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ
- Some of them just resemble Euclidean distance, and centroids (i.e., means) still minimize those

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2\text{-Norm}$)
- The same hard clustering framework can be used with other δ
- Some of them just resemble Euclidean distance, and centroids (i.e., means) still minimize those
 - $\delta = \text{Cosine distance} = \text{Euclidean distance on normalized input points}$
 - $\delta = \text{Correlation} = \text{Euclidean distance on standardized input points}$

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ
- Some of them just resemble Euclidean distance, and centroids (i.e., means) still minimize those
 - $\delta = \text{Cosine distance}$ = Euclidean distance on normalized input points
 - $\delta = \text{Correlation}$ = Euclidean distance on standardized input points
- Others, require specific minimizers
 - $\delta = \text{Manhattan distance}$ (L^1 -Norm) → median is the minimizer (**K-medians**)

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ
- **PAM** (**P**artitioning **A**round **M**edoids) greedy Algorithm, introduced by Kaufman and Rousseeuw in 1987 [[paper](#)] vs. Lloyd-Forgy

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ
- **PAM** (**P**artitioning **A**round **M**edoids) greedy Algorithm, introduced by Kaufman and Rousseeuw in 1987 [[paper](#)] vs. Lloyd-Forgy
- Robust to outliers yet computationally expensive $O(K(N-K)^2)$

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space
- (Strong) Assumption on the shape of clusters:
 - Normally distributed around the centroid
 - Independence between data dimensions

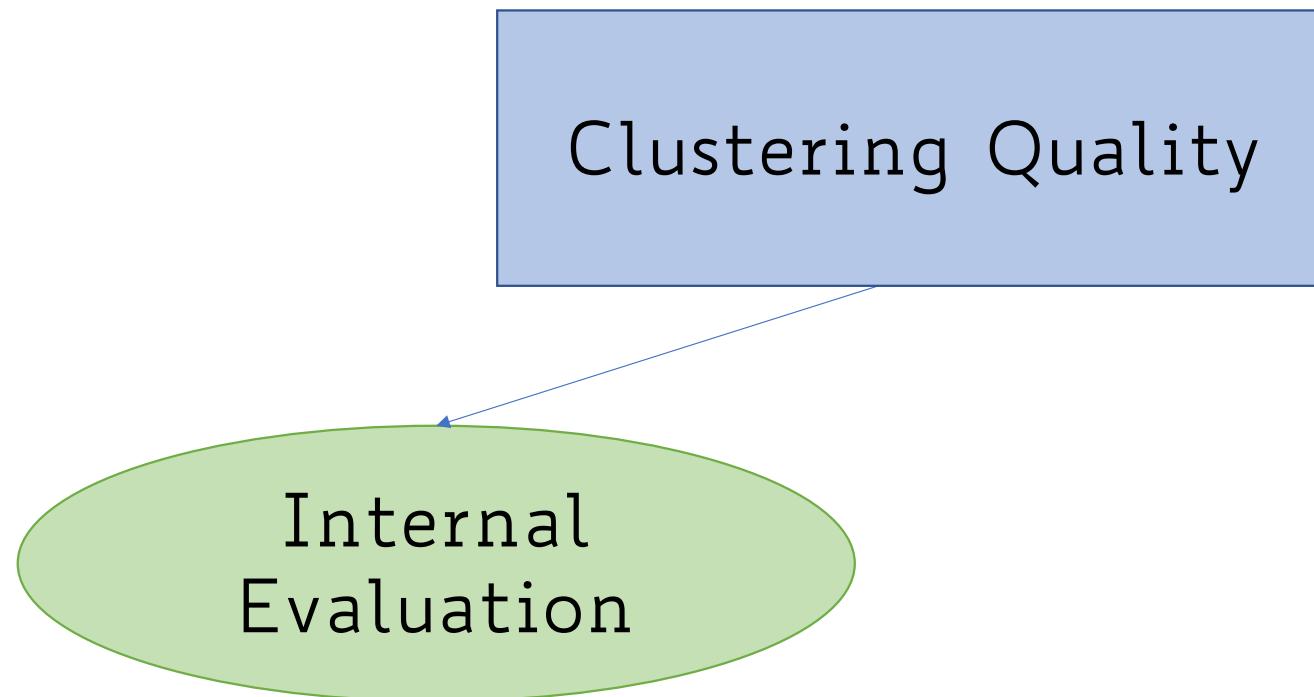
Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space
- (Strong) Assumption on the shape of clusters:
 - Normally distributed around the centroid
 - Independence between data dimensions
- Reference to the original [paper](#)

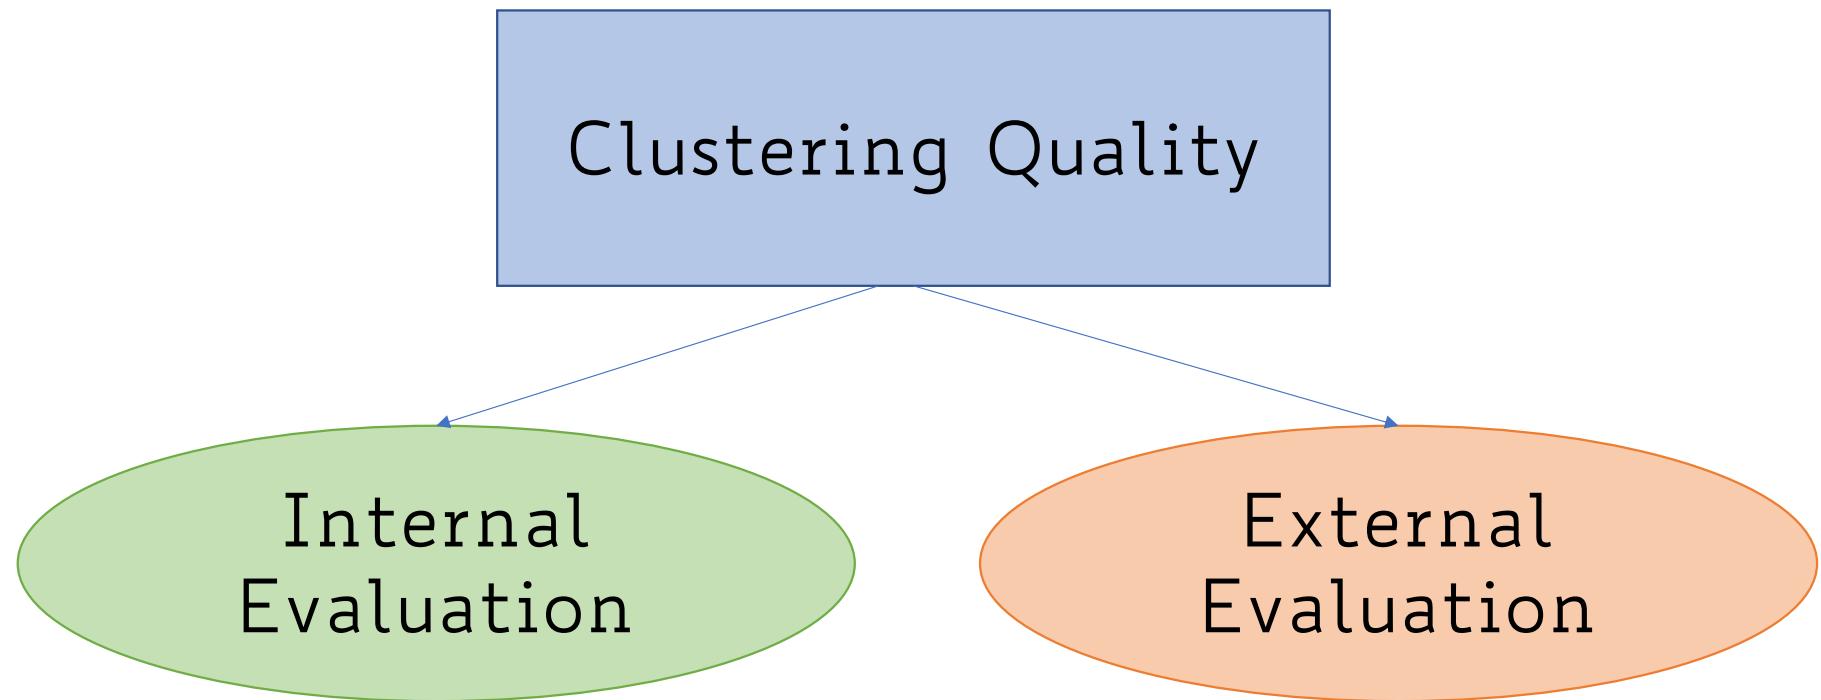
Measures of Clustering Quality

Clustering Quality

Measures of Clustering Quality



Measures of Clustering Quality



Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself

Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself
- A good clustering will produce high quality clusters with:
 - **high intra-cluster similarity**
 - **low inter-cluster similarity**

Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself
- A good clustering will produce high quality clusters with:
 - **high intra-cluster similarity**
 - **low inter-cluster similarity**
- The measured quality of a clustering depends on
 - **data representation**
 - **similarity measure**

Internal Evaluation: Davies-Bouldin Index

$$DB = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} \left(\frac{\sigma_i + \sigma_j}{\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)} \right)$$

K = number of clusters

$\boldsymbol{\mu}_k$ = centroid of cluster C_k

σ_k = avg. distance of all elements of cluster C_k from its centroid $\boldsymbol{\mu}_k$

$\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)$ = distance between centroids of C_i and C_j

Internal Evaluation: Davies-Bouldin Index

$$DB = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} \left(\frac{\sigma_i + \sigma_j}{\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)} \right)$$

K = number of clusters

$\boldsymbol{\mu}_k$ = centroid of cluster C_k

σ_k = avg. distance of all elements of cluster C_k from its centroid $\boldsymbol{\mu}_k$

$\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)$ = distance between centroids of C_i and C_j

The smaller the better

Internal Evaluation: Dunn Index

$$D = \frac{\min_{1 \leq i < j \leq K} \delta(C_i, C_j)}{\max_{1 \leq k \leq K} \delta'(C_k)}$$

K = number of clusters

$\delta(C_i, C_j)$ = distance between cluster C_i and C_j

$\delta'(C_k)$ = intra-cluster distance of cluster C_k

Distance between
centroids

Max distance between
any pair of objects

Internal Evaluation: Dunn Index

$$D = \frac{\min_{1 \leq i < j \leq K} \delta(C_i, C_j)}{\max_{1 \leq k \leq K} \delta'(C_k)}$$

K = number of clusters

$\delta(C_i, C_j)$ = distance between cluster C_i and C_j

$\delta'(C_k)$ = intra-cluster distance of cluster C_k

Distance between
centroids

Max distance between
any pair of objects

The higher the better

Internal Evaluation: Silhouette Coefficient

mean distance between i and all other data points in the same cluster C_i

Internal Evaluation: Silhouette Coefficient

mean distance between i and all other data points in the same cluster C_i

$$a(i) = \frac{1}{|C_i| - 1} \sum_{j \in C_i, j \neq i} \delta(i, j)$$

Internal Evaluation: Silhouette Coefficient

smallest mean distance of i to all points in any other cluster $C_k \neq C_i$

Internal Evaluation: Silhouette Coefficient

smallest mean distance of i to all points in any other cluster $C_k \neq C_i$

$$b(i) = \min_{k \neq i} \frac{1}{|C_k|} \sum_{j \in C_k} \delta(i, j)$$

Internal Evaluation: Silhouette Coefficient

mean distance between i and all other data points in the same cluster C_i

$$a(i) = \frac{1}{|C_i| - 1} \sum_{j \in C_i, j \neq i} \delta(i, j)$$

smallest mean distance of i to all points in any other cluster $C_k \neq C_i$

$$b(i) = \min_{k \neq i} \frac{1}{|C_k|} \sum_{j \in C_k} \delta(i, j)$$

$$s(i) = \begin{cases} 1 - a(i)/b(i) & \text{if } a(i) < b(i) \\ 0 & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1 & \text{if } a(i) > b(i) \end{cases}$$

Internal Evaluation: Silhouette Coefficient

mean distance between i and all other data points in the same cluster C_i

$$a(i) = \frac{1}{|C_i| - 1} \sum_{j \in C_i, j \neq i} \delta(i, j)$$

smallest mean distance of i to all points in any other cluster $C_k \neq C_i$

$$b(i) = \min_{k \neq i} \frac{1}{|C_k|} \sum_{j \in C_k} \delta(i, j)$$

$$s(i) = \begin{cases} 1 - a(i)/b(i) & \text{if } a(i) < b(i) \\ 0 & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1 & \text{if } a(i) > b(i) \end{cases}$$

The higher the better

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)
- Quality measured by the ability to discover some or all of the hidden patterns in gold standard data

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)
- Quality measured by the ability to discover some or all of the hidden patterns in gold standard data
- Hard as it requires labeled data typically provided by human experts

External Evaluation: Purity

C_1, \dots, C_K = set of K clusters

L_1, \dots, L_J = set of J labels

$n_{i,j}$ = number of items with label L_j clustered in C_i

$n_i = \sum_{j=1}^J n_{i,j}$ number of items clustered in C_i

$$\text{purity}(C_i) = \frac{1}{n_i} \max_{j \in \{1, \dots, J\}} n_{i,j}$$

$$\text{purity} = \frac{1}{K} \sum_{i=1}^K \text{purity}(C_i)$$

External Evaluation: Purity

C_1, \dots, C_K = set of K clusters

L_1, \dots, L_J = set of J labels

$n_{i,j}$ = number of items with label L_j clustered in C_i

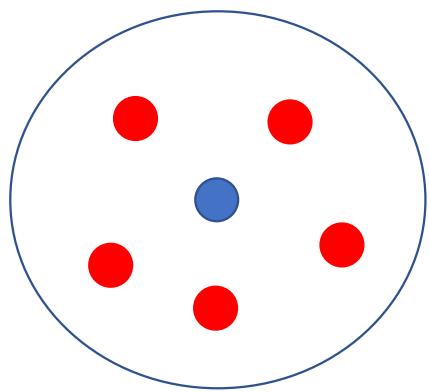
$n_i = \sum_{j=1}^J n_{i,j}$ number of items clustered in C_i

$$\text{purity}(C_i) = \frac{1}{n_i} \max_{j \in \{1, \dots, J\}} n_{i,j}$$

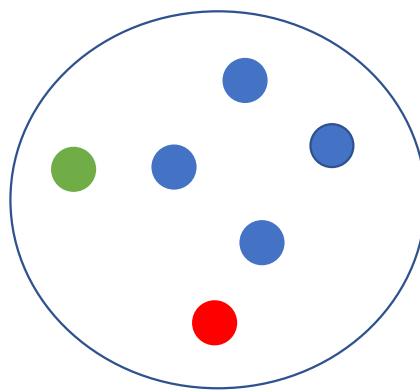
$$\text{purity} = \frac{1}{K} \sum_{i=1}^K \text{purity}(C_i)$$

Biased because having as many clusters as items maximizes purity

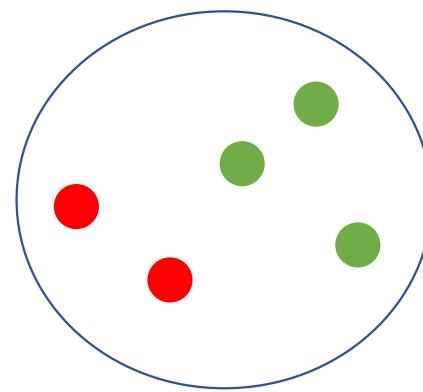
External Evaluation: Purity Example



C_1



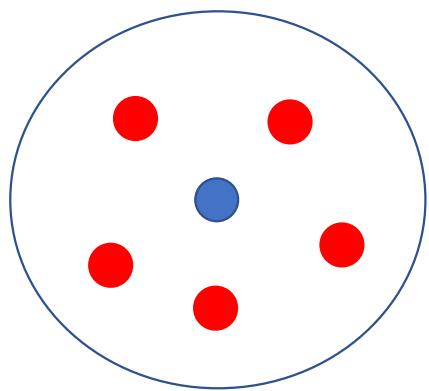
C_2



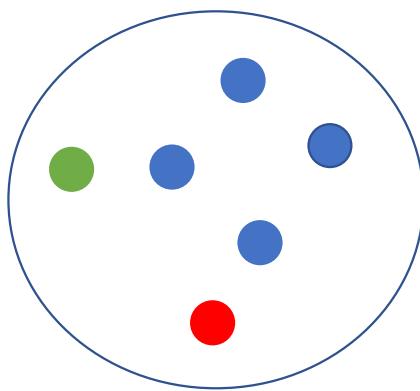
C_3

● L_1 ● L_2 ● L_3

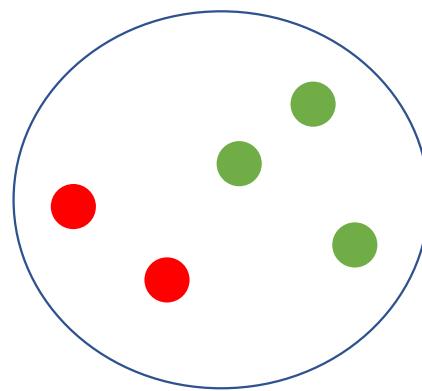
External Evaluation: Purity Example



C_1



C_2

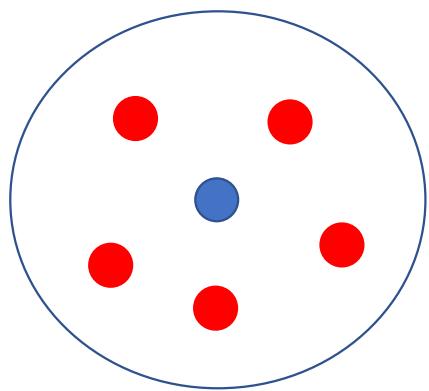


C_3

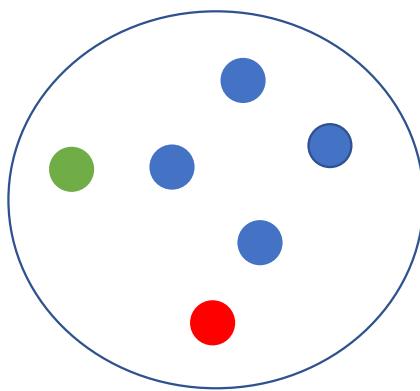
● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

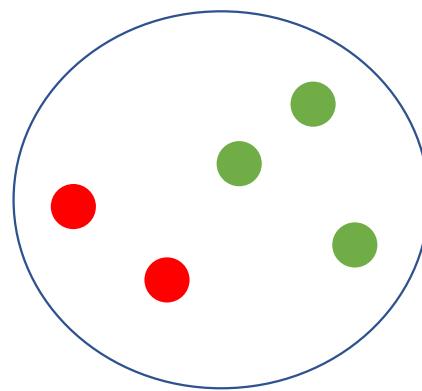
External Evaluation: Purity Example



C_1



C_2



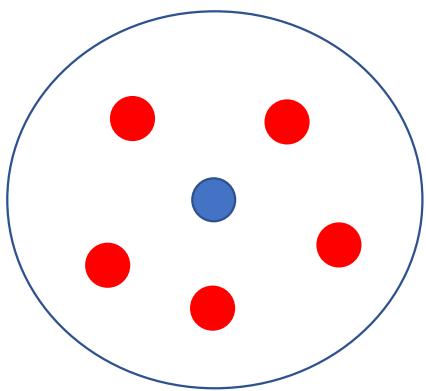
C_3

● L_1 ● L_2 ● L_3

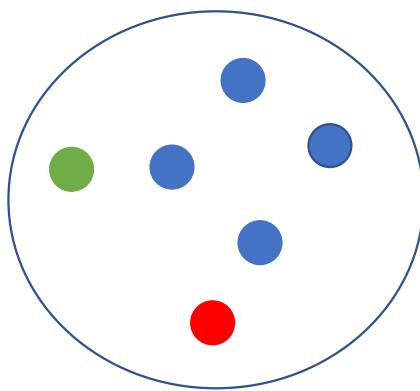
$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$

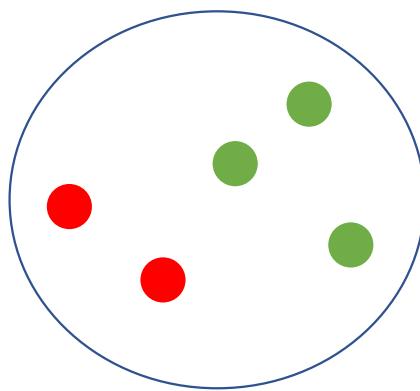
External Evaluation: Purity Example



C_1



C_2



C_3

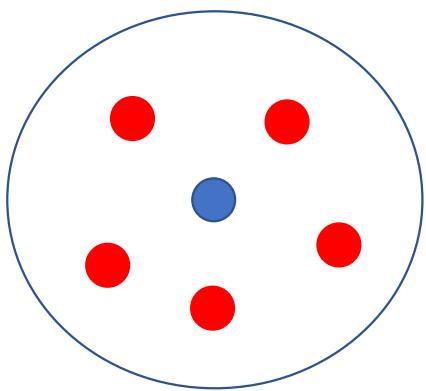
● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

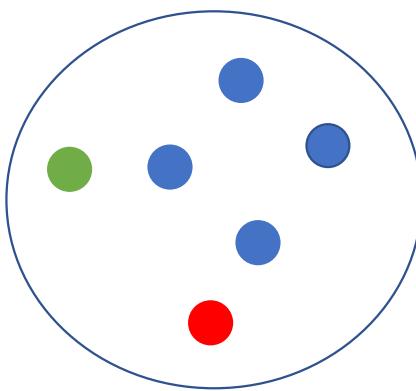
$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$

$$\text{purity}(C_3) = 1/5 * \max\{2, 0, 3\} = 3/5$$

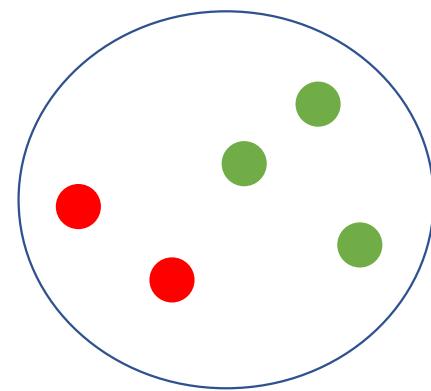
External Evaluation: Purity Example



C_1



C_2



C_3

● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$

$$\text{purity}(C_3) = 1/5 * \max\{2, 0, 3\} = 3/5$$

$$\begin{aligned}\text{purity} &= 1/3 * \text{purity}(C_1) + \text{purity}(C_2) + \text{purity}(C_3) \\ &= 7/10\end{aligned}$$

External Evaluation: Rand Index

$$\text{Rand} = \frac{TP + TN}{TP + TN + FP + FN}$$

TP = number of *true positives*

TN = number of *true negatives*

FP = number of *false positives*

FN = number of *false negatives*

External Evaluation: Rand Index

$$\text{Rand} = \frac{TP + TN}{TP + TN + FP + FN}$$

TP = number of *true positives*

TN = number of *true negatives*

FP = number of *false positives*

FN = number of *false negatives*

All computed from **pairs** of elements

External Evaluation: Rand Index

$$\text{Rand} = \frac{TP + TN}{TP + TN + FP + FN}$$

TP = number of *true positives*

TN = number of *true negatives*

FP = number of *false positives*

FN = number of *false negatives*

All computed from **pairs** of elements

Measures the level of agreement between clustering and ground truth

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth		
Different Clusters in Ground-Truth		

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth	TRUE POSITIVES (TP)	
Different Clusters in Ground-Truth		

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth		
Different Clusters in Ground-Truth		TRUE NEGATIVES (TN)

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth		
Different Clusters in Ground-Truth	FALSE POSITIVES (FP)	

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth		FALSE NEGATIVES (FN)
Different Clusters in Ground-Truth		

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth	TRUE POSITIVES (TP)	FALSE NEGATIVES (FN)
Different Clusters in Ground-Truth	FALSE POSITIVES (FP)	TRUE NEGATIVES (TN)

Confusion Matrix

External Evaluation: Precision, Recall, F-measure

$$P = \frac{TP}{TP + FP} \quad R = \frac{TP}{TP + FN}$$

$$F_{\beta} = \frac{(\beta^2 + 1) \cdot P \cdot R}{\beta^2 \cdot P + R}$$

$$F_1 = \frac{2 \cdot P \cdot R}{P + R}$$

Balances the contribution of false negatives by weighting recall through a parameter β

External Evaluation: Many Other Measures

- Jaccard index
- Dice index
- Fowlkes-Mallows index
- Mutual information
- etc.

Take-Home Message of Today

- **K-means** is an iterative (**approximated**) clustering method that converges to a local minimum

Take-Home Message of Today

- **K-means** is an iterative (**approximated**) clustering method that converges to a local minimum
- Tries to minimize the internal sum of squared Euclidean distances (Lloyd-Forgy Algorithm)

Take-Home Message of Today

- **K-means** is an iterative (**approximated**) clustering method that converges to a local minimum
- Tries to minimize the internal sum of squared Euclidean distances (Lloyd-Forgy Algorithm)
- Many variants:
 - **K-means++**, **K-medoids** (PAM Algorithm), **BFR K-means**, etc.

Take-Home Message of Today

- **K-means** is an iterative (**approximated**) clustering method that converges to a local minimum
- Tries to minimize the internal sum of squared Euclidean distances (Lloyd-Forgy Algorithm)
- Many variants:
 - **K-means++**, **K-medoids** (PAM Algorithm), **BFR K-means**, etc.
- Internal vs. External measures of **clustering quality**