

# Big Data Computing

## Master's Degree in Computer Science

### 2025-2026



SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

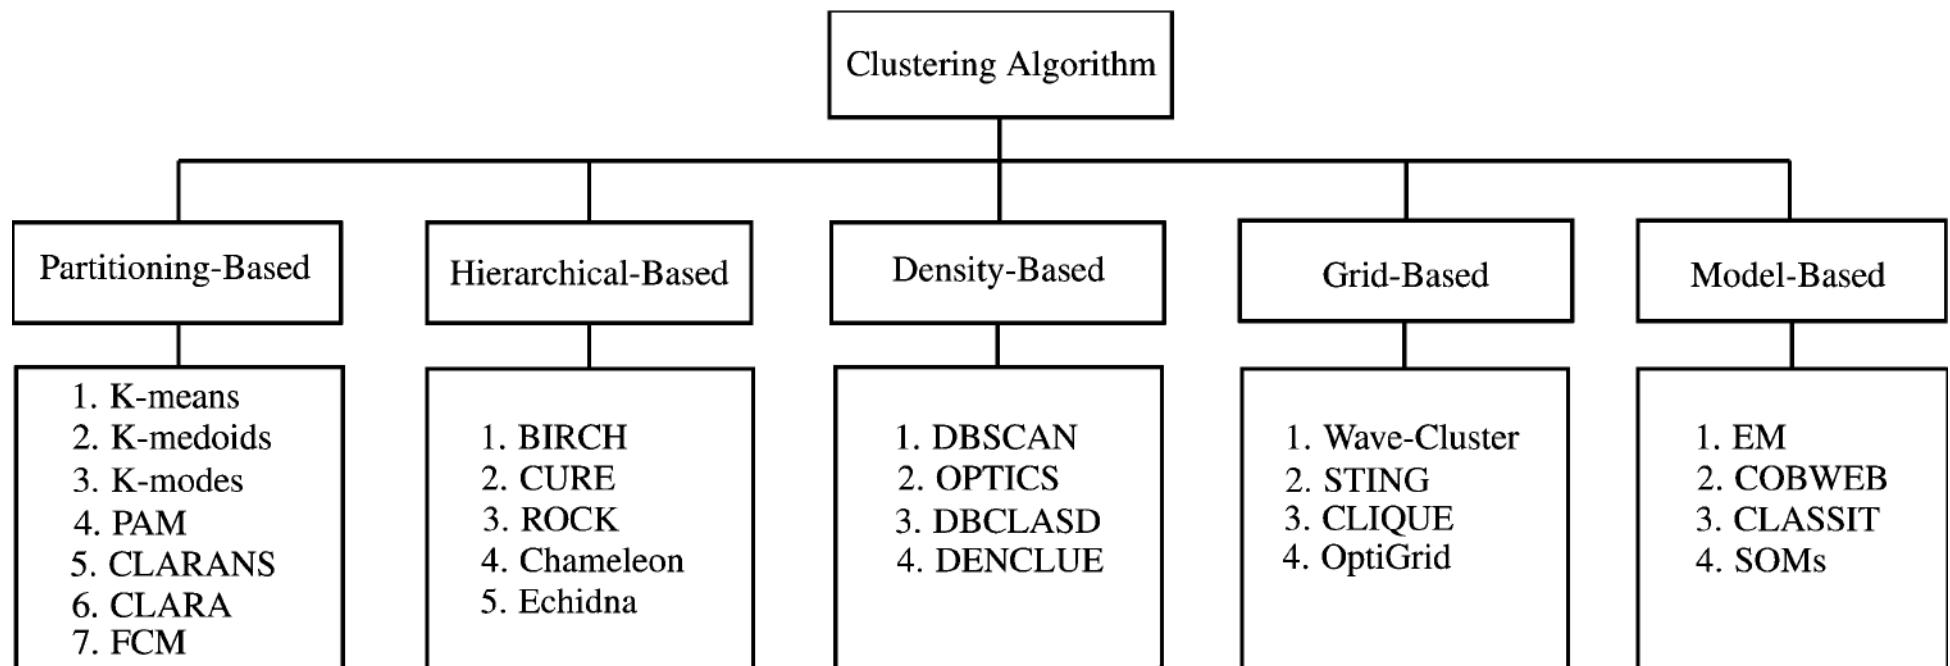
Department of Computer Science  
Sapienza Università di Roma  
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Recap from Last Lecture(s)

- Clustering is an unsupervised learning technique to group "similar" data objects together
- Depends on:
  - object representation
  - similarity measure
- Harder when data dimensionality gets large (**curse of dimensionality**)
- Number of output clusters is part of the problem itself!

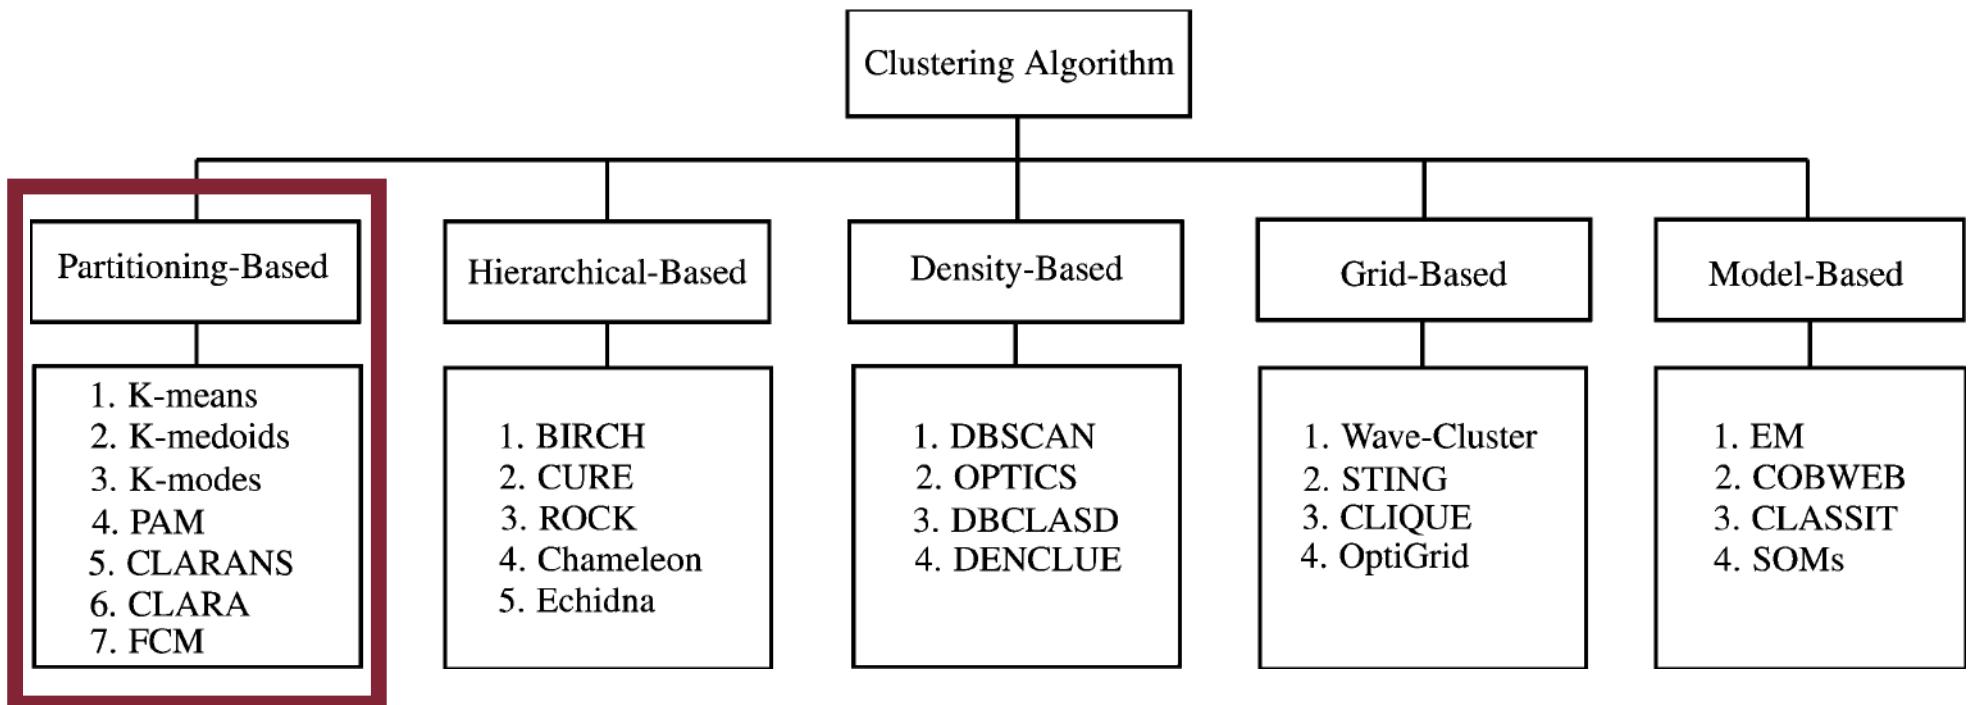
# Clustering Algorithms

# Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEqs2xB>

# Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEqs2xB>

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )
- **Output:** A partition of the  $N$  data points into  $K$  clusters

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )
- **Output:** A partition of the  $N$  data points into  $K$  clusters
- **Goal:** Find the partition which optimizes a certain criterion

# Partitioning: Intuition

Let K=2 for simplicity

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

Here is a possible assignment (i.e., clustering output):

|     | 0 | 1 | 2 | ... | $\dots$ | $N-1$ |   |     |   |   |
|-----|---|---|---|-----|---------|-------|---|-----|---|---|
| $C$ | 0 | 1 | 1 | ... | 0       | 0     | 1 | ... | 0 | 1 |

# Partitioning: Intuition

Let K=2 for simplicity

We can assign each of the N data points to **either** one of the K=2 clusters

We use an N-dimensional assignment vector C,  
where  $C[i] = \{0, 1\}$

Here is another one:

|   | 0 | 1 | 2 | ... | ... | N-1 |   |     |   |   |
|---|---|---|---|-----|-----|-----|---|-----|---|---|
| C | 1 | 1 | 0 | ... | 0   | 0   | 1 | ... | 1 | 1 |

# Partitioning: Intuition

Let K=2 for simplicity

We can assign each of the N data points to **either** one of the K=2 clusters

We use an N-dimensional assignment vector C,  
where  $C[i] = \{0, 1\}$

... And another one:

|   | 0 | 1 | 2 | ... | N-1 |
|---|---|---|---|-----|-----|
| C | 1 | 1 | 0 | ... | 0   |
|   |   |   |   |     |     |
|   |   |   |   |     |     |
|   |   |   |   |     |     |

# Partitioning: Intuition

So, how many possible clustering outputs?

# Partitioning: Intuition

So, how many possible clustering outputs?

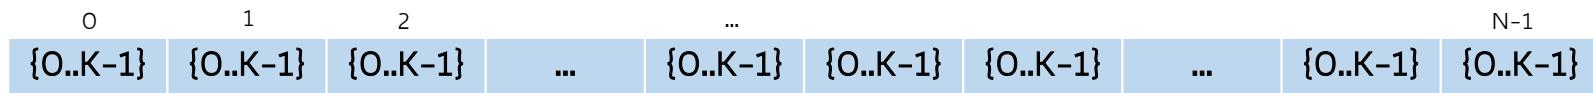
Roughly,  $2^N$ ; More generally,  $K^N$

|          |          |          |     |          |          |
|----------|----------|----------|-----|----------|----------|
| 0        | 1        | 2        | ... | ...      | N-1      |
| {0..K-1} | {0..K-1} | {0..K-1} | ... | {0..K-1} | {0..K-1} |

# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$



Actually, **slightly less** than that because we want each of the  $K$  clusters to contain at least one data point!

# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$

|          |          |          |     |          |          |
|----------|----------|----------|-----|----------|----------|
| 0        | 1        | 2        | ... | ...      | N-1      |
| {0..K-1} | {0..K-1} | {0..K-1} | ... | {0..K-1} | {0..K-1} |

Actually, **slightly less** than that because we want each of the K clusters to contain at least one data point!

In the previous example ( $K=2$ ), the following is **not** allowed

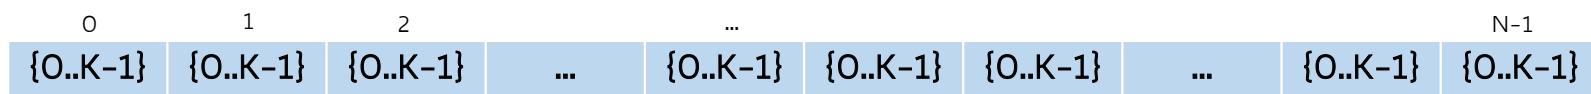
|   |   |   |     |   |   |   |     |   |   |
|---|---|---|-----|---|---|---|-----|---|---|
| 0 | 0 | 0 | ... | 0 | 0 | 0 | ... | 0 | 0 |
|---|---|---|-----|---|---|---|-----|---|---|

|   |   |   |     |     |
|---|---|---|-----|-----|
| 0 | 1 | 2 | ... | N-1 |
| 1 | 1 | 1 | ... | 1   |

# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$



Actually, **slightly less** than that because we want each of the  $K$  clusters to contain at least one data point!

In the previous example ( $K=2$ ), the following is **not** allowed



# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$$

$$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$$

...

$$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$$

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1) \quad \{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$$

$$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1) \quad \{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$$

$$\dots \quad \dots \\ \{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0) \quad \{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$$

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$$\begin{array}{lll} \{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1) & \{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) & \{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) \\ \{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1) & \{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) & \{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) \\ \dots & \dots & \dots \\ \{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0) & \{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) & \{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) \end{array}$$

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

|  |  |  |
|--|--|--|
| $\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$ | $\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$ | $\{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$ |
| $\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$ | $\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$ | $\{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$ |
| ...  | ...  | ...  |
| $\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$ | $\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$ | $\{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$ |

Some combinations are counted twice!

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$$\begin{array}{lll} \{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1) & \{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) & \{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) \\ \{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1) & \{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) & \{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) \\ \dots & \dots & \dots \\ \{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0) & \{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) & \{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) \end{array}$$

Some combinations are counted twice!

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$$\begin{array}{lll} \{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1) & \{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) & \{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1) \\ \{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1) & \{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) & \{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1) \\ \dots & \dots & \dots \\ \{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0) & \{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) & \{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0) \end{array}$$

Some combinations are counted twice!

# Partitioning: NP-Hardness

- **Stirling Partition Number** → K-way non-empty partitions of N elements

# Partitioning: NP-Hardness

- **Stirling Partition Number** → K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

# Partitioning: NP-Hardness

- **Stirling Partition Number** → K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

- Finding the **global optimum** → Intractable for many objective function (enumerate all the possible partitions)\*

\*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

# Partitioning: NP-Hardness

- **Stirling Partition Number** → K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

- Finding the **global optimum** → Intractable for many objective function (enumerate all the possible partitions)\*
- Effective heuristics → K-means, K-medoids, K-means++, etc.

\*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

# Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  the set of  $N$  input data points

$\{C_1, \dots, C_K\}$  the set of  $K$  output clusters

$C_k$  the generic  $k$ -th cluster

$\boldsymbol{\theta}_k$  is the *representative* of cluster  $C_k$

# Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  the set of  $N$  input data points

$\{C_1, \dots, C_K\}$  the set of  $K$  output clusters

$C_k$  the generic  $k$ -th cluster

$\boldsymbol{\theta}_k$  is the *representative* of cluster  $C_k$

## Note:

At this stage we haven't yet specified what a cluster representative actually is

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \quad \forall n \exists! k \text{ such that } \alpha_{n,k} = 1 \wedge \alpha_{n,k'} = 0 \forall k' \neq k$$

hard clustering

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

$\forall n \exists! k \text{ such that } \alpha_{n,k} = 1 \wedge \alpha_{n,k'} = 0 \forall k' \neq k$

**hard clustering**

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore  
exponential search space  
 $S(K, N) \sim O(K^N)$



NP-hard

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore  
exponential search space  
 $S(K, N) \sim O(K^N)$



NP-hard

non-convex due to the  
discrete assignment matrix A



multiple local  
minima

# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)

# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)
- **Non-convexity** doesn't allow us to rely on nice property of convex optimization (unique global optimum)

# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)
- **Non-convexity** doesn't allow us to rely on nice property of convex optimization (unique global optimum)
- A convex objective can be (approximately) solved with numerical methods to find the global optimum

# Iterative Solution: Lloyd-Forgy Algorithm

- **Lloyd-Forgy Algorithm:** 2-step **iterative** approximated solution

# Iterative Solution: Lloyd-Forgy Algorithm

- Lloyd-Forgy Algorithm: 2-step **iterative** approximated solution
- Assignment step
- Update step

# Iterative Solution: Lloyd-Forgy Algorithm

- **Lloyd-Forgy Algorithm:** 2-step **iterative** approximated solution
- **Assignment step**
- **Update step**

Does not guarantee to find the global optimum as it may stuck to a local optimum or a saddle point

# 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

# 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Note:  
Can't take the gradient of  $L$  w.r.t.  $A$   
since  $A$  is discrete!

# 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Intuitively, given a set of fixed representatives,  $L$  is minimized if each data point is assigned to the closest cluster representative according to  $\delta$

( $L$  is the sum of all the distances from each data point to its representative)

# 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Intuitively, given a set of fixed representatives,  $L$  is minimized if each data point is assigned to the closest cluster representative according to  $\delta$

( $L$  is the sum of all the distances from each data point to its representative)

$$\alpha_{n,k} = \begin{cases} 1 & \text{if } \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) = \min_{1 \leq j \leq K} \{\delta(\mathbf{x}_n, \boldsymbol{\theta}_j)\} \\ 0 & \text{otherwise} \end{cases}$$

# 2-Step Optimization: Update Step

Minimize  $L$  w.r.t.  $\Theta$  by fixing  $A$

$L(\Theta|A) = L(\Theta; A) = L$  is a function of  $\Theta$  parametrized by  $A$

# 2-Step Optimization: Update Step

Minimize  $L$  w.r.t.  $\Theta$  by fixing  $A$

$L(\Theta|A) = L(\Theta; A) = L$  is a function of  $\Theta$  parametrized by  $A$

We can minimize  $L$  by taking the **gradient** of  $L$  w.r.t  $\Theta$  (i.e., the vector of partial derivatives), set it to 0 and solve it for  $\Theta$

# 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \left( \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_1}, \dots, \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_K} \right)$$

## 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \left( \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_1}, \dots, \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_K} \right)$$

$$\nabla L(\boldsymbol{\Theta}; A) = \left( \frac{\partial L(\boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_1}, \dots, \frac{\partial L(\boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_K} \right)$$

# 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \left( \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_1}, \dots, \frac{\partial L(\boldsymbol{\Theta}; A)}{\partial \boldsymbol{\theta}_K} \right)$$

$$\nabla L(\boldsymbol{\Theta}; A) = \left( \frac{\partial L(\boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_1}, \dots, \frac{\partial L(\boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_K} \right)$$

$$\boxed{\frac{\partial L(\boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j}}$$

The general j-th partial derivative

# 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

# 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

$$\frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = \boxed{\frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right]}$$

# 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

$$\frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = \boxed{\frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right]}$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_j}$$

To make the notation easier!

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

When computing the partial derivative w.r.t.  $\theta_j$  any other term  $\theta_k$  of the inner summation is treated as constant!

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$= \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \alpha_{n,j} \delta(\mathbf{x}_n, \boldsymbol{\theta}_j) \right] = 0$$

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$= \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \alpha_{n,j} \delta(\mathbf{x}_n, \boldsymbol{\theta}_j) \right] = 0$$

Solve for each  $\theta_j$  independently

Depends on the distance function  $\delta$

# A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)

# A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster

# A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster
- (Re)Assignment of instances to clusters is based on distance/similarity to the current cluster centroids

# A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster
- (Re)Assignment of instances to clusters is based on distance/similarity to the current cluster centroids
- The basic idea is constructing clusters so that the total within-cluster **Sum of Square Distances (SSD)** is minimized

# K-means: Setup

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  the set of  $N$  input data points  
 $\{C_1, \dots, C_K\}$  the set of  $K$  output clusters  
 $C_k$  the generic  $k$ -th cluster

$$\boldsymbol{\theta}_k = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n \in C_k} \mathbf{x}_n$$

$$\text{where } |C_k| = \sum_{n=1}^N \alpha_{n,k}$$

# K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(||\mathbf{x}_n - \boldsymbol{\theta}_k||_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)} \quad \text{Euclidean space}$$

# K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(||\mathbf{x}_n - \boldsymbol{\theta}_k||_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}$$

$$\begin{aligned}\delta(\mathbf{x}_n, \boldsymbol{\theta}_k) &= (||\mathbf{x}_n - \boldsymbol{\theta}_k||_2)^2 = \\ &= \left[ \sqrt{(\mathbf{x}_n - \boldsymbol{\theta}_k)^2} \right]^2 = (\mathbf{x}_n - \boldsymbol{\theta}_k)^2\end{aligned}$$

Sum of Square Distances  
(SSD)

# K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(||\mathbf{x}_n - \boldsymbol{\theta}_k||_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}$$

$$\begin{aligned}\delta(\mathbf{x}_n, \boldsymbol{\theta}_k) &= (||\mathbf{x}_n - \boldsymbol{\theta}_k||_2)^2 = \\ &= \left[ \sqrt{(\mathbf{x}_n - \boldsymbol{\theta}_k)^2} \right]^2 = (\mathbf{x}_n - \boldsymbol{\theta}_k)^2\end{aligned}$$

Sum of Square Distances  
(SSD)

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2$$

# K-means: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

Intuitively, given a set of fixed centroids,  $L$  is minimized if each data point is assigned to the centroid with the smallest SSD

( $L$  is just the SSD from each data point to its assigned centroid)

$$\alpha_{n,k} = \begin{cases} 1 & \text{if } (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 = \min_{1 \leq j \leq K} \{(\mathbf{x}_n - \boldsymbol{\theta}_j)^2\} \\ 0 & \text{otherwise} \end{cases}$$

# K-means: Update Step

Minimize  $L$  w.r.t.  $\Theta$  by fixing  $A$

$$\Theta^* = \operatorname{argmin}_{\Theta} \left\{ \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2}_{L(\Theta; A)} \right\}$$

Compute the gradient w.r.t.  $\Theta$ , set it to 0 and solve it for  $\Theta$

# K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[ \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 \right] = 0 \quad \forall k \in \{1, \dots, K\}$$

# K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[ \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 \right] = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)$$

# K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[ \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 \right] = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)$$

$$\text{Find } \boldsymbol{\theta}_k^* \text{ s.t. } \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k^*) = 0$$

# K-means: Update Step

$$\begin{aligned} \sum_{n=1}^N -2\alpha_{n,k}(\mathbf{x}_n - \boldsymbol{\theta}_k^*) &= 0 \Leftrightarrow \\ 2 \sum_{n=1}^N \alpha_{n,k} \boldsymbol{\theta}_k^* &= 2 \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \\ \boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} &= \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \end{aligned}$$

# K-means: Update Step

$$\sum_{n=1}^N -2\alpha_{n,k}(\mathbf{x}_n - \boldsymbol{\theta}_k^*) = 0 \Leftrightarrow$$
$$2 \sum_{n=1}^N \alpha_{n,k} \boldsymbol{\theta}_k^* = 2 \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$
$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$\boldsymbol{\theta}_k^*$  does not depend on N, therefore it can be factored out

# K-means: Update Step

$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\boldsymbol{\theta}_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n \in C_k} \mathbf{x}_n$$

# K-means: Update Step

$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\boxed{\boldsymbol{\theta}_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n \in C_k} \mathbf{x}_n}$$

# K-means: Update Step

$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\boldsymbol{\theta}_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n \in C_k} \mathbf{x}_n$$

The cluster centroid (i.e., mean) minimizes the objective  
(for a fixed assignment A)

# K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K

# K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids

# K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen

# K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen
4. **Update step:** For each of the K clusters update the centroid by computing the new mean values of all the data points now in the cluster

# K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen
4. **Update step:** For each of the K clusters update the centroid by computing the new mean values of all the data points now in the cluster
5. Iteratively repeat steps 3-4 until a **stopping criterion** is met

# Stopping Criterion

- Several options to choose from:
  - Fixed number of iterations
  - Cluster assignments stop changing (beyond some threshold)
  - Centroid doesn't change (beyond some threshold)

# Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
  - A state in which clusters do not change

# Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
  - A state in which clusters do not change
- Intuitively, in both steps we either improve the objective or not

# Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
  - A state in which clusters do not change
- Intuitively, in both steps we either improve the objective or not
- It is an instance of more general **Expectation Maximization (EM)**
  - EM is known to converge (although not necessarily to a global optimum)

# Lloyd-Forgy's Relationship with EM

- **E-step = Assignment step**

- Each object is assigned to the closest centroid, i.e., to the most likely cluster
- Monotonically decreases SSD

# Lloyd-Forgy's Relationship with EM

- **E-step = Assignment step**

- Each object is assigned to the closest centroid, i.e., to the most likely cluster
- Monotonically decreases SSD

- **M-step = Update step**

- The model (i.e., centroids) are updated (i.e., SSD optimization)
- Monotonically decreases each  $SSD_k$

# Lloyd-Forgy's Complexity Analysis

- Computing the distance between two  $d$ -dimensional data points takes  $O(d)$

# Lloyd-Forgy's Complexity Analysis

- Computing the distance between two  $d$ -dimensional data points takes  $O(d)$
- (Re-)Assigning clusters [E-step]:  $O(KN)$  distance computations or  $O(KNd)$

# Lloyd-Forgy's Complexity Analysis

- Computing the distance between two  $d$ -dimensional data points takes  $O(d)$
- (Re-)Assigning clusters [E-step]:  $O(KN)$  distance computations or  $O(KNd)$
- Computing centroids [M-step]:  $O(Nd)$  as there are  $O(N)$  average computations since each data point is added to a cluster exactly once *at each iteration*, each one taking  $O(d)$

# Lloyd-Forgy's Complexity Analysis

- Computing the distance between two  $d$ -dimensional data points takes  $O(d)$
- (Re-)Assigning clusters [E-step]:  $O(KN)$  distance computations or  $O(KNd)$
- Computing centroids [M-step]:  $O(Nd)$  as there are  $O(N)$  average computations since each data point is added to a cluster exactly once *at each iteration*, each one taking  $O(d)$
- Overall:  $O(RKNd)$  if the 2 steps above are repeated  $R$  times

# K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**

# K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
  - Forgy method **randomly** chooses K data points as the initial means

# K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
  - Forgy method **randomly** chooses K data points as the initial means
  - Random Partition method **randomly** assigns a cluster to each observation

# K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
  - Forgy method **randomly** chooses K data points as the initial means
  - Random Partition method **randomly** assigns a cluster to each observation
- Randomness may converge to **sub-optimal** clusterings

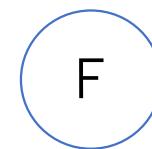
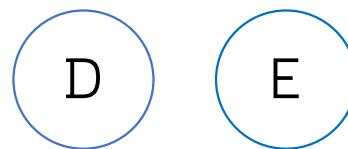
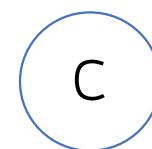
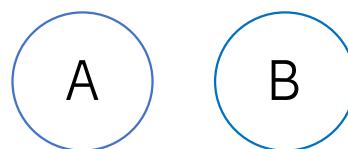
# K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
  - Forgy method **randomly** chooses K data points as the initial means
  - Random Partition method **randomly** assigns a cluster to each observation
- Randomness may converge to **sub-optimal** clusterings

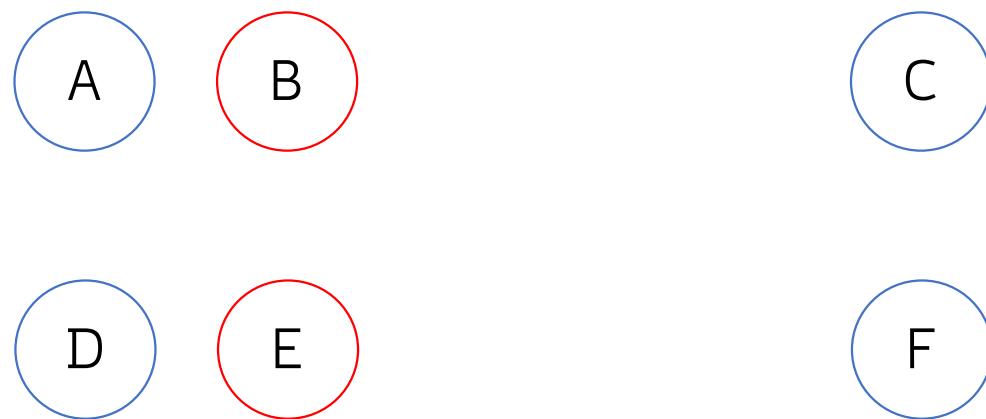
## Problem Mitigation:

Execute several runs of the Lloyd-Forgy algorithm  
with multiple random initialization seeds

# K-means: Seed Choice

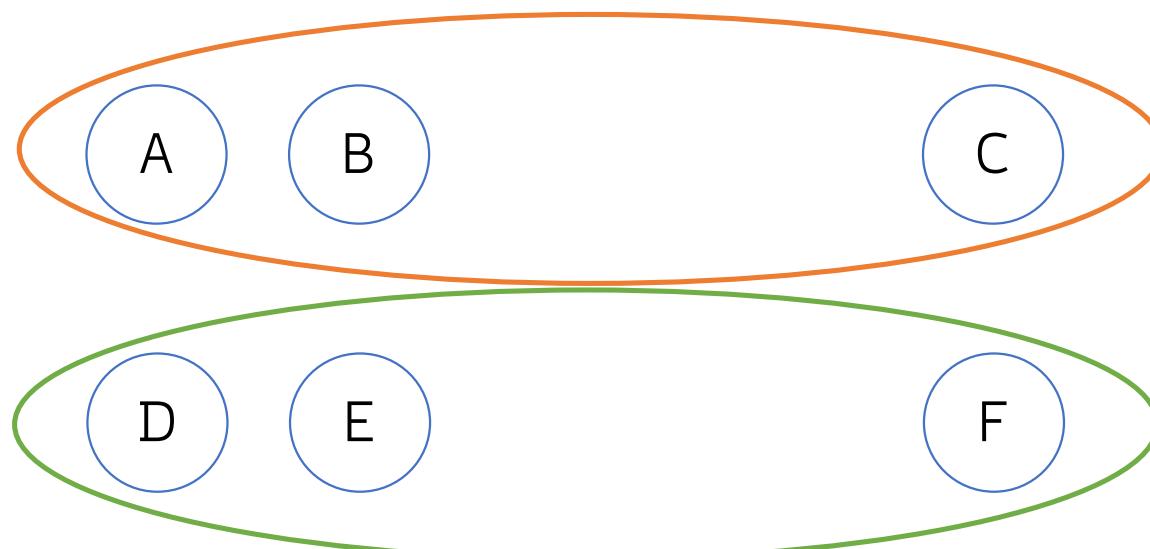


# K-means: Bad (Unlucky) Seed Choice



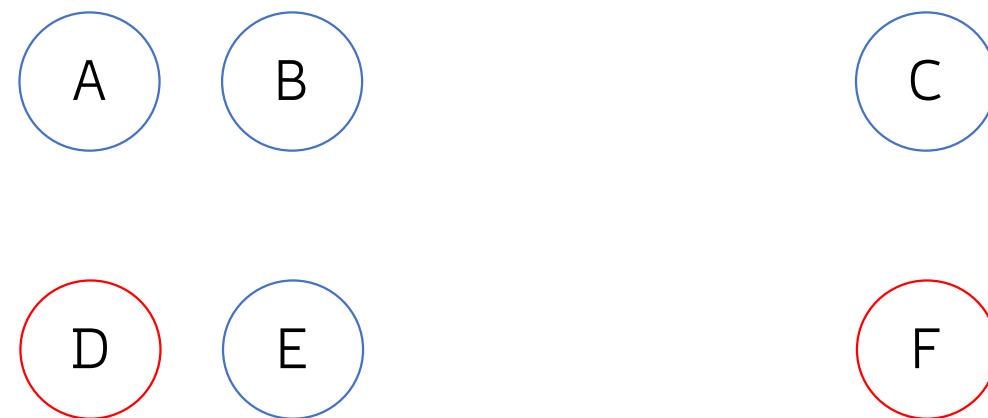
If B and E are randomly chosen as initial centroids...

# K-means: Bad (Unlucky) Seed Choice



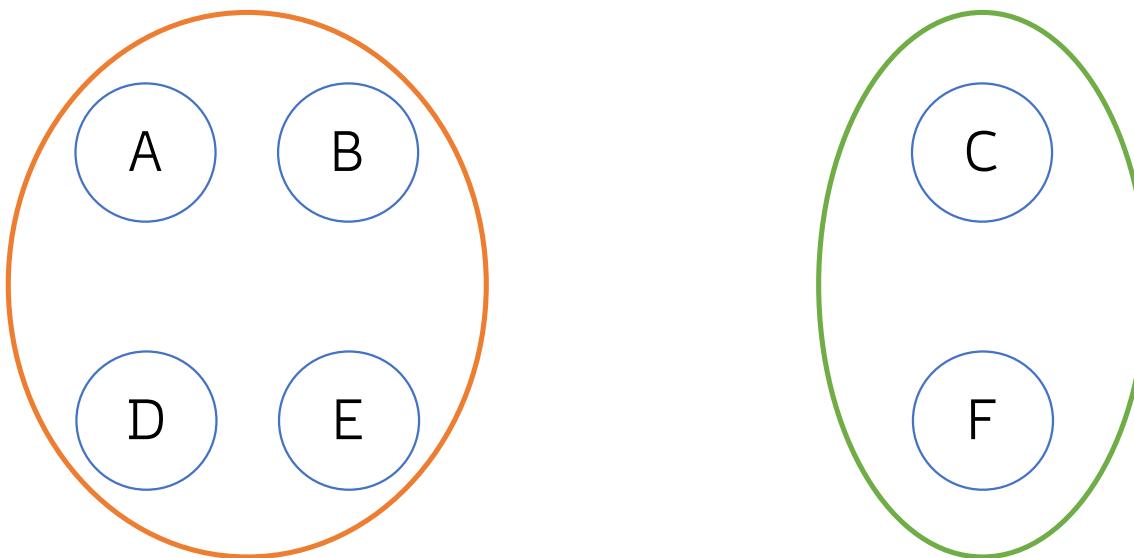
The algorithm converges to the sub-optimal clustering above

# K-means: Good (Lucky) Seed Choice



If D and F are randomly chosen as initial centroids instead...

# K-means: Good (Lucky) Seed Choice



The algorithm converges to a better clustering

# Take-Home Message of Today

- Focus on hard partitioning clustering

# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance

# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance
- Computing exact solution is **NP-hard** due to exponential search space

# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance
- Computing exact solution is **NP-hard** due to exponential search space
- Use an iterative (approximate) solution → e.g., **K-means**