

Big Data Computing

Master's Degree in Computer Science

2025-2026



SAPIENZA
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science
Sapienza Università di Roma
tolomei@di.uniroma1.it

Who Am I?



Who Am I?



UniPI
(1999-2005)



Who Am I?



UniPI
(1999-2005)



UniVE
(2008-2013)

Who Am I?



UniPI
(1999-2005)



UniVE
(2008-2013)



Yahoo! Labs
(2014-2017)
11/12/2025

Who Am I?



UniPI
(1999-2005)



UniVE
(2008-2013)



Yahoo! Labs
(2014-2017)
11/12/2025



UniPD
(2017-2019)

Who Am I?



UniPI
(1999-2005)



UniVE
(2008-2013)



Yahoo! Labs
(2014-2017)
11/12/2025




UniPD
(2017-2019)



Sapienza
(2019-)


My Research Interests



HERCOLE Lab

My Research Interests

Human-Explainable




HERCOLE Lab

My Research Interests


Robust

Human-Explainable




HERCOLE Lab

My Research Interests



My Research Interests

Human-Explainable Robust Collaborative




HERCOLE Lab

LEarning

My Research Interests


Sounds cool?



HERCOLE Lab

My Research Interests

Sounds cool?




HERCOLE Lab

Check out the
lab's [home page](#)

My Research Interests

Sounds cool?



HERCOLE Lab

You can also
follow us on
X (Twitter)
[@HercoleLab](https://twitter.com/HercoleLab)

HERCOLE Lab: People

PhD Students



Flavio Giorgi
PhD Student in Computer Science
Sapienza Università di Roma
[@](#) [home](#) [in](#) [g](#) [Q](#)



Vittoria Vineis
PhD Student in Data Science
Sapienza Università di Roma
[@](#) [in](#) [g](#) [Q](#)



Edoardo Gabrielli
PhD Student in Cybersecurity
Sapienza Università di Roma
[@](#) [home](#) [X](#) [in](#) [g](#) [Q](#)



Fabiano Veglianti
PhD Student in Data Science
Sapienza Università di Roma
[@](#) [X](#) [in](#) [g](#) [Q](#)



Matteo Silvestri
PhD Student in Computer Science
Sapienza Università di Roma
[@](#) [home](#) [X](#) [in](#) [g](#) [Q](#)



Lorenzo Antonelli
PhD Student in Data Science
Sapienza Università di Roma
[@](#) [in](#) [Q](#)

HERCOLE Lab: People

External Collaborators



Fabio Pinelli
Associate Professor of Computer Science
IMT School for Advanced Studies Lucca
[@](#) [X](#) [in](#) [g](#) [o](#)



Fabrizio Silvestri
Full Professor of Computer Science
Sapienza Università di Roma
[@](#) [h](#) [X](#) [in](#) [g](#)



Giuseppe Perelli
Associate Professor of Computer Science
Sapienza Università di Roma
[@](#) [h](#) [o](#) [in](#) [g](#)



Ziheng Chen
Research Scientist
Walmart Labs, Sunnyvale, CA, USA
[@](#) [o](#) [g](#) [id](#)



Federico Siciliano
Postdoctoral researcher
Sapienza Università di Roma
[@](#) [X](#) [o](#) [in](#) [g](#)



Giovanni Trappolini
Assistant Professor of Computer Science
Sapienza Università di Roma
[@](#) [h](#) [X](#) [in](#) [g](#)



Cesare Campagnano
Senior Research Scientist
Pinecone
[@](#) [X](#) [o](#) [in](#) [g](#)

Administrivia

- Class schedule:

- Wednesday from 8:00 a.m. to 11:00 a.m.

Aula Magna

Viale Regina Elena, 295

Administrivia

- Class schedule:

- Wednesday from 8:00 a.m. to 11:00 a.m.
- Thursday from 10:00 a.m. to 12:00 p.m.

Aula Magna

Viale Regina Elena, 295

Aula 1L

Via Castro Laurenziano, 7

Administrivia

- Class schedule:

- Wednesday from 8:00 a.m. to 11:00 a.m.
- Thursday from 10:00 a.m. to 12:00 p.m.

- Office hours:

- Drop me a message to ask for a meeting **online** (Google Meet or Zoom) or in-person at my office (Room 106 @ Viale Regina Elena, 295 – 1st Floor, Building E)

Aula Magna

Viale Regina Elena, 295

Aula 1L

Via Castro Laurenziano, 7

Administrivia

- Contacts:

- Personal homepage: <https://www.di.uniroma1.it/~tolomei>
- Email: tolomei@di.uniroma1.it

Administrivia

- Resources:
 - Course's website: <https://github.com/gtolomei/big-data-computing>
 - Moodle's web page:
<https://elearning.uniroma1.it/enrol/index.php?id=19950>

Administrivia

- Resources:
 - Course's website: <https://github.com/gtolomei/big-data-computing>
 - Moodle's web page:
<https://elearning.uniroma1.it/enrol/index.php?id=19950>
- Class material will be published on the course's website **only**
 - Along with other resources (e.g., suggested readings/books) if needed

Administrivia

- Prerequisites:
 - Familiarity with basics of Data Science and Machine Learning
 - Solid knowledge of Calculus, Linear Algebra, and Probability&Statistics
 - (Python) Programming skills desirable yet not mandatory!

Administrivia


- Prerequisites:
 - Familiarity with basics of Data Science and Machine Learning
 - Solid knowledge of Calculus, Linear Algebra, and Probability&Statistics
 - (Python) Programming skills desirable yet not mandatory!

No worries!

Many subjects will be anyway revisited during class lectures

Administrivia

- Exam: Oral questions on the *whole* program of the course!



Questions?

Outline of the Course

Big Data Infrastructure

Outline of the Course

Big Data Infrastructure

HDFS

Hadoop

Outline of the Course


Big Data Infrastructure

HDFS

Hadoop

Spark

Outline of the Course



Outline of the Course

Big Data Problems and Applications

Big Data Infrastructure

HDFS

Hadoop

Spark

...

Outline of the Course

Big Data Problems and Applications

Similarity in High-Dimensional Spaces

Big Data Infrastructure


HDFS

Hadoop


Spark

...


Outline of the Course




Outline of the Course




Outline of the Course




Outline of the Course




Outline of the Course




Outline of the Course



Outline of the Course



Outline of the Course



Let's Get Started!

The Search Problem

Given a collection S of N items, find if $x \in S$

The Search Problem

Given a collection S of N items, find if $x \in S$

This is a very general problem that occurs frequently

The Search Problem

Given a collection S of N items, find if $x \in S$

This is a very general problem that occurs frequently

Example:

Find if a number occurs in a list of N integers

The Search Problem

If the list is **not** sorted, it takes $O(N)$ steps to respond

The Search Problem

If the list is **not** sorted, it takes $O(N)$ steps to respond

If the list is sorted, it takes $O(\log(N))$ steps to respond
but we must pay the cost of sorting it $O(N\log(N))$

The Search Problem

If the list is **not** sorted, it takes $O(N)$ steps to respond

If the list is sorted, it takes $O(\log(N))$ steps to respond
but we must pay the cost of sorting it $O(N\log(N))$

Still, it might be beneficial to pre-sort the list if we
repeat the find operation several times...

Time Complexity

Let M be the number of times the find operation is called

Time Complexity

Let M be the number of times the find operation is called

Without sorting, the total complexity is $O(M*N)$ steps

Time Complexity

Let M be the number of times the find operation is called

Without sorting, the total complexity is $O(M*N)$ steps

With sorting, the complexity becomes $O(N\log(N)) + O(M\log(N))$

Time Complexity

Let M be the number of times the find operation is called

Without sorting, the total complexity is $O(M*N)$ steps

With sorting, the complexity becomes $O(N\log(N)) + O(M\log(N))$

We must find the value of M for which:

$$O(N\log(N)) + O(M\log(N)) < O(M*N)$$

with sorting

without sorting

Time Complexity Trade-Off

Brutally:

$$N \log(N) + M \log(N) < MN$$

Time Complexity Trade-Off

Brutally:

$$N \log(N) + M \log(N) < MN$$

$$M \log(N) - MN < -N \log(N)$$

Time Complexity Trade-Off

Brutally:

$$N \log(N) + M \log(N) < MN$$

$$M \log(N) - MN < -N \log(N)$$

$$MN - M \log(N) > N \log(N)$$

Time Complexity Trade-Off

Brutally:

$$N \log(N) + M \log(N) < MN$$

$$M \log(N) - MN < -N \log(N)$$

$$MN - M \log(N) > N \log(N)$$

$$M(N - \log(N)) > N \log(N)$$

Time Complexity Trade-Off

Brutally:

$$N \log(N) + M \log(N) < MN$$

$$M \log(N) - MN < -N \log(N)$$

$$MN - M \log(N) > N \log(N)$$

$$M(N - \log(N)) > N \log(N)$$

$$M \geq \lceil N \log(N) / (N - \log(N)) \rceil \text{ [assuming } N \in \mathbb{Z}_{>0} \text{]}$$

Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$

Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000			

Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000	2		

Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	

Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	$1000*3 + 2*3 = 3006$

Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	$1000*3 + 2*3 = 3006$

Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	$1000*3 + 2*3 = 3006$
1000	3	3000	$1000*3 + 3*3 = 3009$
1000	4	4000	$1000*3 + 4*3 = 3012$

Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	1000*3 + 2*3 = 3006
1000	3	3000	1000*3 + 3*3 = 3009
1000	4	4000	1000*3 + 4*3 = 3012

$$M \geq \lceil N \log(N) / (N - \log(N)) \rceil \text{ [assuming } N \in \mathbb{Z}_{>0} \text{]}$$

$$M \geq \lceil 1000 * \log(1000) / (1000 - \log(1000)) \rceil$$

$$M \geq \lceil 3000 / (1000 - 3) \rceil$$

$$M \geq \lceil 3.009 \rceil = 4$$

'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If N grows large and the list is **not** sorted, a linear scan can be too costly

'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If N grows large and the list is **not sorted**, a linear scan can be too costly

On the other hand, if N **does not fit into main memory**, we will need some external R-way merge-sorting

'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If N grows large and the list is **not sorted**, a linear scan can be too costly

On the other hand, if N **does not fit into main memory**, we will need some external R-way merge-sorting

In both cases, the complexity may arise from the fact that we are dealing with a massive 'vertical' input size N

'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If N grows large and the list is **not sorted**, a linear scan can be too costly

On the other hand, if N **does not fit into main memory**, we will need some external R-way merge-sorting

In both cases, the complexity may arise from the fact that we are dealing with a massive 'vertical' input size N

'vertical' here means that the **number** N of input data points is large, yet their **representation** (e.g., 8-byte integers) is not!

'Vertical' vs. 'Horizontal' Scale

Suppose that S is a collection of *images*, not integers

'Vertical' vs. 'Horizontal' Scale

Suppose that S is a collection of **images**, not integers

Our goal is to find if a given image x is in S
(or retrieve an image x' in S that is 'most similar' to x)

'Vertical' vs. 'Horizontal' Scale

Suppose that S is a collection of **images**, not integers

Our goal is to find if a given image x is in S
(or retrieve an image x' in S that is 'most similar' to x)

We may have to cope both with the large number of items N in
the collection **and** the complex representation of each item

'Vertical' vs. 'Horizontal' Scale

Suppose that S is a collection of **images**, not integers

Our goal is to find if a given image x is in S
(or retrieve an image x' in S that is 'most similar' to x)

We may have to cope both with the large number of items N in
the collection **and** the complex representation of each item

Each image must be represented by a high-dimensional vector
of RGB pixels

'Vertical' vs. 'Horizontal' Scale

Suppose that S is a collection of **images**, not integers

Our goal is to find if a given image x is in S
(or retrieve an image x' in S that is 'most similar' to x)

We may have to cope both with the large number of items N in
the collection **and** the complex representation of each item


Each image must be represented by a high-dimensional vector
of RGB pixels

For example, a 100x100 pixel image requires a 30000-dimensional
real-value vector


'Vertical' vs. 'Horizontal' Scale

27
1290
-45
...
32
-670
...
42
...
810

'Vertical' vs. 'Horizontal' Scale



'Vertical' vs. 'Horizontal' Scale




'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810

vertical



...



...



...




'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810

vertical



...



...



...


`find(``)`

'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810

vertical



...



...



`find(`)



horizontal

'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

Very large **number** of input data points
(vertical)

AND

Very large **representation** of input data points
(horizontal)

'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

Very large **number** of input data points
(vertical)

AND

Very large **representation** of input data points
(horizontal)

We will focus mostly on the second and still assume we deal with very large number of input data points

Similarity Measures

- What does "**similar**" mean?

Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data

Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data
- Domain and representation specific (e.g., similar images vs. similar text documents)

Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data
- Domain and representation specific (e.g., similar images vs. similar text documents)
- Crucial for several big data tasks (e.g., search/retrieval/filter, clustering, classification, etc.)

Notion of Similarity

- We implicitly assumed data live in a d -dimensional **Euclidean space**

Notion of Similarity

- We implicitly assumed data live in a d -dimensional **Euclidean space**
- Similarity between data is computed using:
 - **Euclidean metric** (i.e., distance)
 - **Cosine similarity**
 - **Jaccard coefficient**
 - ...

Metric and Metric Space

X is a set

δ is a function $\delta : X \times X \rightarrow [0, \infty)$, where:

1. $\delta(x, y) \geq 0$ (**non-negativity**)
2. $\delta(x, y) = 0 \Leftrightarrow x = y$ (**identity** of indiscernibles)
3. $\delta(x, y) = \delta(y, x)$ (**symmetry**)
4. $\delta(x, y) \leq \delta(x, z) + \delta(z, y)$ (**triangle inequality**)

Then δ is called a **metric** (or distance function) and X a **metric space**

Euclidean Metric (Distance) & Euclidean Space

$$X = \mathbb{R}^d$$

$$\delta : \mathbb{R}^d \times \mathbb{R}^d \rightarrow [0, \infty)$$

$\mathbf{x} = (x_1, \dots, x_d)$ and $\mathbf{y} = (y_1, \dots, y_d)$ are 2 points in \mathbb{R}^d

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + \dots + (x_d - y_d)^2} = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$$

Euclidean Norm (L^2 -Norm)

- The position of a point in a Euclidean d -space is a **Euclidean vector**

Euclidean Norm (L^2 -Norm)

- The position of a point in a Euclidean d -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

Euclidean Norm (L^2 -Norm)

- The position of a point in a Euclidean d -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

$$\|\mathbf{x}\|_2 = \sqrt{x_1^2 + \dots + x_d^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$$

where \cdot indicates the **dot product**

Euclidean Norm (L²-Norm)

- The position of a point in a Euclidean d -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

$$\|\mathbf{x}\|_2 = \sqrt{x_1^2 + \dots + x_d^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$$

where \cdot indicates the **dot product**

This can be just seen as the Euclidean distance between
vector's tail and tip

Euclidean Norm & Euclidean Metric

Let $\mathbf{x} - \mathbf{y} = (x_1 - y_1, \dots, x_d - y_d)$ the **displacement vector** between \mathbf{x} and \mathbf{y}

Euclidean Norm & Euclidean Metric

Let $\mathbf{x} - \mathbf{y} = (x_1 - y_1, \dots, x_d - y_d)$ the **displacement vector** between \mathbf{x} and \mathbf{y}

The Euclidean distance between x and y is just the Euclidean norm of the displacement vector

Euclidean Norm & Euclidean Metric

Let $\mathbf{x} - \mathbf{y} = (x_1 - y_1, \dots, x_d - y_d)$ the **displacement vector** between \mathbf{x} and \mathbf{y}

The Euclidean distance between x and y is just the Euclidean norm of the displacement vector

$$\delta(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_2 = \sqrt{(\mathbf{x} - \mathbf{y}) \cdot (\mathbf{x} - \mathbf{y})}$$

Euclidean Distance: 1-dimensional Case

$$d = 1$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}$$

$\mathbf{x} = x, \mathbf{y} = y$ both \mathbf{x} and \mathbf{y} are scalars

$$\delta(\mathbf{x}, \mathbf{y}) = \delta(x, y) = \sqrt{(x - y)^2} = |x - y|$$

Euclidean Distance: 1-dimensional Case

$$d = 1$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}$$

$\mathbf{x} = x, \mathbf{y} = y$ both \mathbf{x} and \mathbf{y} are scalars

$$\delta(\mathbf{x}, \mathbf{y}) = \delta(x, y) = \sqrt{(x - y)^2} = |x - y|$$

The Euclidean distance between any two 1-d points on the real line is the **absolute value** of the numerical difference of their coordinates

Euclidean Distance: 2-dimensional Case

$$d = 2$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}^2$$

$$\mathbf{x} = (x_1, x_2), \mathbf{y} = (y_1, y_2)$$

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} = \|\mathbf{x} - \mathbf{y}\|_2$$

Euclidean Distance: 2-dimensional Case

$$d = 2$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}^2$$

$$\mathbf{x} = (x_1, x_2), \mathbf{y} = (y_1, y_2)$$

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} = \|\mathbf{x} - \mathbf{y}\|_2$$

The Euclidean distance between any two 2-d points on the Euclidean plane equals to the Pythagorean theorem

Minkowski Distance (L^p -Norm)

Generalization of the Euclidean distance

$\mathbf{x} = (x_1, \dots, x_d)$ and $\mathbf{y} = (y_1, \dots, y_d) \in \mathbb{R}^d$

$$\delta_p(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^d |x_i - y_i|^p \right)^{\frac{1}{p}}$$

Minkowski Distance (L^p -Norm): $p=1$

L^1 -Norm or Manhattan Distance

$$\delta_1(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^d |x_i - y_i|^1 \right)^{\frac{1}{1}} = \sum_{i=1}^d |x_i - y_i|$$

Minkowski Distance (L^p -Norm): $p=2$

L^2 -Norm or Euclidean Distance

$$\delta_2(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^d |x_i - y_i|^2 \right)^{\frac{1}{2}} = \sqrt{\sum_{i=1}^d |x_i - y_i|^2}$$

Minkowski Distance (L^p -Norm): $p=\infty$

L^∞ -Norm or Chebyshev Distance

$$\delta_\infty(\mathbf{x}, \mathbf{y}) = \lim_{p \rightarrow \infty} \left(\sum_{i=1}^d |x_i - y_i|^p \right)^{\frac{1}{p}} = \\ = \max\{|x_1 - y_1|, |x_2 - y_2|, \dots, |x_d - y_d|\}$$

Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space

Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors


Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors
- It ranges between [-1,1]

Cosine Similarity


- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors
- It ranges between [-1,1]
- It captures the **orientation** and not the magnitude

Cosine Similarity




θ is close to 0°
 $\cos(\theta) \approx 1$

similar vectors



θ is close to 90°
 $\cos(\theta) \approx 0$


orthogonal vectors



θ is close to 180°
 $\cos(\theta) \approx -1$

opposite vectors

Cosine Similarity: 2-dimensional Case




$$\theta = \beta - \alpha$$


$$x = (\underbrace{\|x\|\cos\alpha}_{x_1}, \underbrace{\|x\|\sin\alpha}_{x_2})$$

$$y = (\underbrace{\|y\|\cos\beta}_{y_1}, \underbrace{\|y\|\sin\beta}_{y_2})$$


Cosine Similarity: 2-dimensional Case




Cosine Similarity: 2-dimensional Case




Cosine Similarity: 2-dimensional Case




Cosine Similarity: 2-dimensional Case




Cosine Similarity: 2-dimensional Case



Cosine Similarity: 2-dimensional Case



Cosine Similarity: 2-dimensional Case



$$\mathbf{x} \cdot \mathbf{y} = \|\mathbf{x}\| \|\mathbf{y}\| \cos \theta$$



$$\cos \theta = \mathbf{x} \cdot \mathbf{y} / \|\mathbf{x}\| \|\mathbf{y}\|$$

Cosine Similarity: d -dimensional Case

- Computed as in the case of 2-dimensional vectors

Cosine Similarity: d -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two d -dimensional vectors are not collinear then they span a 2-dimensional plane $E \subset \mathbb{R}^d$

Cosine Similarity: d -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two d -dimensional vectors are not collinear then they span a 2-dimensional plane $E \subset \mathbb{R}^d$
- This plane E inherits the dot product in \mathbb{R}^d and so becomes an ordinary Euclidean plane

Cosine Similarity: d -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two d -dimensional vectors are not collinear then they span a 2-dimensional plane $E \subset \mathbb{R}^d$
- This plane E inherits the dot product in \mathbb{R}^d and so becomes an ordinary Euclidean plane
- The angles in this plane are related to the dot product as they are in 2-dimensional vector geometry

Jaccard Index (Coefficient)

Measures similarity between finite sample sets

Jaccard Index (Coefficient)


Measures similarity between finite sample sets

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$$

$$J(A, B) = 1 \text{ if } A = B = \emptyset$$

$$0 \leq J(A, B) \leq 1$$

Jaccard Index (Coefficient): Interpretation



source: [Wikipedia](#)

Jaccard Distance

Complementary to the Jaccard coefficient

$$\delta_J(A, B) = 1 - J(A, B) = \frac{|A \cup B| - |A \cap B|}{|A \cup B|}$$

This distance is a **metric** on the collection of all finite sets

Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
 - number of input data points vs. high-dimensional representation

Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
 - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items

Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
 - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items
- Different similarity measures:
 - Euclidean, Cosine, Jaccard, etc.

Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
 - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items
- Different similarity measures:
 - Euclidean, Cosine, Jaccard, etc.
- We'll see the notion of similarity can be flawed in high-dimensional spaces