

# Big Data Computing

## Master's Degree in Computer Science

### 2025-2026



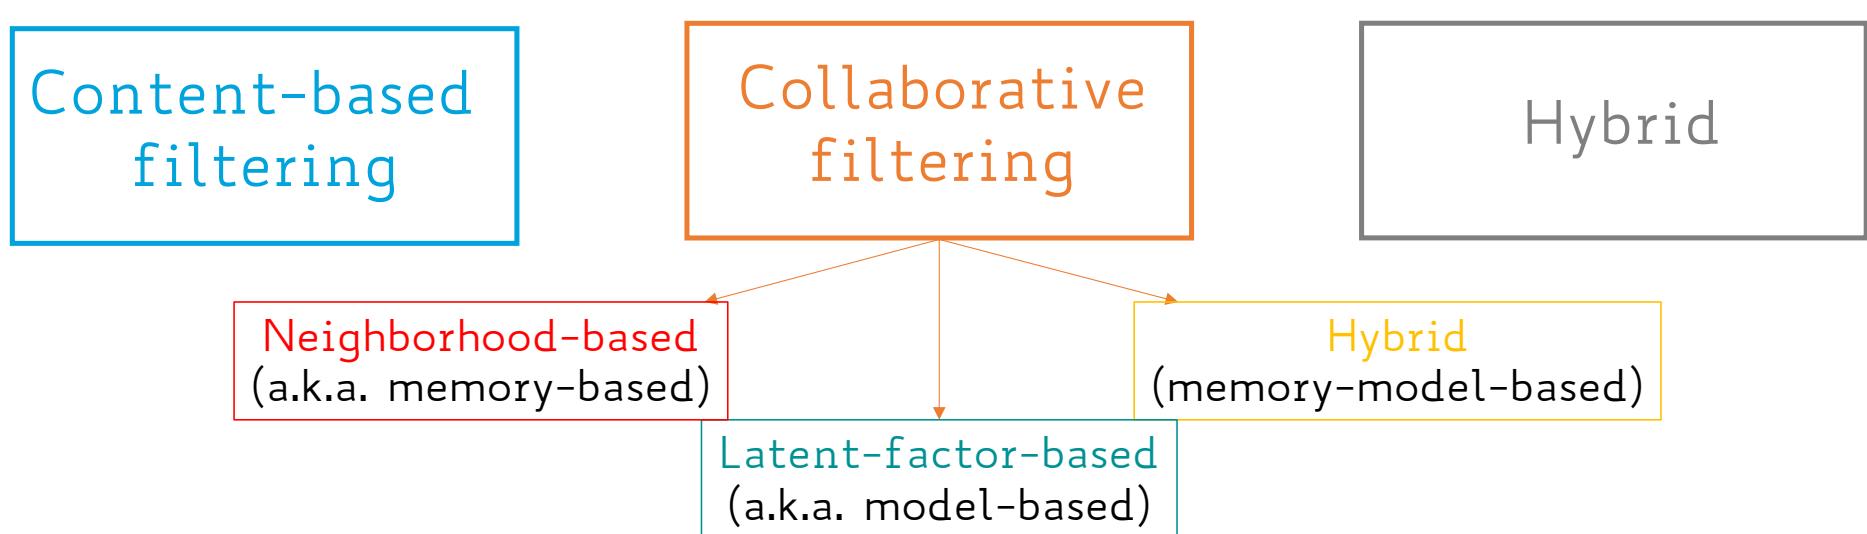
SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science  
Sapienza Università di Roma  
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Recommendation Strategies

3 approaches to recommender systems



# COLLABORATIVE FILTERING

# Collaborative Filtering (CF)

## Idea

Recommend items to user  $u$  based on preferences of other users similar to  $u$

# Collaborative Filtering (CF)

## Idea

Recommend items to user  $u$  based on preferences of other users similar to  $u$

Core concept:  
User-to-User or Item-to-Item similarity

# Collaborative Filtering (CF)

## Idea

Recommend items to user  $u$  based on preferences of other users similar to  $u$

Core concept:  
User-to-User or Item-to-Item similarity

No need for explicit creation of user/item profiles

# Collaborative Filtering: Approaches

3 main approaches to collaborative filtering

# Collaborative Filtering: Approaches

3 main approaches to collaborative filtering

Neighborhood-based  
(a.k.a. memory-based)

# Collaborative Filtering: Approaches

3 main approaches to collaborative filtering

Neighborhood-based  
(a.k.a. memory-based)

Latent-factor-based  
(a.k.a. model-based)

# Collaborative Filtering: Approaches

3 main approaches to collaborative filtering

Neighborhood-based  
(a.k.a. memory-based)

Hybrid  
(memory-model-based)

Latent-factor-based  
(a.k.a. model-based)

# Neighborhood-based (Memory-based) CF

Compute the relationship between users or items

# Neighborhood-based (Memory-based) CF

Compute the relationship between users or items

## User-based

Evaluates a user's preference for an item based on ratings of "neighboring" users for that item

# Neighborhood-based (Memory-based) CF

Compute the relationship between **users** or **items**

## User-based

Evaluates a user's preference for an item based on ratings of "neighboring" users for that item

## Item-based

Evaluates a user's preference for an item based on ratings of "neighboring" items by the same user

# USER-BASED COLLABORATIVE FILTERING

# User-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

# User-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of users  $\{u': u' \neq u\}$  who have already rated  $i$   
extract a subset of  $k$  **neighbours** of  $u$

# User-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of users  $\{u': u' \neq u\}$  who have already rated  $i$   
extract a subset of  $k$  **neighbours** of  $u$

$k$ -neighborhood of  $u$  is found on the basis of the similarity  
between user ratings without explicit user profiles

# User-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of users  $\{u': u' \neq u\}$  who have already rated  $i$   
extract a subset of  $k$  **neighbours** of  $u$

$k$ -neighborhood of  $u$  is found on the basis of the similarity  
between user ratings without explicit user profiles

Estimate  $r(u, i)$  based on the ratings of users in the  
 $k$ -neighborhood of  $u$

# User-based Neighborhood

In theory, rating prediction  $r(u,i)$  could be defined on  
any item  $i$  not rated by  $u$

# User-based Neighborhood

In theory, rating prediction  $r(u,i)$  could be defined on  
any item  $i$  not rated by  $u$

In practice, we are interested only in estimating  $r(u,i)$  for  
those items  $i$  which have been rated by the  $u$ 's  $k$ -neighborhood

# User-based Neighborhood

In theory, rating prediction  $r(u,i)$  could be defined on  
any item  $i$  not rated by  $u$

In practice, we are interested only in estimating  $r(u,i)$  for  
those items  $i$  which have been rated by the  $u$ 's  $k$ -neighborhood

Intuitively, if a user  $v$  is not in the  $u$ 's  $k$ -neighborhood  
then very likely  $u$  will not be interested in any item  
that only  $v$  has rated

# User-based Neighborhood

In theory, rating prediction  $r(u,i)$  could be defined on  
any item  $i$  not rated by  $u$

In practice, we are interested only in estimating  $r(u,i)$  for  
those items  $i$  which have been rated by the  $u$ 's  $k$ -neighborhood

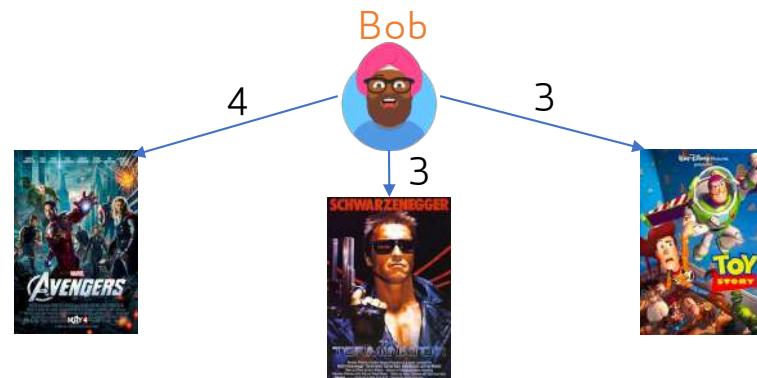
Intuitively, if a user  $v$  is not in the  $u$ 's  $k$ -neighborhood  
then very likely  $u$  will not be interested in any item  
that only  $v$  has rated

In other words, the  $u$ 's  $k$ -neighborhood must be  
computed first to narrow down the set of items which  
we must predict the rating of

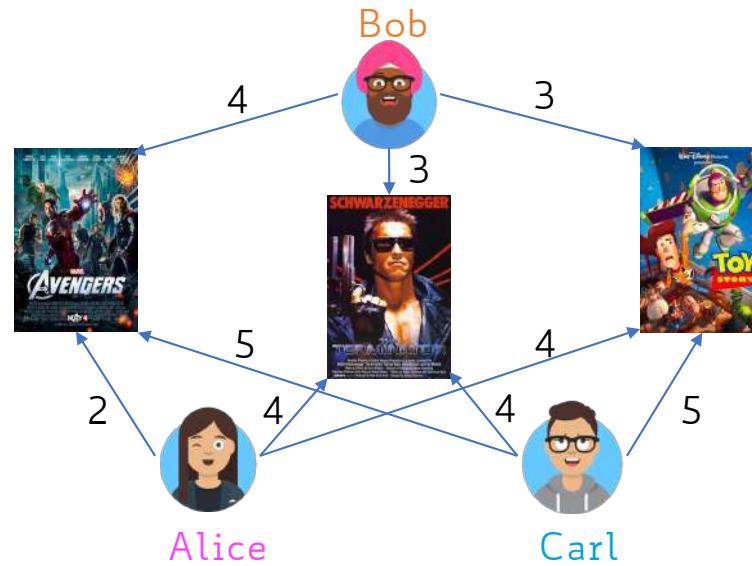
# User-based Neighborhood: Example

		MOVIES							
USERS	Alice	2		5	4	5	4	4	
	Bob	4				3		3	
	Carl	5	5	3	4	5	4	5	
	...	...	...	...	...	...	...	...	
	Zoe		1	3				5	
								4	

# User-based Neighborhood: Example

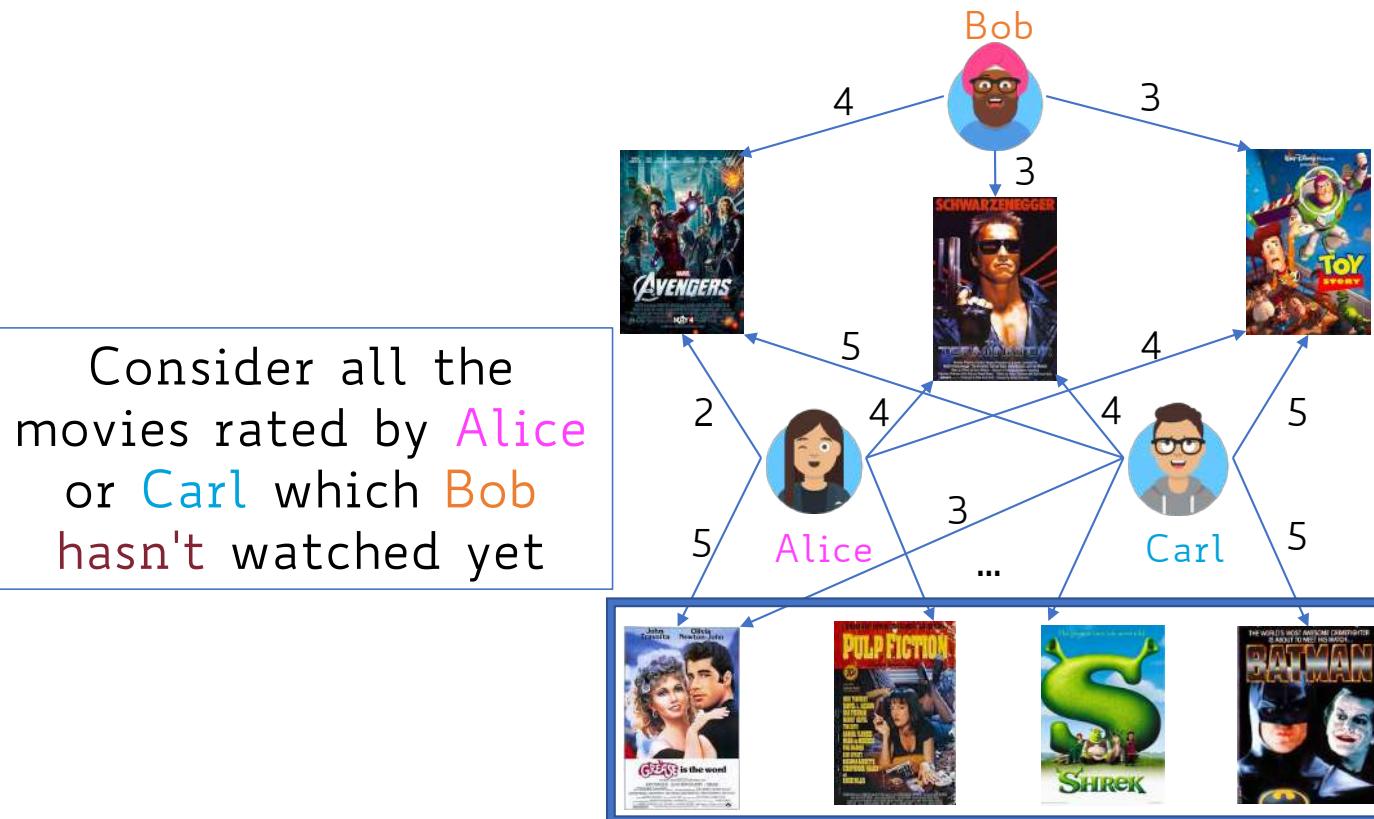


# User-based Neighborhood: Example



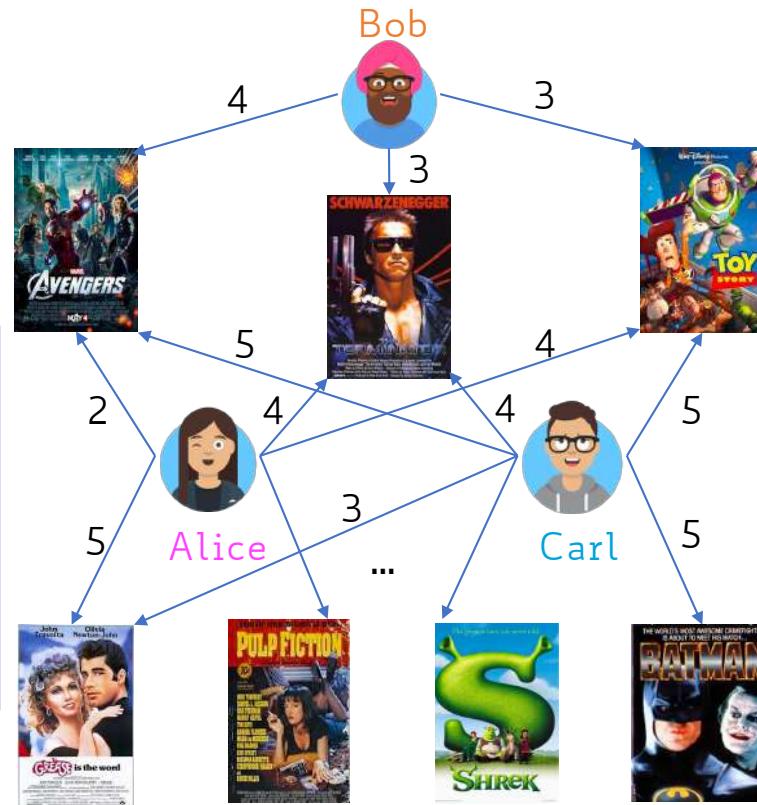
Alice and Carl are the 2-nearest neighbours of Bob if we look at their rating behaviours

# User-based Neighborhood: Example



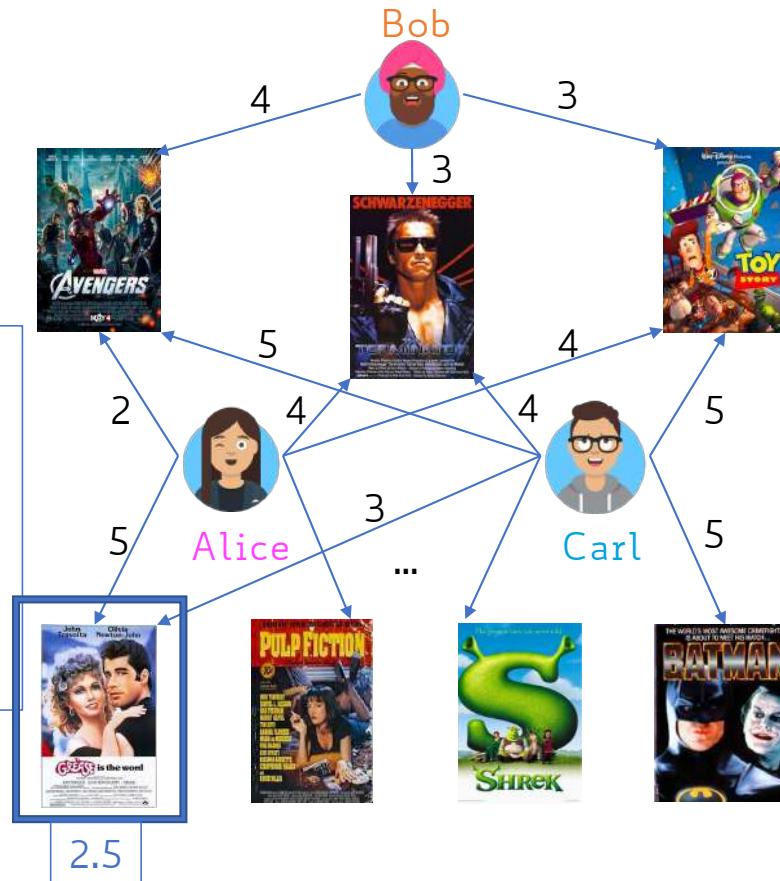
# User-based Neighborhood: Example

Predict the rating that Bob would give to each of those movies on the basis of Alice's and Carl's ratings



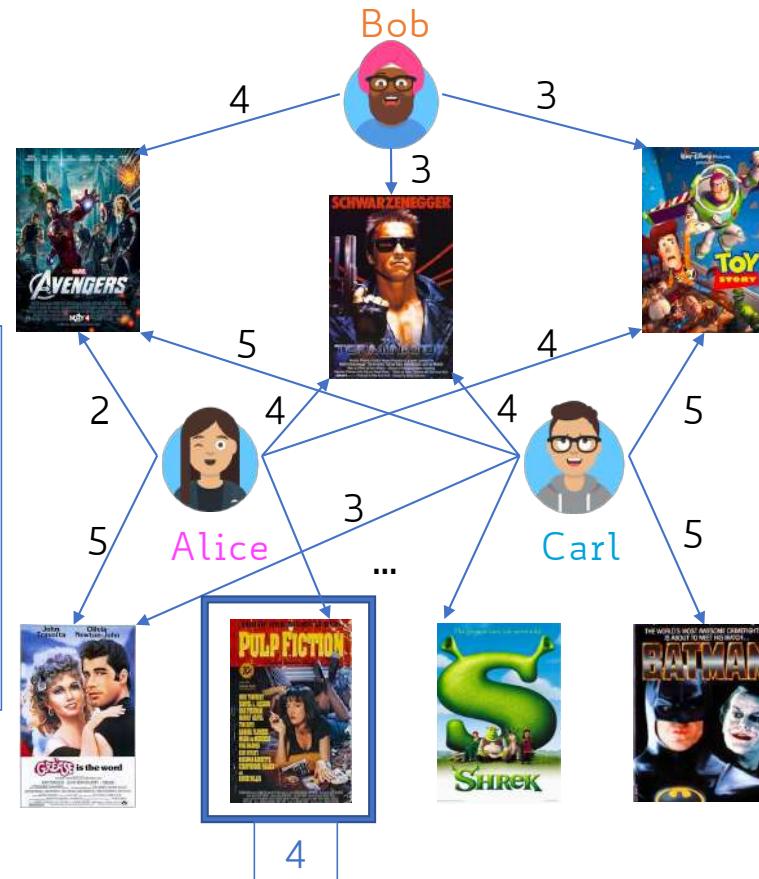
# User-based Neighborhood: Example

Predict the rating that Bob would give to each of those movies on the basis of Alice's and Carl's ratings



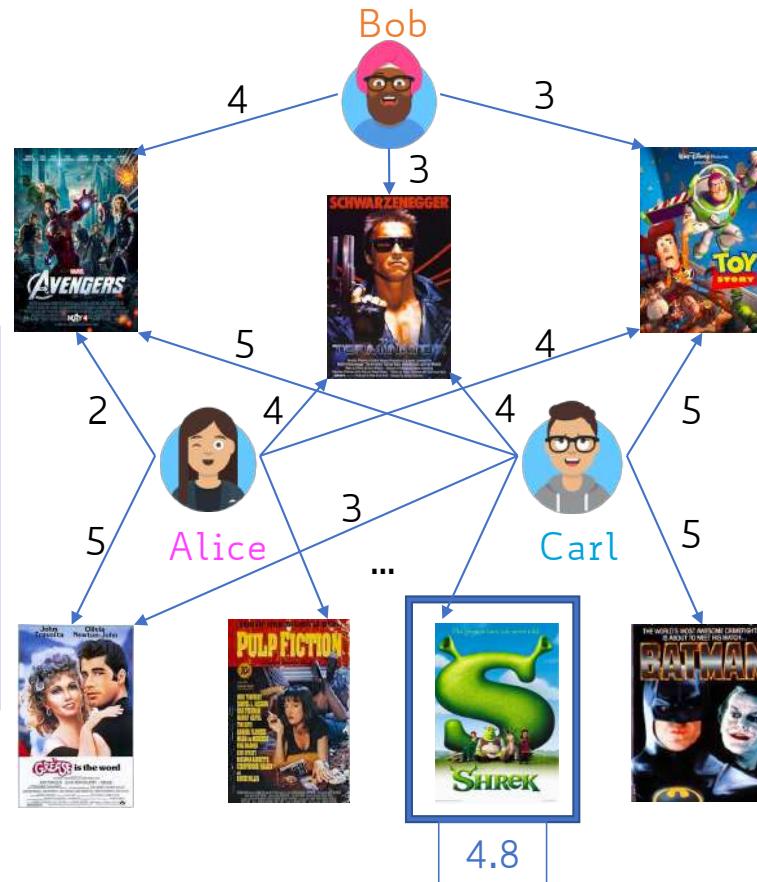
# User-based Neighborhood: Example

Predict the rating that Bob would give to each of those movies on the basis of Alice's and Carl's ratings



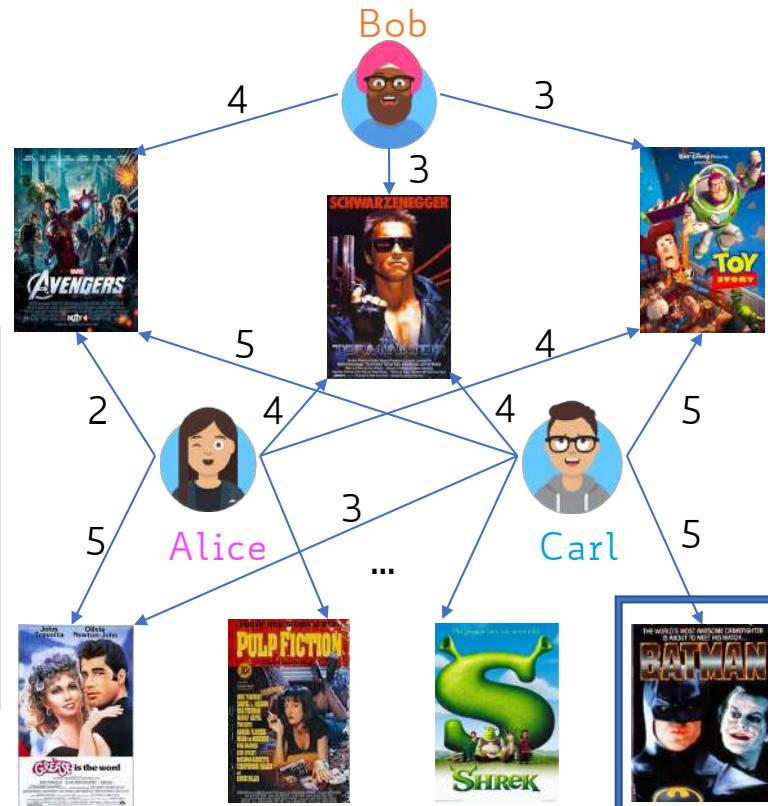
# User-based Neighborhood: Example

Predict the rating that Bob would give to each of those movies on the basis of Alice's and Carl's ratings



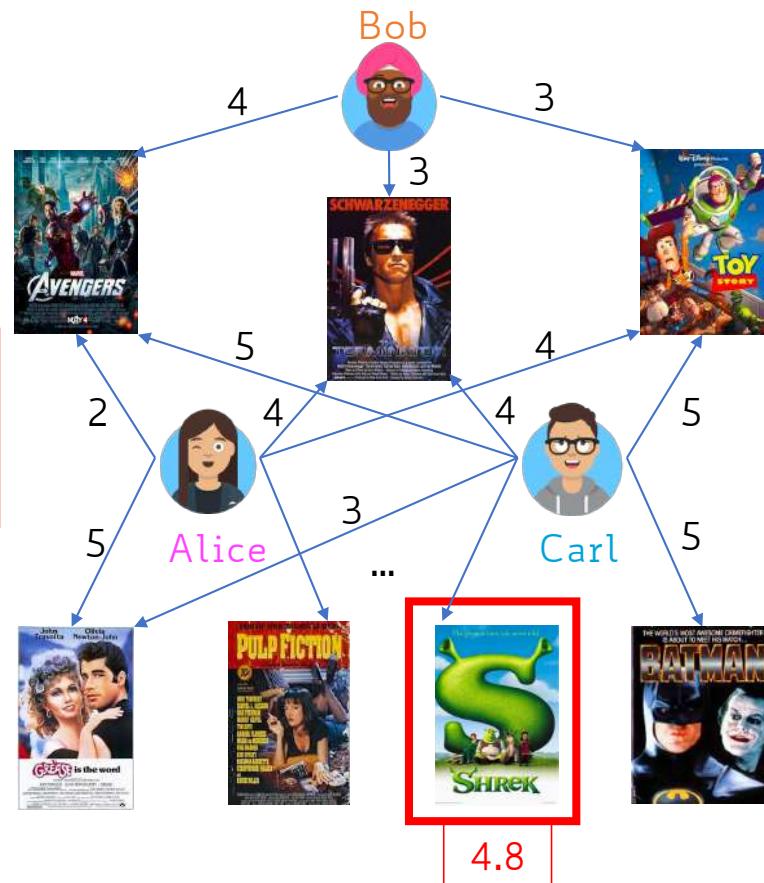
# User-based Neighborhood: Example

Predict the rating that Bob would give to each of those movies on the basis of Alice's and Carl's ratings

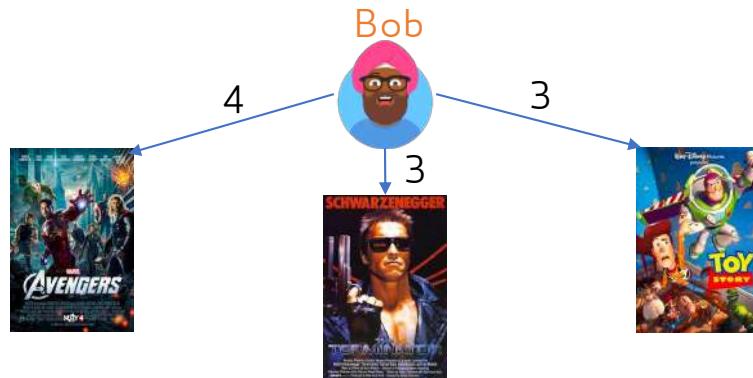


# User-based Neighborhood: Example

Recommend the highest rated movie(s) to Bob!



# User-based Neighborhood: Example



There is no point in predicting the rating of a movie which has only been rated by a user (Zoe) who is **not** in the **Bob's** neighborhood



# User-to-User Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given user  $u$  is the ability of finding users  $u'$  that are "similar" to  $u$

# User-to-User Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given user  $u$  is the ability of finding users  $u'$  that are "similar" to  $u$
- Remember that we are **not** building any content-based user/item profile

# User-to-User Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given user  $u$  is the ability of finding users  $u'$  that are "similar" to  $u$
- Remember that we are **not** building any content-based user/item profile
- Intuitively, 2 users  $u_1$  and  $u_2$  are similar to each other if their ratings (of items) are similar

# User-to-User Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given user  $u$  is the ability of finding users  $u'$  that are "similar" to  $u$
- Remember that we are **not** building any content-based user/item profile
- Intuitively, 2 users  $u_1$  and  $u_2$  are similar to each other if their ratings (of items) are similar
- Each user represented by her/his rating vector and similarity between them is measured in the item (rating) space

# User-to-User Similarity

$\text{sim}(u, v)$  Similarity metric between any pair of users

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# User-to-User Similarity

$\text{sim}(u, v)$  Similarity metric between any pair of users

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

Must capture the intuition:  $\text{sim}(\text{Alice}, \text{Carl}) > \text{sim}(\text{Alice}, \text{Bob})$

# User-to-User Similarity

$\mathbf{r}_u$  n-dimensional vector of ratings provided by user u ( $n = \# \text{movies}$ )

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# User-to-User Similarity

$\mathbf{r}_u$  n-dimensional vector of ratings provided by user u ( $n = \# \text{movies}$ )

		MOVIES							
		AVENGERS	BATMAN	CRUISE IN THE WOOD	PULP FICTION	SHREK	TERMINATOR	WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4				3		3	
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$\mathbf{r}_{\text{Bob}}$

# User-to-User Similarity: Jaccard Similarity

$$\text{sim}(u, v) = J(\mathbf{r}_u, \mathbf{r}_v) = \frac{|\mathbf{r}_u \cap \mathbf{r}_v|}{|\mathbf{r}_u \cup \mathbf{r}_v|}$$

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# User-to-User Similarity: Jaccard Similarity

$$\text{sim}(u, v) = J(\mathbf{r}_u, \mathbf{r}_v) = \frac{|\mathbf{r}_u \cap \mathbf{r}_v|}{|\mathbf{r}_u \cup \mathbf{r}_v|}$$

		MOVIES							
		AVENGERS	BATMAN	ROSE	PULP FICTION	SHREK	TERMINATOR	WOLF	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$$\begin{aligned} \text{sim}(\text{Alice}, \text{Bob}) &= \frac{|\mathbf{r}_{\text{Alice}} \cap \mathbf{r}_{\text{Bob}}|}{|\mathbf{r}_{\text{Alice}} \cup \mathbf{r}_{\text{Bob}}|} \\ &= \frac{3}{6} = 0.5 \end{aligned}$$

# User-to-User Similarity: Jaccard Similarity

$$\text{sim}(u, v) = J(\mathbf{r}_u, \mathbf{r}_v) = \frac{|\mathbf{r}_u \cap \mathbf{r}_v|}{|\mathbf{r}_u \cup \mathbf{r}_v|}$$

		MOVIES							
		AVENGERS	BATMAN	ROOMMATE	PULP FICTION	SHREK	TERMINATOR	WOLF	TOY
USERS	Alice	2		5	4	5	4		4
	Bob		4				3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$$\text{sim}(\text{Alice}, \text{Carl}) = \frac{|\mathbf{r}_{\text{Alice}} \cap \mathbf{r}_{\text{Carl}}|}{|\mathbf{r}_{\text{Alice}} \cup \mathbf{r}_{\text{Carl}}|}$$

$$= \frac{6}{7} \approx 0.86$$

# User-to-User Similarity: Jaccard Similarity

$$\text{sim}(u, v) = J(\mathbf{r}_u, \mathbf{r}_v) = \frac{|\mathbf{r}_u \cap \mathbf{r}_v|}{|\mathbf{r}_u \cup \mathbf{r}_v|}$$

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

**Problem!**  
Jaccard ignores  
rating values

# User-to-User Similarity: Cosine Similarity

$$\text{sim}(u, v) = \text{cosine}(\mathbf{r}_u, \mathbf{r}_v) = \frac{\mathbf{r}_u \cdot \mathbf{r}_v}{\|\mathbf{r}_u\| \|\mathbf{r}_v\|}$$

		MOVIES							
		AVENGERS	BATMAN	ROOMMATE	PULP FICTION	SHREK	TERMINATOR	WOLF	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4				3			3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$$\text{sim}(\text{Alice}, \text{Bob}) = \frac{\mathbf{r}_{\text{Alice}} \cdot \mathbf{r}_{\text{Bob}}}{\|\mathbf{r}_{\text{Alice}}\| \|\mathbf{r}_{\text{Bob}}\|}$$

$$= \frac{32}{\sqrt{102}\sqrt{44}} \approx 0.48$$

# User-to-User Similarity: Cosine Similarity

$$\text{sim}(u, v) = \text{cosine}(\mathbf{r}_u, \mathbf{r}_v) = \frac{\mathbf{r}_u \cdot \mathbf{r}_v}{\|\mathbf{r}_u\| \|\mathbf{r}_v\|}$$

		MOVIES							
		AVENGERS	BATMAN	ROOMMATE	PULP FICTION	SHREK	KONG ISLAND	WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4				3		3	3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$$\begin{aligned} \text{sim}(\text{Alice}, \text{Carl}) &= \frac{\mathbf{r}_{\text{Alice}} \cdot \mathbf{r}_{\text{Carl}}}{\|\mathbf{r}_{\text{Alice}}\| \|\mathbf{r}_{\text{Carl}}\|} \\ &= \frac{102}{\sqrt{102}\sqrt{141}} \approx 0.85 \end{aligned}$$

# User-to-User Similarity: Cosine Similarity

$$\text{sim}(u, v) = \text{cosine}(\mathbf{r}_u, \mathbf{r}_v) = \frac{\mathbf{r}_u \cdot \mathbf{r}_v}{\|\mathbf{r}_u\| \|\mathbf{r}_v\|}$$

		MOVIES							
		AVENGERS	BATMAN	ROSEMARY'S BABY	PULP FICTION	SHREK	KOINNEE LEEWOOD	THE WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

**Problem!**  
Missing rating values  
are treated as 0s  
and have a negative  
effect

# User-to-User Similarity: Pearson Correlation

$$\text{sim}(u, v) = \text{Pearson}(\mathbf{r}_u, \mathbf{r}_v) = \frac{(\mathbf{r}_u - \bar{\mathbf{r}}_u) \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}{\sqrt{(\mathbf{r}_u - \bar{\mathbf{r}}_u)^T \cdot (\mathbf{r}_u - \bar{\mathbf{r}}_u)} \times \sqrt{(\mathbf{r}_v - \bar{\mathbf{r}}_v)^T \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}}$$

		MOVIES							
USERS	Alice	-2		1	0	1	0		0
	Bob	2/3					-1/3		-1/3
	Carl	4/7	4/7	-10/7	-3/7	4/7	-3/7		4/7
	...	...	...	...	...	...	...	...	...
	Zoe			-9/4	-1/4			7/4	-1/4

**Solution:**  
Normalize ratings by subtracting the mean rating

Now 0 means neutral, and if we treat missing ratings as 0, it doesn't mean it's negative

# User-to-User Similarity: Pearson Correlation

$\mathbf{r}'_u = \mathbf{r}_u - \bar{\mathbf{r}}_u$  mean-scaled rating vector of u

$\mathbf{r}'_v = \mathbf{r}_v - \bar{\mathbf{r}}_v$  mean-scaled rating vector of v

# User-to-User Similarity: Pearson Correlation

$\mathbf{r}'_u = \mathbf{r}_u - \bar{\mathbf{r}}_u$  mean-scaled rating vector of u

$\mathbf{r}'_v = \mathbf{r}_v - \bar{\mathbf{r}}_v$  mean-scaled rating vector of v

$$\text{cosine}(\mathbf{r}'_u, \mathbf{r}'_v) = \frac{\mathbf{r}'_u \cdot \mathbf{r}'_v}{\|\mathbf{r}'_u\| \|\mathbf{r}'_v\|} = \frac{(\mathbf{r}_u - \bar{\mathbf{r}}_u) \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}{\|\mathbf{r}_u - \bar{\mathbf{r}}_u\| \|\mathbf{r}_v - \bar{\mathbf{r}}_v\|} =$$

# User-to-User Similarity: Pearson Correlation

$\mathbf{r}'_u = \mathbf{r}_u - \bar{\mathbf{r}}_u$  mean-scaled rating vector of u

$\mathbf{r}'_v = \mathbf{r}_v - \bar{\mathbf{r}}_v$  mean-scaled rating vector of v

$$\begin{aligned} \text{cosine}(\mathbf{r}'_u, \mathbf{r}'_v) &= \frac{\mathbf{r}'_u \cdot \mathbf{r}'_v}{\|\mathbf{r}'_u\| \|\mathbf{r}'_v\|} = \frac{(\mathbf{r}_u - \bar{\mathbf{r}}_u) \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}{\|\mathbf{r}_u - \bar{\mathbf{r}}_u\| \|\mathbf{r}_v - \bar{\mathbf{r}}_v\|} = \\ &= \frac{(\mathbf{r}_u - \bar{\mathbf{r}}_u) \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}{\sqrt{(\mathbf{r}_u - \bar{\mathbf{r}}_u)^T \cdot (\mathbf{r}_u - \bar{\mathbf{r}}_u)} \times \sqrt{(\mathbf{r}_v - \bar{\mathbf{r}}_v)^T \cdot (\mathbf{r}_v - \bar{\mathbf{r}}_v)}} = \text{Pearson}(\mathbf{r}_u, \mathbf{r}_v) \end{aligned}$$

# User-based Neighborhood: Predictions

$\mathbf{r}_u$  Vector of ratings provided by user  $u$

# User-based Neighborhood: Predictions

$\mathbf{r}_u$  Vector of ratings provided by user  $u$

$$\mathcal{U}^k = \operatorname{argmax}_{\mathcal{U}' \subseteq \mathcal{U} \setminus u, |\mathcal{U}'|=k} \sum_{u' \in \mathcal{U}'} \operatorname{sim}(u, u')$$

Top-k most "similar" users to  $u$   
 $u$ 's k-neighborhood

# User-based Neighborhood: Predictions

$\mathbf{r}_u$  Vector of ratings provided by user  $u$

$$\mathcal{U}^k = \operatorname{argmax}_{\mathcal{U}' \subseteq \mathcal{U} \setminus u, |\mathcal{U}'|=k} \sum_{u' \in \mathcal{U}'} \operatorname{sim}(u, u')$$

Top-k most "similar" users to  $u$   
 $u$ 's k-neighborhood

Set of items rated by  $u$ 's neighbors

$$\mathcal{I}^k = \{i \in \mathcal{I} : \mathbf{r}_{u',i} = \downarrow \wedge u' \in \mathcal{U}^k\}$$

# User-based Neighborhood: Predictions

$\mathbf{r}_u$  Vector of ratings provided by user  $u$

$$\mathcal{U}^k = \operatorname{argmax}_{\mathcal{U}' \subseteq \mathcal{U} \setminus u, |\mathcal{U}'|=k} \sum_{u' \in \mathcal{U}'} \operatorname{sim}(u, u')$$

Top-k most "similar" users to  $u$   
 $u$ 's k-neighborhood

Set of items rated by  $u$ 's neighbors

$$\mathcal{I}^k = \{i \in \mathcal{I} : \mathbf{r}_{u',i} \neq \downarrow \wedge u' \in \mathcal{U}^k\}$$

Predicted rating given by user  $u$  to item  $i$

$$\mathbf{r}_u[i] = r(u, i) = r_{u,i}$$

# User-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

# User-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

$$\forall i \in \mathcal{I}^k$$

# User-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

$$\forall i \in \mathcal{I}^k$$

$$r_{u,i} = \frac{1}{k} \sum_{u' \in \mathcal{U}^k} r_{u',i}$$

plain average

# User-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

$$\forall i \in \mathcal{I}^k$$

$$r_{u,i} = \frac{1}{k} \sum_{u' \in \mathcal{U}^k} r_{u',i}$$

plain average

$$r_{u,i} = \frac{1}{k} \sum_{u' \in \mathcal{U}^k} \text{sim}(u, u') \cdot r_{u',i}$$

weighted average

# User-based CF: Drawbacks

3 main issues with user-based CF

# User-based CF: Drawbacks

3 main issues with user-based CF

**Sparsity**  
systems performed poorly when they had many items but comparatively few ratings

# User-based CF: Drawbacks

3 main issues with user-based CF

**Sparsity**  
systems performed poorly when they had many items but comparatively few ratings

**Efficiency**  
computing similarities between all pairs of users is expensive

# User-based CF: Drawbacks

3 main issues with user-based CF

**Sparsity**  
systems performed poorly when they had many items but comparatively few ratings

**Efficiency**  
computing similarities between all pairs of users is expensive

**Aging**  
user profiles changed quickly and the entire system model had to be recomputed

# ITEM-BASED COLLABORATIVE FILTERING

# Item-based CF

- Introduced by Amazon to overcome the 3 issues of user-based CF

# Item-based CF

- Introduced by Amazon to overcome the 3 issues of user-based CF
- Since systems have typically more users than items, each item has more ratings than each user

# Item-based CF

- Introduced by Amazon to overcome the 3 issues of user-based CF
- Since systems have typically more users than items, each item has more ratings than each user
- As such, an item's rating average is **more stable** over time

# Item-based CF

- Introduced by Amazon to overcome the 3 issues of user-based CF
- Since systems have typically more users than items, each item has more ratings than each user
- As such, an item's rating average is **more stable** over time
- The model doesn't suffer from aging and therefore it does not need to be recomputed frequently

# Item-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

# Item-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of items already rated by  $u$  ( $I_u$ )  
extract a subset of  $k$  neighbours of  $i$

# Item-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of items already rated by  $u$  ( $I_u$ )  
extract a subset of  $k$  neighbours of  $i$

$k$ -neighborhood is found on the basis of the similarity  
between items without explicit item's content or metadata

# Item-based Neighborhood

Given a user  $u$  and an item  $i$  not rated by  $u$ , we want to estimate  $r(u, i)$

From the set of items already rated by  $u$  ( $I_u$ )  
extract a subset of  $k$  neighbours of  $i$

$k$ -neighborhood is found on the basis of the similarity  
between items without explicit item's content or metadata

Estimate  $r(u, i)$  based on the ratings of items in the  
 $k$ -neighborhood of  $i$

# Item-to-Item Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given item  $i$  is the ability of finding items that are "similar" to  $i$

# Item-to-Item Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given item  $i$  is the ability of finding items that are "similar" to  $i$
- Remember that we are not building any content-based user/item profile

# Item-to-Item Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given item  $i$  is the ability of finding items that are "similar" to  $i$
- Remember that we are not building any content-based user/item profile
- Intuitively, 2 items  $i_1$  and  $i_2$  are similar to each other if users who rated them are similar

# Item-to-Item Similarity

- The key "trick" to discover the  $k$ -neighborhood of a given item  $i$  is the ability of finding items that are "similar" to  $i$
- Remember that we are not building any content-based user/item profile
- Intuitively, 2 items  $i_1$  and  $i_2$  are similar to each other if users who rated them are similar
- Each item represented by the user ratings vector and similarity between them is measured in the user (rating) space

# Item-to-Item Similarity

$\mathbf{r}_i$  m-dimensional vector of ratings provided for item i (m = #users)

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# Item-to-Item Similarity

$\mathbf{r}_i$  m-dimensional vector of ratings provided for item i (m = #users)

		MOVIES							
		AVENGERS	BATMAN	FRIDAY	PULP FICTION	SHREK	SCHWANZENBERGER	WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4				3			3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$\mathbf{r}_{\text{Shrek}}$

# Item-based Neighborhood: Example

Let's consider again Bob!

		MOVIES							
USERS	Alice	2		5	4	5	4	4	
	Bob	4				3		3	
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# Item-based Neighborhood: Example

Suppose we want to predict the rating Bob would give to [Shrek](#)

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4				?	3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

# Item-based Neighborhood: Example

We first extract the subset of k most similar items to Shrek which have been rated by Bob

		MOVIES							
		AVENGERS	BATMAN	CELESTE	PULP FICTION	SHREK	SCHWARZENEGGER	WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4					3		3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe		1	3				5	4

$r_{\text{Shrek}}$

# Item-based Neighborhood: Example

Suppose those are: The Avengers and The Terminator

		MOVIES							
USERS	Alice	2		5	4	5	4		4
	Bob	4				3			3
	Carl	5	5	3	4	5	4		5
	...	...	...	...	...	...	...	...	...
	Zoe	1	3					5	4

For example, item similarity is measured using Pearson's correlation

# Item-based Neighborhood: Example

The predicted rating is computed as an aggregating function of the ratings that Bob gave to the k most similar movies to Shrek

		MOVIES							
		AVENGERS	BATMAN	CELOPHANE	PULP FICTION	SHREK	SCHWARZENEGGER	WOLF OF WALL STREET	TOY STORY
USERS	Alice	2		5	4	5	4		4
	Bob	4				?	3		3
	Carl	5		5	3	4	5		5
	...	...	...	...	...	...	...	...	...
	Zoe	1	3					5	4

# Item-based Neighborhood: Predictions

$\mathbf{r}_i$  Vector of ratings given to item i

# Item-based Neighborhood: Predictions

$\mathbf{r}_i$  Vector of ratings given to item i

$$\mathcal{I}_u = \{i \in \mathcal{I} : r_{u,i} \neq \downarrow\}$$

Set of items rated by u

# Item-based Neighborhood: Predictions

$\mathbf{r}_i$  Vector of ratings given to item i

$$\mathcal{I}_u = \{i \in \mathcal{I} : r_{u,i} \neq \text{None}\}$$

Set of items rated by u

$$\mathcal{I}_u^k = \operatorname{argmax}_{\mathcal{I}'_u \subseteq \mathcal{I}_u, |\mathcal{I}'_u|=k} \sum_{i' \in \mathcal{I}'_u} \text{sim}(i, i')$$

Top-k most "similar" items  
to i among those rated by u  
i's k-neighborhood

# Item-based Neighborhood: Predictions

$\mathbf{r}_i$  Vector of ratings given to item i

$$\mathcal{I}_u = \{i \in \mathcal{I} : r_{u,i} \neq \downarrow\}$$

Set of items rated by u

$$\mathcal{I}_u^k = \operatorname{argmax}_{\mathcal{I}'_u \subseteq \mathcal{I}_u, |\mathcal{I}'_u| = k} \sum_{i' \in \mathcal{I}'_u} \text{sim}(i, i')$$

Top-k most "similar" items  
to i among those rated by u  
i's k-neighborhood

Predicted rating given by user u to item i

$$\mathbf{r}_u[i] = r(u, i) = r_{u,i}$$

# Item-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

# Item-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

$$\forall i' \in \mathcal{I}_u^k$$

$$r_{u,i} = \frac{1}{k} \sum_{i' \in \mathcal{I}_u^k} r_{u,i'}$$

plain average

# Item-based Neighborhood: Predictions

2 possible ways of aggregating neighbors ratings

$$\forall i' \in \mathcal{I}_u^k$$

$$r_{u,i} = \frac{1}{k} \sum_{i' \in \mathcal{I}_u^k} r_{u,i'}$$

plain average

$$r_{u,i} = \frac{1}{k} \sum_{i' \in \mathcal{I}_u^k} \text{sim}(i, i') \cdot r_{u,i'}$$

weighted average

# Item-to-Item Collaborative Filtering

- Item similarity can be computed from rating vectors (in user space)

# Item-to-Item Collaborative Filtering

- Item similarity can be computed from rating vectors (in user space)
- Analogous to user similarity of rating vectors (in item space):
  - Jaccard index
  - Cosine similarity (normalized = Pearson's correlation)

# Item-to-Item Collaborative Filtering

- Item similarity can be computed from rating vectors (in user space)
- Analogous to user similarity of rating vectors (in item space):
  - Jaccard index
  - Cosine similarity (normalized = Pearson's correlation)
- Rating prediction using the same methods proposed for user-based CF
  - Plain average of ratings
  - Weighted average of ratings (taking item similarity into account)

# Item-to-Item Collaborative Filtering

In general, item-based works better than user-based CF

# Memory-based CF: Implementation

- The most expensive step is finding the  $k$  most similar users (or items)

# Memory-based CF: Implementation

- The most expensive step is finding the  $k$  most similar users (or items)
- This computation is too expensive to do online (for every user/item)

# Memory-based CF: Implementation

- The most expensive step is finding the  $k$  most similar users (or items)
- This computation is too expensive to do online (for every user/item)
- Finding the  $k$  most similar users/items should be pre-computed (offline)

# Memory-based CF: Implementation

- The most expensive step is finding the  $k$  most similar users (or items)
- This computation is too expensive to do online (for every user/item)
- Finding the  $k$  most similar users/items should be pre-computed (offline)
- $k$ -nearest neighbors search in high dimensions (i.e., quickly find the set of  $k$  nearest data points)

# Memory-based CF: Implementation

The curse of dimensionality (again!)



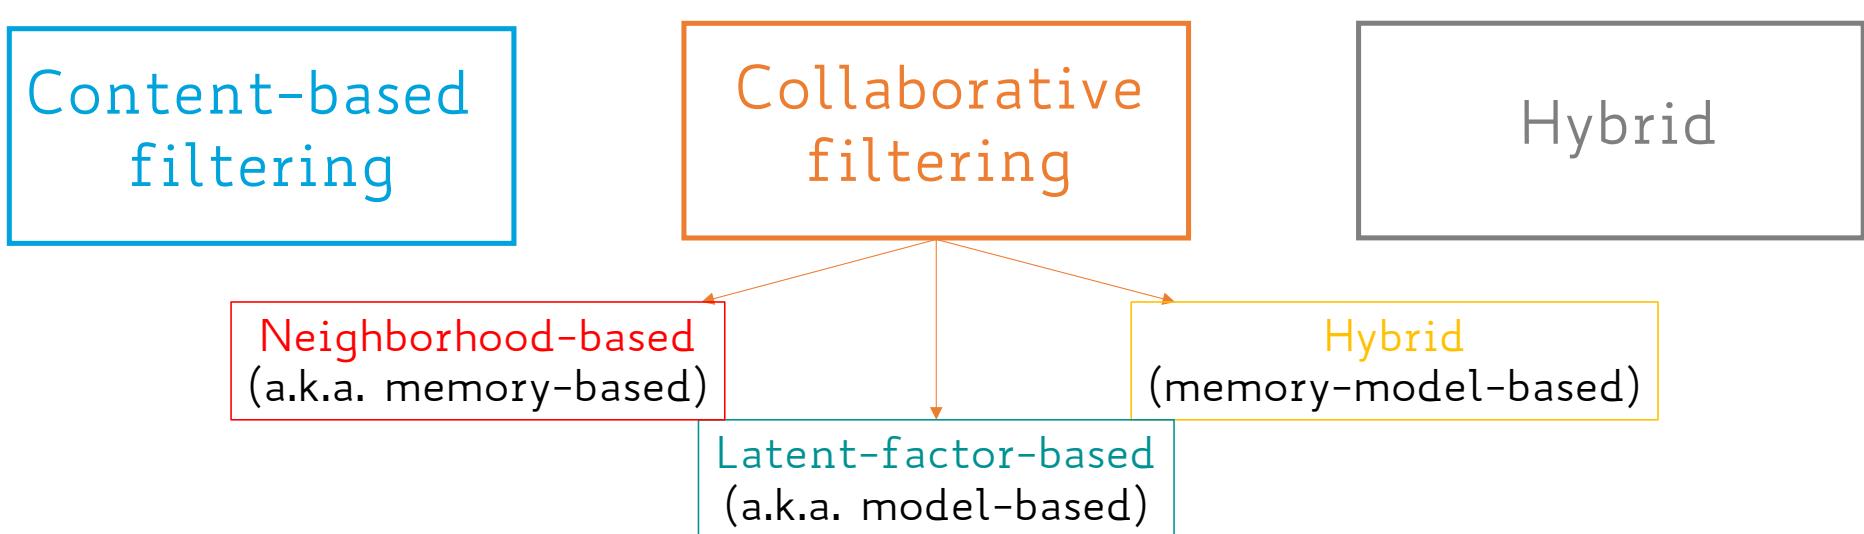
# Memory-based CF: Implementation

Locality-Sensitive Hashing (LSH) approximation



# Recommendation Strategies

3 approaches to recommender systems



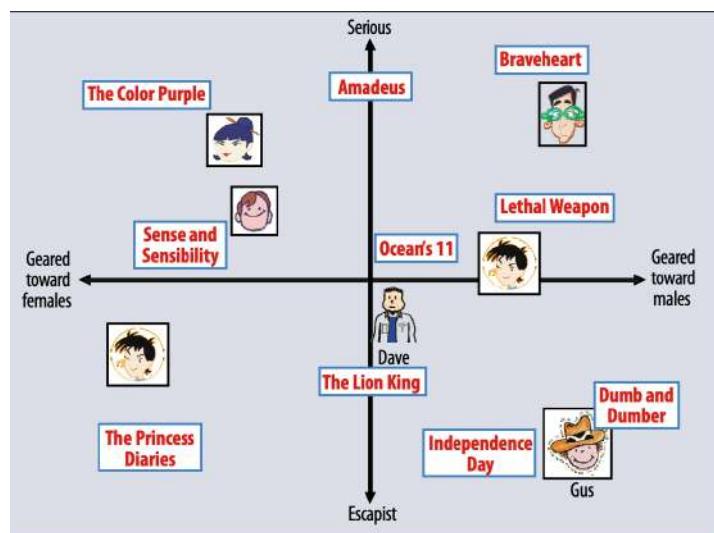
# LATENT FACTOR MODELS

# Latent Factor (Model-based) CF

Tries to predict ratings by representing both items and users with a number of **hidden factors** inferred from observed ratings

# Latent Factor (Model-based) CF

Tries to predict ratings by representing both items and users with a number of **hidden factors** inferred from observed ratings

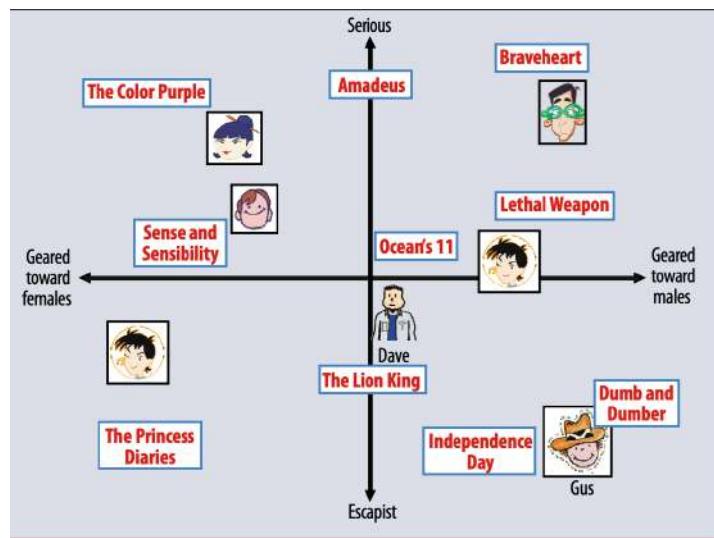


Example: 2 hidden factors

- Dim. 1: Male vs. Female
- Dim. 2: Serious vs. Escapist

# Latent Factor (Model-based) CF

Tries to predict ratings by representing both items and users with a number of **hidden factors** inferred from observed ratings



**Example: 2 hidden factors**

- Dim. 1: Male vs. Female
- Dim. 2: Serious vs. Escapist

A user's predicted rating for an item (movie) would equal the **dot product** of the movie and user vectors on the plot

# Matrix Factorization

- Some of the most successful realizations of latent factor models are based on **matrix factorization** (MF)

# Matrix Factorization

- Some of the most successful realizations of latent factor models are based on **matrix factorization** (MF)
- The original idea behind MF is to represent users and items in a lower dimensional latent space (i.e., as **vectors of latent factors**)

# Matrix Factorization

- Some of the most successful realizations of latent factor models are based on **matrix factorization** (MF)
- The original idea behind MF is to represent users and items in a lower dimensional latent space (i.e., as **vectors of latent factors**)
- Such vectors are inferred (i.e., learned) from observed item ratings

# Matrix Factorization

- Some of the most successful realizations of latent factor models are based on **matrix factorization** (MF)
- The original idea behind MF is to represent users and items in a lower dimensional latent space (i.e., as **vectors of latent factors**)
- Such vectors are inferred (i.e., learned) from observed item ratings
- High correspondence between item and user factors leads to a recommendation

# Matrix Factorization Framework

- Map both items and users to a **joint latent factor**  $d$ -dimensional space

# Matrix Factorization Framework

- Map both items and users to a **joint latent factor**  $d$ -dimensional space
- User-Item interactions are modeled as **inner products** in that space

# Matrix Factorization Framework

- Map both items and users to a **joint latent factor**  $d$ -dimensional space
- User-Item interactions are modeled as **inner products** in that space

$\mathbf{x}_u \in \mathbb{R}^d$   $d$ -dimensional vector representing **user**  $u$

Each  $x_u[k]$  measures the extent of interest user  $u$  has in items exhibiting the  $k$ -th factor

# Matrix Factorization Framework

- Map both items and users to a **joint latent factor**  $d$ -dimensional space
- User-Item interactions are modeled as **inner products** in that space

$\mathbf{x}_u \in \mathbb{R}^d$  d-dimensional vector representing **user**  $u$   
 $\mathbf{w}_i \in \mathbb{R}^d$  d-dimensional vector representing **item**  $i$

Each  $x_u[k]$  measures the extent of interest user  $u$  has in items exhibiting the  $k$ -th factor

Each  $w_i[k]$  measures the extent to which the item  $i$  has the  $k$ -th factor

# What Are Those Factors?

- Essentially,  $d$  hidden features for describing both users and items

# What Are Those Factors?

- Essentially,  $d$  hidden features for describing both users and items
- In the user-movie example, a feature  $f$  may refer to:
  - how much each a user likes Disney movies (in the case of user vectors)
  - how close a movie is to a Disney movie (in the case of item, i.e., movie, vectors)

# What Are Those Factors?

- Essentially,  $d$  hidden features for describing both users and items
- In the user-movie example, a feature  $f$  may refer to:
  - how much each a user likes Disney movies (in the case of user vectors)
  - how close a movie is to a Disney movie (in the case of item, i.e., movie, vectors)
- We do not know what these features are nor do we have to determine them beforehand!

# What Are Those Factors?

- Essentially,  $d$  hidden features for describing both users and items
- In the user-movie example, a feature  $f$  may refer to:
  - how much each a user likes Disney movies (in the case of user vectors)
  - how close a movie is to a Disney movie (in the case of item, i.e., movie, vectors)
- We do not know what these features are nor do we have to determine them beforehand!
- That is why these features are often refer to as **latent features**

# Matrix Factorization Framework

$r(u, i) = r_{u,i}$  rating of user  $u$  for the item  $i$

# Matrix Factorization Framework

$r(u, i) = r_{u,i}$  rating of user  $u$  for the item  $i$

$\hat{r}_{u,i} = \mathbf{x}_u^T \cdot \mathbf{w}_i = \sum_{j=1}^d x_{u,j} w_{j,i}$  estimated (i.e., predicted) rating of user  $u$  for the item  $i$

# Matrix Factorization Framework

$r(u, i) = r_{u,i}$  rating of user  $u$  for the item  $i$

$\hat{r}_{u,i} = \mathbf{x}_u^T \cdot \mathbf{w}_i = \sum_{j=1}^d x_{u,j} w_{j,i}$  estimated (i.e., predicted) rating of user  $u$  for the item  $i$

The major challenge is computing the mapping of each item and user to latent factor vectors  $\mathbf{x}_u$  and  $\mathbf{w}_i$

# Matrix Factorization Framework

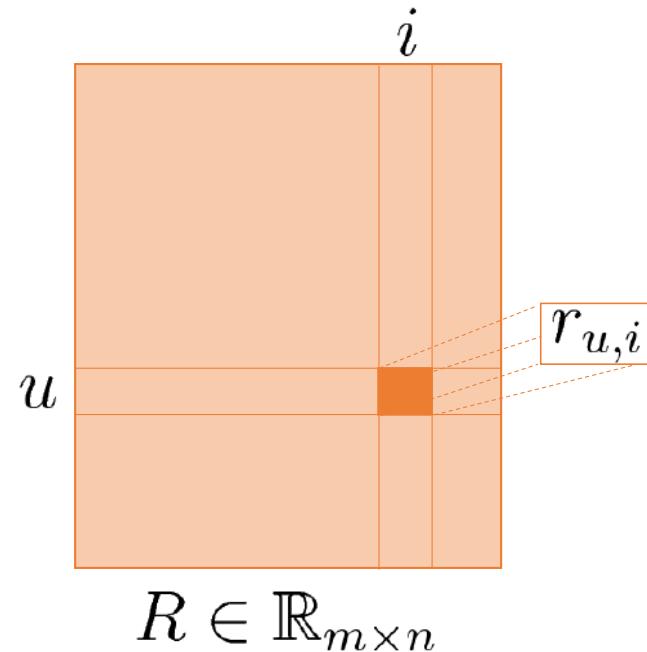
$r(u, i) = r_{u,i}$  rating of user  $u$  for the item  $i$

$\hat{r}_{u,i} = \mathbf{x}_u^T \cdot \mathbf{w}_i = \sum_{j=1}^d x_{u,j} w_{j,i}$  estimated (i.e., predicted) rating of user  $u$  for the item  $i$

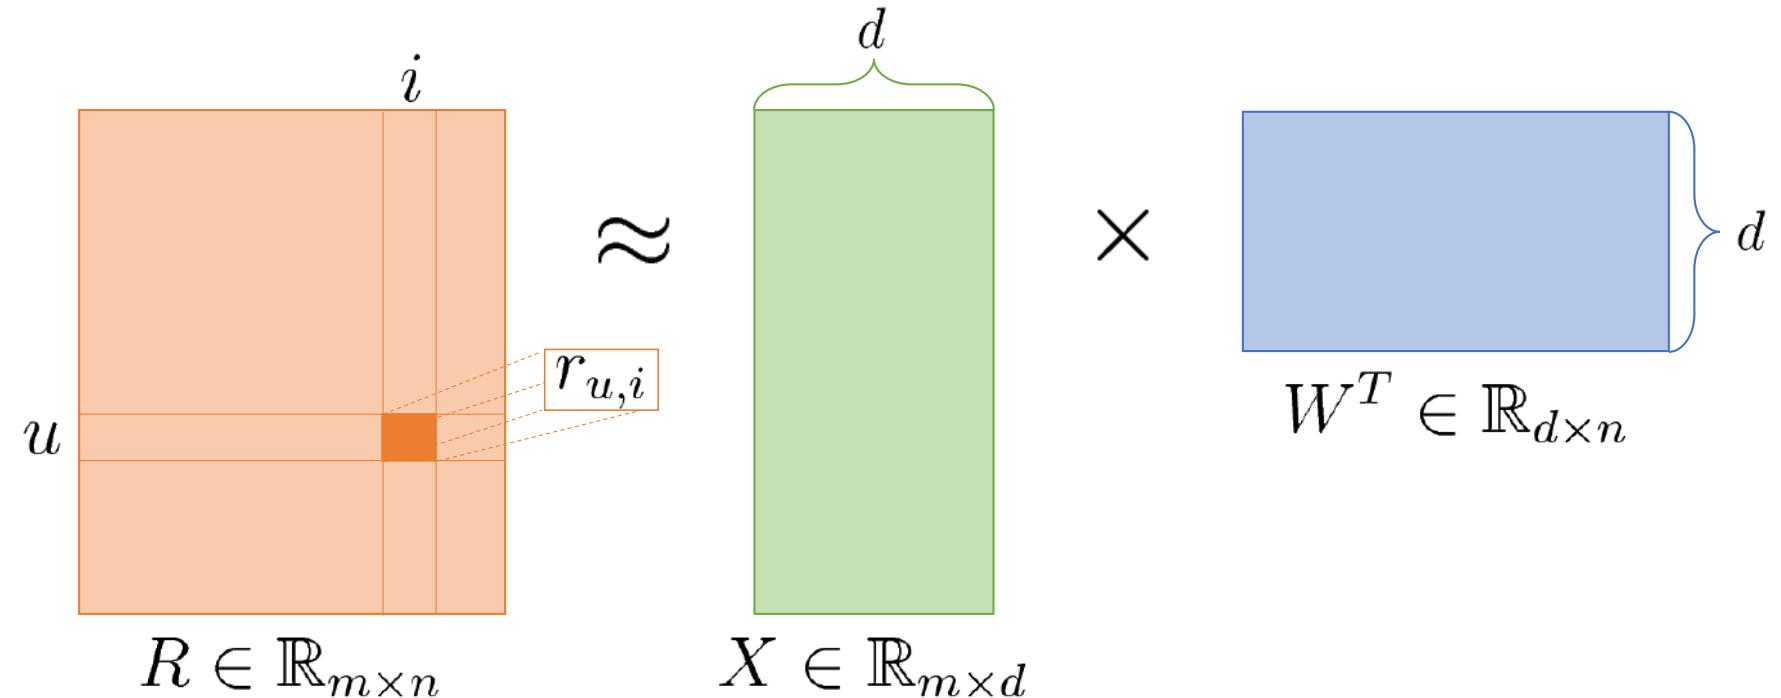
The major challenge is computing the mapping of each item and user to latent factor vectors  $\mathbf{x}_u$  and  $\mathbf{w}_i$

Recommendations for a user are generated by computing the estimated ratings for unseen items, and by taking the top-k highest rated ones

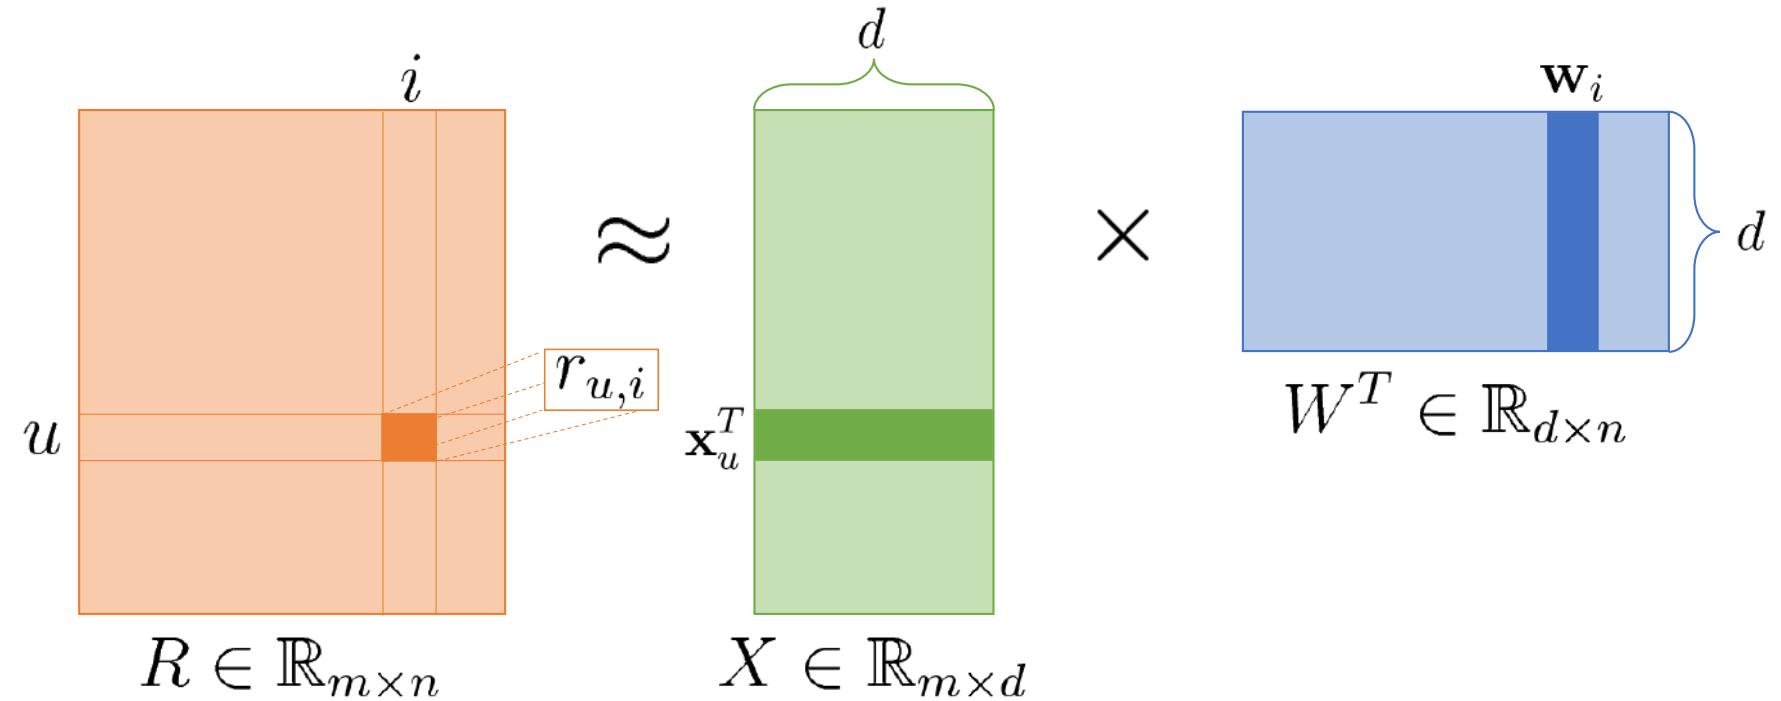
# Matrix Factorization Framework



# Matrix Factorization Framework



# Matrix Factorization Framework



Approximate the user-item rating matrix  $R$  with the product of  $X \times W^T$

# How Do We Learn $X$ and $W$ ?

Assuming we have access to a dataset of observed ratings

# How Do We Learn $X$ and $W$ ?

Assuming we have access to a dataset of observed ratings

The matrix  $R$  is partially known and filled with those observations

# How Do We Learn X and W?

Assuming we have access to a dataset of observed ratings

The matrix R is partially known and filled with those observations

To actually learn the latent factor representations  $\mathbf{x}_u$  and  $\mathbf{w}_i$  we minimize the following loss function

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

Training set of  
observed ratings

# How Do We Learn X and W?

Assuming we have access to a dataset of observed ratings

The matrix R is partially known and filled with those observations

To actually learn the latent factor representations  $\mathbf{x}_u$  and  $\mathbf{w}_i$  we minimize the following loss function

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

squared error term

↑

Training set of observed ratings

12/11/2024 129

# How Do We Learn X and W?

Assuming we have access to a **dataset** of **observed ratings**

The matrix R is **partially known** and filled with those observations

To actually learn the latent factor representations  $\mathbf{x}_u$  and  $\mathbf{w}_i$  we minimize the following loss function

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

squared error term      regularization term

↑

Training set of observed ratings

12/11/2024

130

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right) \right\}$$

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right) \right\}$$

Mathematically convenient

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right) \right\}$$

Still, how do we solve this?

# Take-Home Message of Today

- Collaborative-filtering may require computing the all-pair user/item similarity

# Take-Home Message of Today

- Collaborative-filtering may require computing the all-pair user/item similarity
- In large scale systems, this must be pre-computed offline

# Take-Home Message of Today

- Collaborative-filtering may require computing the all-pair user/item similarity
- In large scale systems, this must be pre-computed offline
- At inference time, make use of ad hoc data structures (e.g., k-d trees) to efficiently compute the set of (approximated) nearest neighbors for a query user/item

# Take-Home Message of Today

- Collaborative-filtering may require computing the all-pair user/item similarity
- In large scale systems, this must be pre-computed offline
- At inference time, make use of ad hoc data structures (e.g., k-d trees) to efficiently compute the set of (approximated) nearest neighbors for a query user/item
- Latent Factor Models overcome this need