

# Big Data Computing

## Master's Degree in Computer Science

### 2025-2026

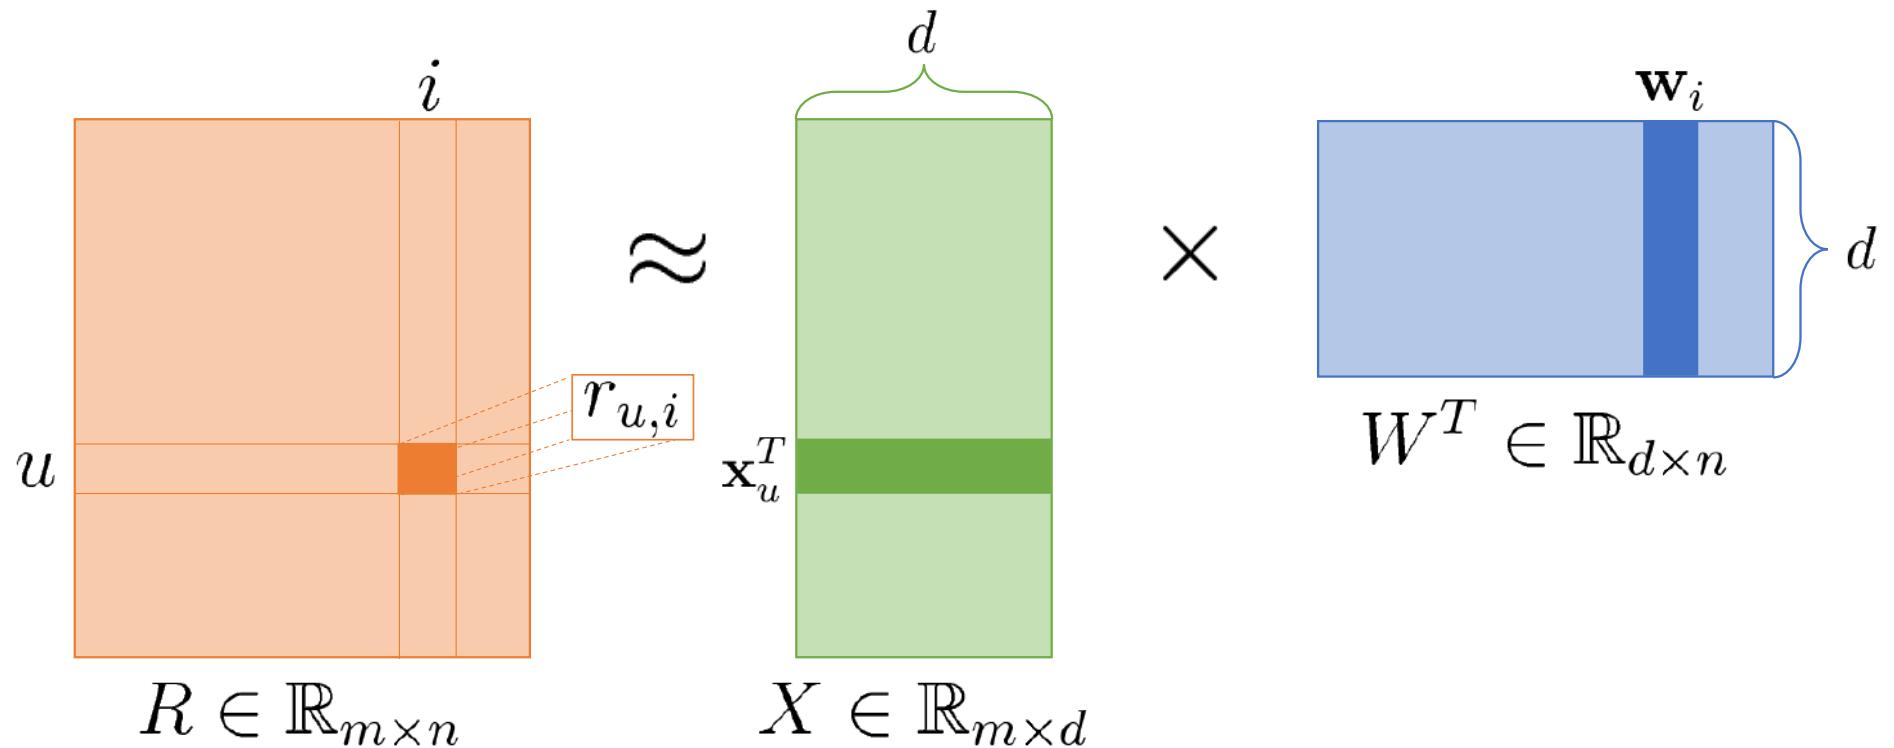


SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science  
Sapienza Università di Roma  
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Matrix Factorization Framework



Approximate the user-item rating matrix  $R$  with the product of  $X \times W^T$

# How Do We Learn X and W?

Assuming we have access to a dataset of observed ratings

The matrix R is partially known and filled with those observations

To actually learn the latent factor representations  $\mathbf{x}_u$  and  $\mathbf{w}_i$  we minimize the following loss function

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

↑  
Training set of observed ratings

squared error term

regularization term

# Matrix Factorization: Optimization

$$X^*, W^* = \operatorname{argmin}_{X,W} L(X, W)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right) \right\}$$

Still, how do we solve this?

# Learning Algorithms

2 main optimization methods

# Learning Algorithms

2 main optimization methods

Stochastic Gradient Descent (SGD)

# Learning Algorithms

2 main optimization methods

```
graph TD; A[2 main optimization methods] -- blue arrow --> B[Stochastic Gradient Descent (SGD)]; A -- green arrow --> C[Alternating Least Squares (ALS)]
```

Stochastic Gradient Descent (SGD) Alternating Least Squares (ALS)

# Stochastic Gradient Descent (SGD)

For each training instance  $(u, i)$ , let's compute the gradient of the loss with respect to  $x_u$  and  $w_i$

# Stochastic Gradient Descent (SGD)

For each training instance  $(u, i)$ , let's compute the gradient of the loss with respect to  $x_u$  and  $w_i$

$$\nabla L(\mathbf{x}_u; \mathbf{w}_i) = \frac{1}{2} \left[ -2(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + 2\lambda \mathbf{x}_u \right] = -(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u$$

$$\nabla L(\mathbf{w}_i; \mathbf{x}_u) = \frac{1}{2} \left[ -2(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{x}_u + 2\lambda \mathbf{w}_i \right] = -(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{x}_u + \lambda \mathbf{w}_i$$

# Stochastic Gradient Descent (SGD)

We know that the updating strategy for SGD is as follows:

$$\mathbf{x}_u^{(t+1)} \leftarrow \mathbf{x}_u^{(t)} - \eta \nabla L(\mathbf{x}_u^{(t)}; \mathbf{w}_i^{(t)})$$

$$\mathbf{w}_i^{(t+1)} \leftarrow \mathbf{w}_i^{(t)} - \eta \nabla L(\mathbf{w}_i^{(t)}; \mathbf{x}_u^{(t)})$$

# Stochastic Gradient Descent (SGD)

We know that the updating strategy for SGD is as follows:

$$\begin{aligned}\mathbf{x}_u^{(t+1)} &\leftarrow \mathbf{x}_u^{(t)} - \eta \nabla L(\mathbf{x}_u^{(t)}; \mathbf{w}_i^{(t)}) & \mathbf{x}_u^{(0)}, \mathbf{w}_i^{(0)} \\ \mathbf{w}_i^{(t+1)} &\leftarrow \mathbf{w}_i^{(t)} - \eta \nabla L(\mathbf{w}_i^{(t)}; \mathbf{x}_u^{(t)}) & \text{typically randomly initialized}\end{aligned}$$

# Stochastic Gradient Descent (SGD)

We know that the updating strategy for SGD is as follows:

$$\begin{aligned}\mathbf{x}_u^{(t+1)} &\leftarrow \mathbf{x}_u^{(t)} - \eta \nabla L(\mathbf{x}_u^{(t)}; \mathbf{w}_i^{(t)}) & \mathbf{x}_u^{(0)}, \mathbf{w}_i^{(0)} \\ \mathbf{w}_i^{(t+1)} &\leftarrow \mathbf{w}_i^{(t)} - \eta \nabla L(\mathbf{w}_i^{(t)}; \mathbf{x}_u^{(t)}) & \text{typically randomly initialized}\end{aligned}$$

At each iteration, both user and item latent vectors are updated by a magnitude proportional to  $\eta$  in the **opposite direction** of the gradient

# Stochastic Gradient Descent (SGD)

We define the **prediction error** associated with each training instance  $(u, i)$

$$e_{u,i} = r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i$$

# Stochastic Gradient Descent (SGD)

We define the **prediction error** associated with each training instance  $(u, i)$

$$e_{u,i} = r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i$$

$$\nabla L(\mathbf{x}_u; \mathbf{w}_i) = -(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = \boxed{-e_{u,i} \mathbf{w}_i + \lambda \mathbf{x}_u}$$

$$\mathbf{x}_u^{(t+1)} \leftarrow \mathbf{x}_u^{(t)} - \eta \boxed{\nabla L(\mathbf{x}_u^{(t)}; \mathbf{w}_i^{(t)})}$$

$$\mathbf{x}_u^{(t+1)} \leftarrow \mathbf{x}_u^{(t)} + \eta (e_{u,i} \mathbf{w}_i^{(t)} - \lambda \mathbf{x}_u^{(t)}) \boxed{\quad}$$

# Stochastic Gradient Descent (SGD)

We define the **prediction error** associated with each training instance  $(u, i)$

$$e_{u,i} = r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i$$

$$\nabla L(\mathbf{w}_i; \mathbf{x}_u) = -(r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{x}_u + \lambda \mathbf{w}_i = \boxed{-e_{u,i} \mathbf{x}_u + \lambda \mathbf{w}_i}$$

$$\mathbf{w}_i^{(t+1)} \leftarrow \mathbf{w}_i^{(t)} - \eta \nabla L(\mathbf{w}_i^{(t)}; \mathbf{x}_u^{(t)})$$

$$\boxed{\mathbf{w}_i^{(t+1)} \leftarrow \mathbf{w}_i^{(t)} + \eta (e_{u,i} \mathbf{x}_u^{(t)} - \lambda \mathbf{w}_i^{(t)})}$$

# Stochastic Gradient Descent (SGD)

- Key Properties:

- Incremental → updates occur sample by sample
- Easy to implement
- Scales well when data is sparse
- Order matters → the sequence of updates influences convergence
- Learning rate ( $\eta$ ) must be tuned carefully
- Converges to a local optimum, but sometimes slowly

# Stochastic Gradient Descent (SGD)

- Good for:

- Very large, sparse datasets (Netflix-like settings)
- Online/streaming updates (new ratings can be incorporated incrementally)
- Systems where simplicity and speed per update matter

# Stochastic Gradient Descent (SGD)

- SGD has been shown to work well optimizing MF models

# Stochastic Gradient Descent (SGD)

- SGD has been shown to work well optimizing MF models
- However, it is not a popular choice if the dimensionality of the original rating matrix  $R$  is high

# Stochastic Gradient Descent (SGD)

- SGD has been shown to work well optimizing MF models
- However, it is not a popular choice if the dimensionality of the original rating matrix  $R$  is high
- Indeed, there are  $d(m+n)$  parameters to optimize

# Stochastic Gradient Descent (SGD)

- SGD has been shown to work well optimizing MF models
- However, it is not a popular choice if the dimensionality of the original rating matrix  $R$  is high
- Indeed, there are  $d(m+n)$  parameters to optimize
- In real life problems, this number can get very large quite often, requiring both a parallelization mechanism or an alternative optimizer

# Alternative Least Squares (ALS): Intuition

- The original objective is **non-convex**, as both  $x_u$  and  $w_i$  are unknown

# Alternative Least Squares (ALS): Intuition

- The original objective is **non-convex**, as both  $\mathbf{x}_u$  and  $\mathbf{w}_i$  are unknown
- ALS alternately fixes (i.e., assumes constant) one latent vector (e.g., item vector) and updates the other one (e.g., user vector)

# Alternative Least Squares (ALS): Intuition

- The original objective is **non-convex**, as both  $\mathbf{x}_u$  and  $\mathbf{w}_i$  are unknown
- ALS alternately fixes (i.e., assumes constant) one latent vector (e.g., item vector) and updates the other one (e.g., user vector)
- When one latent vector is fixed, the objective becomes quadratic (i.e., convex) and therefore can be solved optimally

# Alternative Least Squares (ALS): Intuition

- The original objective is **non-convex**, as both  $\mathbf{x}_u$  and  $\mathbf{w}_i$  are unknown
- ALS alternately fixes (i.e., assumes constant) one latent vector (e.g., item vector) and updates the other one (e.g., user vector)
- When one latent vector is fixed, the objective becomes quadratic (i.e., convex) and therefore can be solved optimally
- Each alternating iteration reduces to traditional least squares and can be solved using OLS or its regularized variant (e.g., pseudo-inverse)

# Alternating Least Squares (ALS)

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

# Alternating Least Squares (ALS)

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

Let's assume we fix the item latent vector  $\mathbf{w}_i$  and we take the gradient with respect to the user latent vector  $\mathbf{x}_u$

# Alternating Least Squares (ALS)

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} \left( r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i \right)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

Let's assume we fix the item latent vector  $\mathbf{w}_i$  and we take the gradient with respect to the user latent vector  $\mathbf{x}_u$

$$\nabla L(\mathbf{x}_u; \mathbf{w}_i) = \frac{1}{2} \left[ -2 \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + 2\lambda \mathbf{x}_u \right] = - \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u$$

# Alternating Least Squares (ALS)

$$L(X, W) = \sum_{(u,i) \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i)^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 \right)$$

Let's assume we fix the item latent vector  $\mathbf{w}_i$  and we take the gradient with respect to the user latent vector  $\mathbf{x}_u$

$$\nabla L(\mathbf{x}_u; \mathbf{w}_i) = \frac{1}{2} \left[ -2 \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + 2\lambda \mathbf{x}_u \right] = - \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u$$

We want to set this to

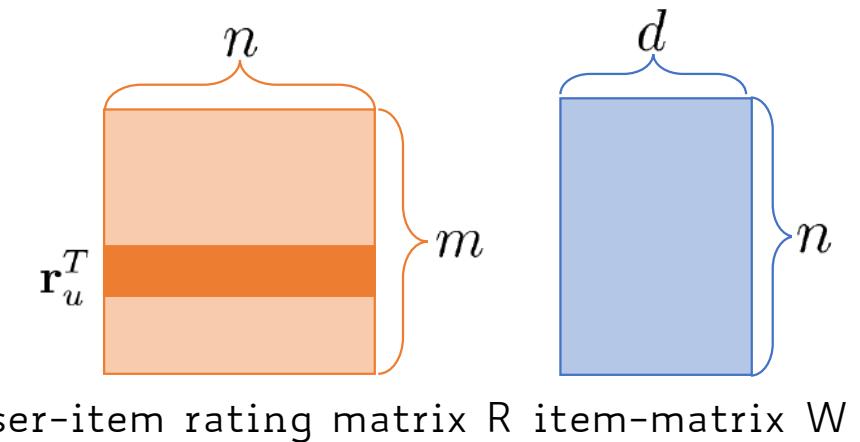
$$- \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

# ALS: Item Vector Fixed

$$-\sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

# ALS: Item Vector Fixed

$$\begin{aligned} & - \sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0 \\ &= -W^T(\mathbf{r}_u - W \cdot \mathbf{x}_u) + \lambda \mathbf{x}_u = 0 \end{aligned}$$

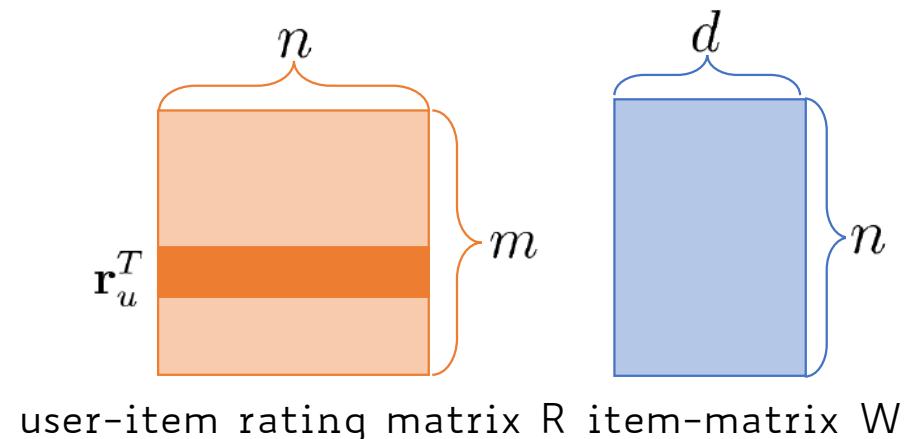


# ALS: Item Vector Fixed

$$-\sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

$$= -W^T(\mathbf{r}_u - W \cdot \mathbf{x}_u) + \lambda \mathbf{x}_u = 0$$

$$= W^T \cdot \mathbf{r}_u = W^T W \cdot \mathbf{x}_u + \lambda \mathbf{x}_u$$



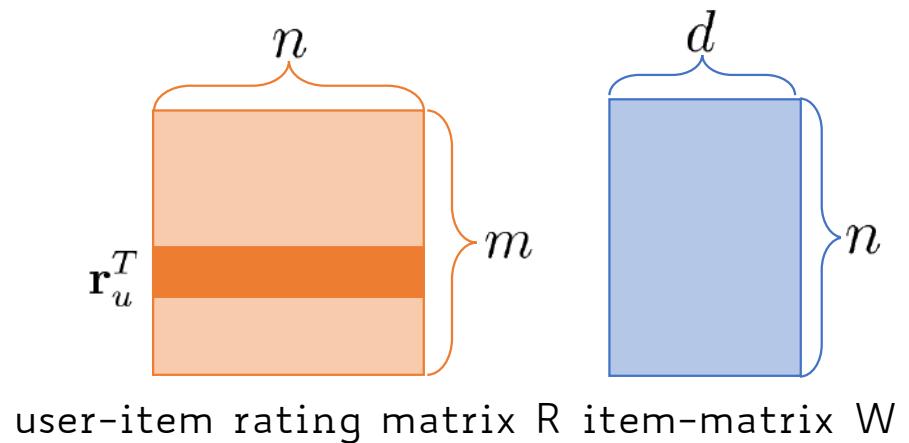
# ALS: Item Vector Fixed

$$-\sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

$$= -W^T(\mathbf{r}_u - W \cdot \mathbf{x}_u) + \lambda \mathbf{x}_u = 0$$

$$= W^T \cdot \mathbf{r}_u = W^T W \cdot \mathbf{x}_u + \lambda \mathbf{x}_u$$

$$= W^T \cdot \mathbf{r}_u = \mathbf{x}_u (W^T W + \lambda I) \quad I \in \mathbb{R}_{d \times d} \text{ identity matrix}$$



# ALS: Item Vector Fixed

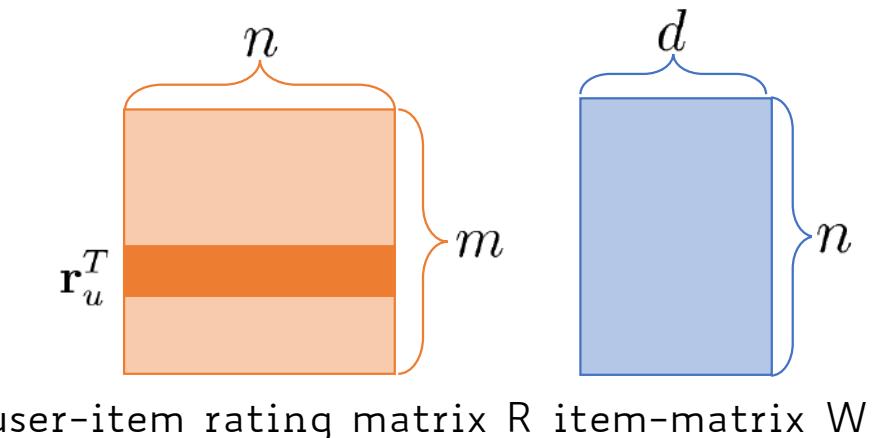
$$-\sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

$$= -W^T(\mathbf{r}_u - W \cdot \mathbf{x}_u) + \lambda \mathbf{x}_u = 0$$

$$= W^T \cdot \mathbf{r}_u = W^T W \cdot \mathbf{x}_u + \lambda \mathbf{x}_u$$

$$= W^T \cdot \mathbf{r}_u = \mathbf{x}_u (W^T W + \lambda I) \quad I \in \mathbb{R}_{d \times d} \text{ identity matrix}$$

$$= (W^T W + \lambda I)^{-1} \cdot W^T \cdot \mathbf{r}_u = \mathbf{x}_u (W^T W + \lambda I) \cdot (W^T W + \lambda I)^{-1}$$



# ALS: Item Vector Fixed

$$-\sum_{i \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{w}_i + \lambda \mathbf{x}_u = 0$$

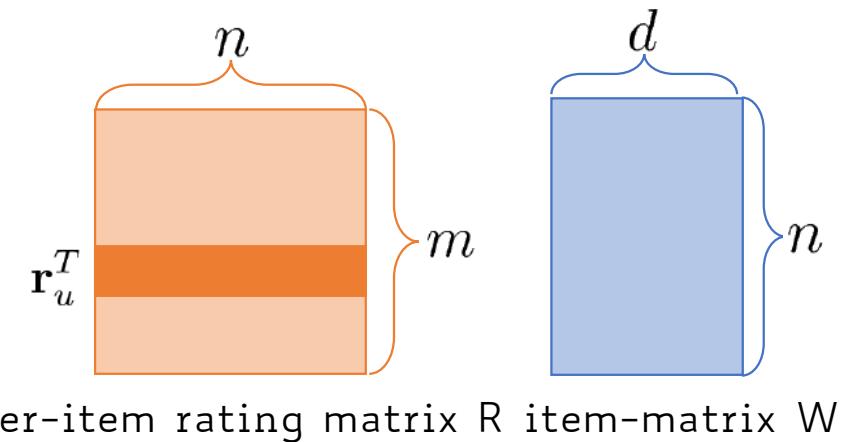
$$= -W^T(\mathbf{r}_u - W \cdot \mathbf{x}_u) + \lambda \mathbf{x}_u = 0$$

$$= W^T \cdot \mathbf{r}_u = W^T W \cdot \mathbf{x}_u + \lambda \mathbf{x}_u$$

$$= W^T \cdot \mathbf{r}_u = \mathbf{x}_u (W^T W + \lambda I) \quad I \in \mathbb{R}_{d \times d} \text{ identity matrix}$$

$$= (W^T W + \lambda I)^{-1} \cdot W^T \cdot \mathbf{r}_u = \mathbf{x}_u (W^T W + \lambda I) \cdot (W^T W + \lambda I)^{-1}$$

$$\boxed{\mathbf{x}_u = (W^T W + \lambda I)^{-1} \cdot W^T \cdot \mathbf{r}_u}$$



# ALS: User Vector Fixed

$$-\sum_{u \in \mathcal{D}} (r_{u,i} - \mathbf{x}_u^T \cdot \mathbf{w}_i) \mathbf{x}_u + \lambda \mathbf{w}_i = 0$$

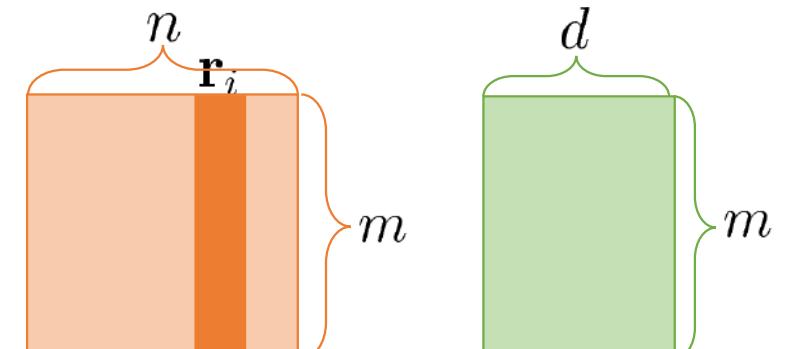
$$= -X^T(\mathbf{r}_i - X \cdot \mathbf{w}_i) - \lambda \mathbf{w}_i = 0$$

$$= X^T \cdot \mathbf{r}_i = X^T X \cdot \mathbf{w}_i + \lambda \mathbf{w}_i$$

$$= X^T \cdot \mathbf{r}_i = \mathbf{w}_i(X^T X + \lambda I) \quad I \in \mathbb{R}_{d \times d} \text{ identity matrix}$$

$$= (X^T X + \lambda I)^{-1} \cdot X^T \cdot \mathbf{r}_i = \mathbf{w}_i(X^T X + \lambda I) \cdot (X^T X + \lambda I)^{-1}$$

$$\boxed{\mathbf{w}_i = (X^T X + \lambda I)^{-1} \cdot X^T \cdot \mathbf{r}_i}$$



user-item rating matrix R user-matrix X

# ALS: Pseudocode

1. Initialize all the user latent vectors  $X$  and all the item latent vectors  $W$  randomly

# ALS: Pseudocode

1. Initialize all the user latent vectors  $X$  and all the item latent vectors  $W$  randomly
2. Fix all the item vectors  $W$  and solve for  $X$  (users)

# ALS: Pseudocode

1. Initialize all the user latent vectors  $X$  and all the item latent vectors  $W$  randomly
2. Fix all the item vectors  $W$  and solve for  $X$  (users)
3. Fix all the user vectors  $X$  and solve for  $W$  (items)

# ALS: Pseudocode

1. Initialize all the user latent vectors  $X$  and all the item latent vectors  $W$  randomly
2. Fix all the item vectors  $W$  and solve for  $X$  (users)
3. Fix all the user vectors  $X$  and solve for  $W$  (items)
4. Repeat step 2 and 3 until convergence

# ALS: Pseudocode

1. Initialize all the user latent vectors  $X$  and all the item latent vectors  $W$  randomly
2. Fix all the item vectors  $W$  and solve for  $X$  (users)
3. Fix all the user vectors  $X$  and solve for  $W$  (items)
4. Repeat step 2 and 3 until convergence

Convergence is guaranteed because in each step the loss function either decreases or stays unchanged, never increases

# ALS: Pseudocode

- **Key Properties:**

- Deterministic updates → each alternating step is a guaranteed minimizer for the constrained objective (closed-form)
- No learning rate → much easier to tune
- Parallelizable because each user (or item) update is independent
- Typically converges in fewer iterations (but each iteration is heavier computationally)
- Solves large linear systems → uses more memory and computational power

# ALS: Pseudocode

- Good for:

- Distributed systems (Spark, Hadoop) where parallel computation is essential
- Dense or moderately sparse datasets
- Scenarios where stable, monotonic convergence is important

# SGD vs. ALS

- In general, SGD is easier and faster than ALS

# SGD vs. ALS

- In general, SGD is easier and faster than ALS
- However, ALS is favorable in at least **2** cases:

# SGD vs. ALS

- In general, SGD is easier and faster than ALS
- However, ALS is favorable in at least **2** cases:
  - **Parallelization:** each  $x_u$  and  $w_i$  is computed independently of user/item factors

# SGD vs. ALS

- In general, SGD is easier and faster than ALS
- However, ALS is favorable in at least **2** cases:
  - **Parallelization:** each  $x_u$  and  $w_i$  is computed independently of user/item factors
  - **Implicit Data:** the training set is dense and looping over each single instance - as SGD does - would be unfeasible

# SGD vs. ALS

Aspect	SGD	ALS
Update style	Incremental (rating by rating)	Block-wise (solve full least-squares problems)
Convergence	Potentially slow, noisy	Faster in iterations, monotonic
Parallelism	Hard to parallelize safely	Highly parallelizable
Hyperparameters	Learning rate must be tuned	Only regularization ( $\lambda$ )
Implementation complexity	Very simple	More complex (matrix inversions)
Dataset size	Very large-scale sparse	Large-scale, especially distributed
Adaptability	Good for online/streaming	Batch-oriented
Stability	Sensitive to $\eta$	Very stable, no oscillations
Hardware	Works on a single CPU/GPU easily	Thrives on clusters / distributed systems

# Including Biases

- One benefit of the matrix factorization approach to CF is its **flexibility** in dealing with various data aspects

# Including Biases

- One benefit of the matrix factorization approach to CF is its **flexibility** in dealing with various data aspects
- The basic learning framework tries to capture the interactions between users and items that produce the different rating values

# Including Biases

- However, much of the observed variation in ratings depends on **biases** associated with users or items, independent of any interactions

# Including Biases

- However, much of the observed variation in ratings depends on **biases** associated with users or items, independent of any interactions
- For example, some users systematically tend to give higher ratings than others, and some items receive higher ratings than others

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate
- We separate the latent factor modeling from the **bias modeling**

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate
- We separate the latent factor modeling from the **bias modeling**
- Given a rating  $r_{u,i}$  a first-order approximation of the bias involved with it ( $b_{u,i}$ ) can be defined as follows:

$$b_{u,i} = \mu + b_u + b_i$$

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate
- We separate the latent factor modeling from the **bias modeling**
- Given a rating  $r_{u,i}$  a first-order approximation of the bias involved with it ( $b_{u,i}$ ) can be defined as follows:

$$b_{u,i} = \mu + b_u + b_i$$

Overall avg. rating

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate
- We separate the latent factor modeling from the **bias modeling**
- Given a rating  $r_{u,i}$  a first-order approximation of the bias involved with it ( $b_{u,i}$ ) can be defined as follows:

$$b_{u,i} = \mu + b_u + b_i$$

Observed deviations of user u from the avg.

# Including Biases

- Trying to explain observed ratings from user-item interactions only may not be enough accurate
- We separate the latent factor modeling from the **bias modeling**
- Given a rating  $r_{u,i}$  a first-order approximation of the bias involved with it ( $b_{u,i}$ ) can be defined as follows:

$$b_{u,i} = \mu + b_u + b_i$$

# Including Biases: Example

$$b_{u,i} = \mu + b_u + b_i$$

We want a first-order estimate for user Joe's rating of the movie Titanic

# Including Biases: Example

$$b_{u,i} = \mu + b_u + b_i$$

We want a first-order estimate for user Joe's rating of the movie Titanic

$$\mu = 3.7$$

The average rating over all movies (i.e., items)

# Including Biases: Example

$$b_{u,i} = \mu + b_u + b_i$$

We want a first-order estimate for user Joe's rating of the movie Titanic

$\mu = 3.7$  The average rating over all movies (i.e., items)

$b_{\text{Titanic}} = 0.5$  Titanic is a 0.5 stars above the avg. rated movie

# Including Biases: Example

$$b_{u,i} = \mu + b_u + b_i$$

We want a first-order estimate for user Joe's rating of the movie Titanic

$\mu = 3.7$  The average rating over all movies (i.e., items)

$b_{\text{Titanic}} = 0.5$  Titanic is a 0.5 stars above the avg. rated movie

$b_{\text{Joe}} = -0.3$  Joe is a critical user who tends to give 0.3 less stars than avg.

# Including Biases: Example

$$b_{u,i} = \mu + b_u + b_i$$

We want a first-order estimate for user Joe's rating of the movie Titanic

$\mu = 3.7$  The average rating over all movies (i.e., items)

$b_{\text{Titanic}} = 0.5$  Titanic is a 0.5 stars above the avg. rated movie

$b_{\text{Joe}} = -0.3$  Joe is a critical user who tends to give 0.3 less stars than avg.

$$b_{\text{Joe,Titanic}} = 3.7 - 0.3 + 0.5 = 3.9$$

Bias term

# Including Bias into the Optimization

$$\hat{r}_{u,i} = \underbrace{\mathbf{x}_u^T \cdot \mathbf{w}_i}_{\text{latent factors}} + \underbrace{\mu + b_u + b_i}_{\text{bias}}$$

The estimated rating of an item  $i$  for the user  $u$  is now made of **2 components**

# Including Bias into the Optimization

$$\hat{r}_{u,i} = \underbrace{\mathbf{x}_u^T \cdot \mathbf{w}_i}_{\text{latent factors}} + \underbrace{\mu + b_u + b_i}_{\text{bias}}$$

The estimated rating of an item  $i$  for the user  $u$  is now made of 2 components

Latent factor term

models user-item interaction

# Including Bias into the Optimization

$$\hat{r}_{u,i} = \underbrace{\mathbf{x}_u^T \cdot \mathbf{w}_i}_{\text{latent factors}} + \underbrace{\mu + b_u + b_i}_{\text{bias}}$$

The estimated rating of an item  $i$  for the user  $u$  is now made of 2 components

Latent factor term

models user-item interaction

Bias term

models global average,  
user and item bias

# Including Bias into the Optimization

Overall, the original optimization problem becomes as follows

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left[ r_{u,i} - (\mathbf{x}_u^T \cdot \mathbf{w}_i + \mu + b_u + b_i) \right]^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 + \sum_{u \in \mathcal{D}} b_u^2 + \sum_{i \in \mathcal{D}} b_i^2 \right) \right\}$$

# Including Bias into the Optimization

Overall, the original optimization problem becomes as follows

$$X^*, W^* = \operatorname{argmin}_{X,W} \left\{ \frac{1}{2} \sum_{(u,i) \in \mathcal{D}} \left[ r_{u,i} - (\mathbf{x}_u^T \cdot \mathbf{w}_i + \mu + b_u + b_i) \right]^2 + \lambda \left( \sum_{u \in \mathcal{D}} \|\mathbf{x}_u\|^2 + \sum_{i \in \mathcal{D}} \|\mathbf{w}_i\|^2 + \sum_{u \in \mathcal{D}} b_u^2 + \sum_{i \in \mathcal{D}} b_i^2 \right) \right\}$$

Can still be solved using ALS

# Collaborative Filtering: Drawbacks

CF methods suffer from 3 main problems

# Collaborative Filtering: Drawbacks

CF methods suffer from 3 main problems

## cold start

for a new user/item  
entering the system  
there is not enough  
data to make  
recommendations

# Collaborative Filtering: Drawbacks

CF methods suffer from 3 main problems

## cold start

for a new user/item entering the system there is not enough data to make recommendations

## scalability

million users/items systems may require extraordinary computational power to generate recommendations

# Collaborative Filtering: Drawbacks

CF methods suffer from 3 main problems

## cold start

for a new user/item entering the system there is not enough data to make recommendations

## scalability

million users/items systems may require extraordinary computational power to generate recommendations

## sparsity

the vast majority of items are not rated by users

# Hybrid: Content-based + Collaborative Filtering

- Most recommender systems now use a **hybrid** approach, combining **content-based** and **collaborative filtering**

# Hybrid: Content-based + Collaborative Filtering

- Most recommender systems now use a **hybrid** approach, combining **content-based** and **collaborative filtering**
- Hybrid approaches can be implemented in several ways:

# Hybrid: Content-based + Collaborative Filtering

- Most recommender systems now use a **hybrid** approach, combining **content-based** and **collaborative filtering**
- Hybrid approaches can be implemented in several ways:
  - Combining individual recommendations from content-based and collaborative filtering systems, separately

# Hybrid: Content-based + Collaborative Filtering

- Most recommender systems now use a **hybrid** approach, combining **content-based** and **collaborative filtering**
- Hybrid approaches can be implemented in several ways:
  - Combining individual recommendations from content-based and collaborative filtering systems, separately
  - Adding content-based capabilities to a collaborative-based approach (and vice versa)

# Hybrid: Content-based + Collaborative Filtering

- Most recommender systems now use a **hybrid** approach, combining **content-based** and **collaborative filtering**
- Hybrid approaches can be implemented in several ways:
  - Combining individual recommendations from content-based and collaborative filtering systems, separately
  - Adding content-based capabilities to a collaborative-based approach (and vice versa)
  - Unifying the two approaches into one model

# Hybrid: Content-based + Collaborative Filtering

- Hybrid methods have shown to provide **more accurate** recommendations than pure content-based or collaborative filtering

# Hybrid: Content-based + Collaborative Filtering

- Hybrid methods have shown to provide **more accurate** recommendations than pure content-based or collaborative filtering
- They can also be used to overcome common problems in recommender systems such as cold start and the sparseness of user-item matrix

# Hybrid: Content-based + Collaborative Filtering

- Hybrid methods have shown to provide **more accurate** recommendations than pure content-based or collaborative filtering
- They can also be used to overcome common problems in recommender systems such as cold start and the sparseness of user-item matrix
- **Netflix** is a good example of hybrid recommender systems

# Netflix's Hybrid Recommender System

Recommendations are generated

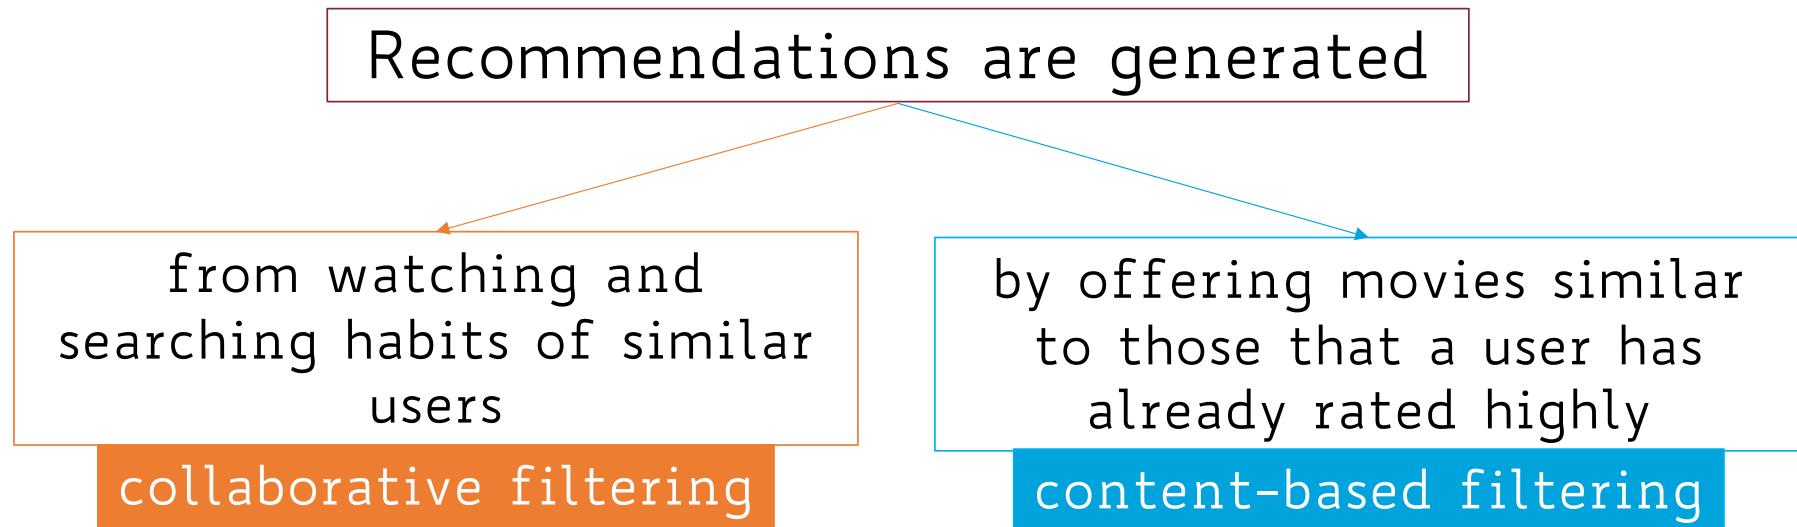
# Netflix's Hybrid Recommender System

Recommendations are generated

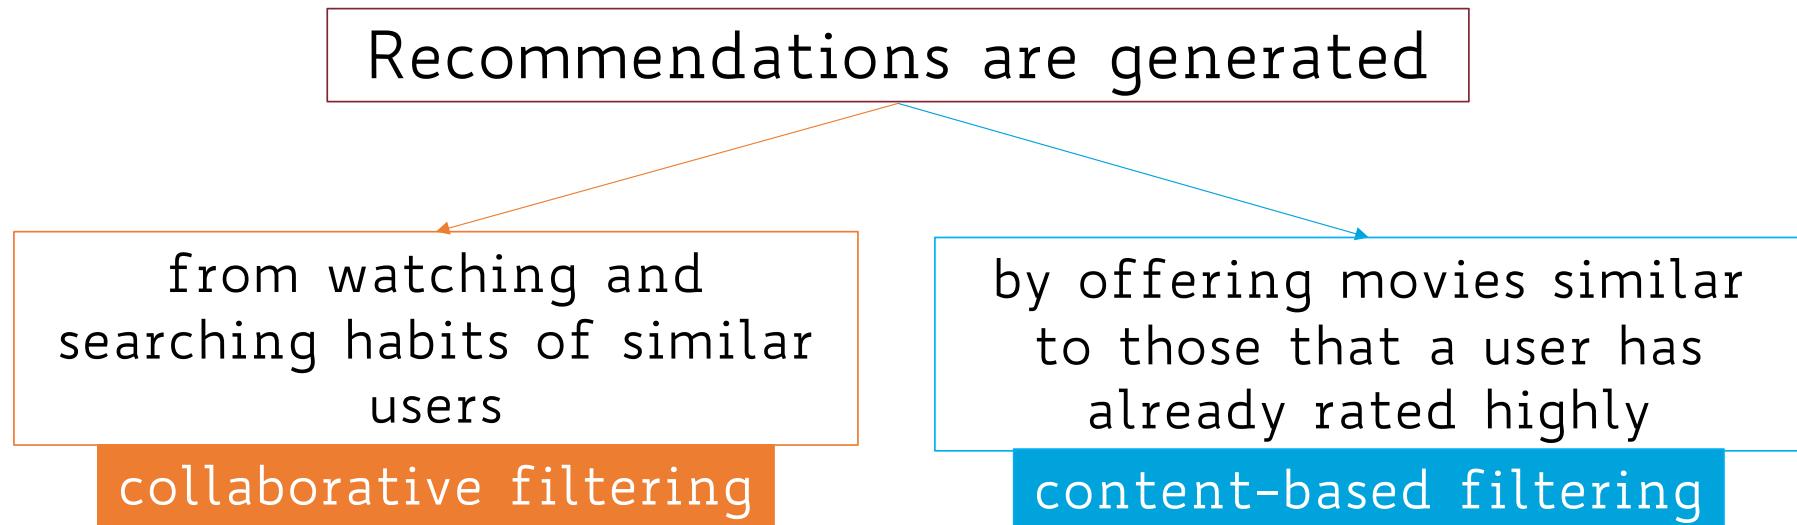
from watching and  
searching habits of similar  
users

collaborative filtering

# Netflix's Hybrid Recommender System



# Netflix's Hybrid Recommender System



## [Netflix: What Happens When You Press Play?](#)

For more details about how Netflix actually works

# The Netflix Prize

- In 2006, the online DVD rental company Netflix announced a [contest](#) to improve the state of its recommender system, called Cinematch

# The Netflix Prize

- In 2006, the online DVD rental company Netflix announced a [contest](#) to improve the state of its recommender system, called Cinematch
- Released a training set of ~100M ratings from about 500K anonymous customers and their ratings on more than 17K movies (1 to 5 stars)

# The Netflix Prize

- In 2006, the online DVD rental company Netflix announced a [contest](#) to improve the state of its recommender system, called Cinematch
- Released a training set of ~100M ratings from about 500K anonymous customers and their ratings on more than 17K movies (1 to 5 stars)
- Participating teams submit predicted ratings for a test set of approximately 3M ratings

# The Netflix Prize

- Netflix calculates a Root Mean Squared Error (RMSE) based on the held-out truth

# The Netflix Prize

- Netflix calculates a Root Mean Squared Error (RMSE) based on the held-out truth
- The first team that can improve on the Netflix algorithm's RMSE performance by 10% or more wins a \$1 million prize

# The Netflix Prize

- If no team reaches the 10 percent goal, Netflix gives a \$50,000 Progress Prize to the team in first place at the end of each year

# The Netflix Prize

- If no team reaches the 10 percent goal, Netflix gives a \$50,000 Progress Prize to the team in first place at the end of each year
- According to the [contest website](#), more than 48,000 teams from 182 different countries have downloaded the data

# The Netflix Prize: The Winners

- The BellKor team (AT&T Labs) took over the top spot in the competition in the summer of 2007

# The Netflix Prize: The Winners

- The BellKor team (AT&T Labs) took over the top spot in the competition in the summer of 2007
- It firstly won the 2007 Progress Prize with the best score at the time: 8.43% better than Netflix

# The Netflix Prize: The Winners

- The BellKor team (AT&T Labs) took over the top spot in the competition in the summer of 2007
- It firstly won the 2007 Progress Prize with the best score at the time: 8.43% better than Netflix
- On June 26, 2009 the team "BellKor's Pragmatic Chaos", a merger of "Bellkor in BigChaos" and "Pragmatic Theory", achieved a 10.05% lift

# The Netflix Prize: The Winners

- The BellKor team (AT&T Labs) took over the top spot in the competition in the summer of 2007
- It firstly won the 2007 Progress Prize with the best score at the time: 8.43% better than Netflix
- On June 26, 2009 the team "BellKor's Pragmatic Chaos", a merger of "Bellkor in BigChaos" and "Pragmatic Theory", achieved a 10.05% lift

A combination of 100 different predictor sets, mostly factorization models

# Evaluation Metrics

How do we evaluate recommendations generated?

# Evaluation Metrics

How do we evaluate recommendations generated?

## Offline

RMSE, MAE,  
MAP@K, MAR@K,  
Coverage,  
Personalization

# Evaluation Metrics

How do we evaluate recommendations generated?

## Offline

RMSE, MAE,  
MAP@K, MAR@K,  
Coverage,  
Personalization

## Online

A/B testing measuring  
CTR, ROI, and other  
"live" metrics

# Evaluation Metrics: RMSE

$$\text{RMSE} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sqrt{\sum_{(u,i) \in \mathcal{D}_{\text{test}}} (r_{u,i} - \hat{r}_{u,i})^2}$$

# Evaluation Metrics: RMSE

$$\text{RMSE} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sqrt{\sum_{(u,i) \in \mathcal{D}_{\text{test}}} (r_{u,i} - \hat{r}_{u,i})^2}$$

Narrow focus on accuracy may penalize prediction diversity  
(all recommendations are too similar to the item)

# Evaluation Metrics: RMSE

$$\text{RMSE} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sqrt{\sum_{(u,i) \in \mathcal{D}_{\text{test}}} (r_{u,i} - \hat{r}_{u,i})^2}$$

Narrow focus on accuracy may penalize prediction diversity  
(all recommendations are too similar to the item)

The order of recommendations should also be taken into account

# Evaluation Metrics: RMSE

$$\text{RMSE} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sqrt{\sum_{(u,i) \in \mathcal{D}_{\text{test}}} (r_{u,i} - \hat{r}_{u,i})^2}$$

Narrow focus on accuracy may penalize prediction diversity  
(all recommendations are too similar to the item)

The order of recommendations should also be taken into account

The RMSE might penalize a method that does well for high ratings and badly for others

# Evaluation Metrics: Precision & Recall

For a binary classifier predicting a condition ( $y = 1$ ) or not, we define

$$P = \frac{TP}{TP + FP} \quad R = \frac{TP}{TP + FN}$$

# Evaluation Metrics: Precision & Recall

For a binary classifier predicting a condition ( $y = 1$ ) or not, we define

$$P = \frac{TP}{TP + FP} \quad R = \frac{TP}{TP + FN}$$

Mapping of binary classification terminology to recommender systems

binary classifier	recommender system
# with condition ( $y = 1$ )	# of all possible relevant items for a user
# predicted positive ( $TP + FP$ )	# of recommended items
# correct positives ( $TP$ )	# of recommended items that are relevant

# Evaluation Metrics: Precision & Recall

For a recommender system, we can therefore define

$$P = \frac{\# \text{ relevant item recommendations}}{\# \text{ items recommended}}$$

$$R = \frac{\# \text{ relevant item recommendations}}{\# \text{ items actually relevant}}$$

# Evaluation Metrics: Precision & Recall

For a recommender system, we can therefore define

$$P = \frac{\# \text{ relevant item recommendations}}{\# \text{ items recommended}} \quad R = \frac{\# \text{ relevant item recommendations}}{\# \text{ items actually relevant}}$$

A recommender system generates k=5 items to recommend

There are only 3 relevant items

The success/failure of our recommendations: [0, 1, 1, 0, 0] 0=not relevant/1=relevant

$$P = \frac{2}{5} \quad R = \frac{2}{3}$$

# Evaluation Metrics: Precision & Recall @ k

- Precision and Recall don't seem to care about ordering

# Evaluation Metrics: Precision & Recall @ k

- Precision and Recall don't seem to care about ordering
- Consider Precision and Recall at cutoff k (i.e., P@k and R@k)

# Evaluation Metrics: Precision & Recall @ k

- Precision and Recall don't seem to care about ordering
- Consider Precision and Recall at cutoff k (i.e., P@k and R@k)
- Imagine taking our list of N recommendations and considering only the first element, then only the first two, then only the first three, and so on

# Evaluation Metrics: Precision & Recall @ k

- Precision and Recall don't seem to care about ordering
- Consider Precision and Recall at cutoff k (i.e., P@k and R@k)
- Imagine taking our list of N recommendations and considering only the first element, then only the first two, then only the first three, and so on
- P@k and R@k are simply the precision and recall calculated only from the subset of the first k recommendations

# P@k: Example

$$k = 3 \\ P@3 = \frac{1}{3}$$

Rank	Product Recommended	Result
1	Credit card	Correct positive
2	Christmas Fund	False positive
3	Debit Card	False positive
4	Auto loan	False positive
5	HELOC	Correct Positive
6	College Fund	Correct positive
7	Personal loan	False positive

# P@k: Example

$$k = 3 \\ P@3 = \frac{1}{3}$$

Rank	Product Recommended	Result
1	Credit card	Correct positive
2	Christmas Fund	False positive
3	Debit Card	False positive
4	Auto loan	False positive
5	HELOC	Correct Positive
6	College Fund	Correct positive
7	Personal loan	False positive

Rank	Product Recommended	Result
1	Credit card	Correct positive
2	Christmas Fund	False positive
3	Debit Card	False positive
4	Auto loan	False positive
5	HELOC	Correct Positive
6	College Fund	Correct positive
7	Personal loan	False positive

$$k = 6 \\ P@6 = \frac{3}{6}$$

# Average Precision (AP)

Suppose our recommender system must return  $N$  items,  
with  $|Rel|$  actually relevant items

# Average Precision (AP)

Suppose our recommender system must return  $N$  items,  
with  $|\text{Rel}|$  actually relevant items

We define the Average Precision (AP) as follows:

$$AP@N = \frac{1}{|\text{Rel}|} \sum_{k=1}^N P@k \times \mathbf{1}_{\text{Rel}}(k)$$

# Average Precision (AP)

Suppose our recommender system must return  $N$  items,  
with  $|\text{Rel}|$  actually relevant items

We define the Average Precision (AP) as follows:

$$AP@N = \frac{1}{|\text{Rel}|} \sum_{k=1}^N P@k \times \mathbf{1}_{\text{Rel}}(k)$$

indicator function     $\mathbf{1}_{\text{Rel}}(k) = \begin{cases} 1 & \text{if item } k \in \text{Rel} \\ 0 & \text{otherwise} \end{cases}$

# Mean Average Precision (MAP)

AP@N is computed for a single data point (i.e., user)

# Mean Average Precision (MAP)

AP@N is computed for a single data point (i.e., user)

We define the Mean Average Precision (MAP) as follows:

$$MAP@N = \frac{1}{|\mathcal{U}|} \sum_{u=1}^{|\mathcal{U}|} AP@N(u) = \frac{1}{|\mathcal{U}|} \sum_{u=1}^{|\mathcal{U}|} \frac{1}{|\text{Rel}|} \sum_{k=1}^N P@k(u) \times \mathbf{1}_{\text{Rel}}(k, u)$$

# Personalization

We may need a way to assess if a model recommends many of the same items to different users

# Personalization

We may need a way to assess if a model recommends many of the same items to different users

It is defined as the dissimilarity (i.e., 1-cosine similarity) between user's lists of recommendations

# Personalization

We may need a way to assess if a model recommends many of the same items to different users

It is defined as the dissimilarity (i.e., 1-cosine similarity) between user's lists of recommendations

Intuitively, a high personalization score indicates the recommender system is able to provide a **highly personalized** experience to the users

# Personalization

Suppose 3 users are recommended the following lists  
of items

$$u_1 = [A, B, C, D] \quad u_2 = [A, B, C, E] \quad u_3 = [A, B, F, G]$$

# Personalization

Suppose 3 users are recommended the following lists of items

$$u_1 = [A, B, C, D] \quad u_2 = [A, B, C, E] \quad u_3 = [A, B, F, G]$$

	A	B	C	D	E	F	G
$u_1$	1	1	1	1	0	0	0
$u_2$	1	1	1	0	1	0	0
$u_3$	1	1	0	0	0	1	1

# Personalization

Compute the 3-by-3 triangular matrix containing the cosine similarity between each pair of user's recommendation binary vector

$$M_{i,j} = \text{cosine}(\mathbf{u}_i, \mathbf{u}_j)$$

	$\mathbf{u}_1$	$\mathbf{u}_2$	$\mathbf{u}_3$
$\mathbf{u}_1$	1	0.75	0.58
$\mathbf{u}_2$	0.75	1	0.58
$\mathbf{u}_3$	0.58	0.58	1

# Personalization

Compute the 3-by-3 triangular matrix containing the cosine similarity between each pair of user's recommendation binary vector

$$M_{i,j} = \text{cosine}(\mathbf{u}_i, \mathbf{u}_j)$$

	$\mathbf{u}_1$	$\mathbf{u}_2$	$\mathbf{u}_3$
$\mathbf{u}_1$	1	0.75	0.58
$\mathbf{u}_2$	0.75	1	0.58
$\mathbf{u}_3$	0.58	0.58	1

$\sim 0.64$

Take the average of the upper triangle of the matrix M above

# Personalization

Compute the 3-by-3 triangular matrix containing the cosine similarity between each pair of user's recommendation binary vector

$$M_{i,j} = \text{cosine}(\mathbf{u}_i, \mathbf{u}_j)$$

	$\mathbf{u}_1$	$\mathbf{u}_2$	$\mathbf{u}_3$
$\mathbf{u}_1$	1	0.75	0.58
$\mathbf{u}_2$	0.75	1	0.58
$\mathbf{u}_3$	0.58	0.58	1

~0.64

$$\text{Personalization} = 1 - 0.64 = 0.36$$

# Take-Home Message of Today

- 2 main approaches:
  - Content-based (explicitly creating user and item profiles)
  - Collaborative-filtering (extract patterns from past observed ratings)

# Take-Home Message of Today

- 2 main approaches:
  - Content-based (explicitly creating user and item profiles)
  - Collaborative-filtering (extract patterns from past observed ratings)
- Hybrid approaches combining both usually work better in practice

# Take-Home Message of Today

- 2 main approaches:
  - Content-based (explicitly creating user and item profiles)
  - Collaborative-filtering (extract patterns from past observed ratings)
- Hybrid approaches combining both usually work better in practice
- New Neural-Network-based approaches have been proposed recently

# Recommended Readings and Information :)

- A huge body of work on recommender systems is available out there!
- Surveys:
  - [Adomavicius & Tuzhilin](#) [2005]
  - [Koren & Volinsky](#) [2009]
  - [Bobadilla \*et al.\*](#) [2013]
  - [Zhang \*et al.\*](#) [2019]
- Well-renowned series of Conferences: [RecSys](#), [KDD](#), [SIGIR](#), [TheWebConf](#)