

Big Data Computing

Master's Degree in Computer Science

2025-2026



SAPIENZA
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science
Sapienza Università di Roma
tolomei@di.uniroma1.it

Similarity: Why Bother?

- We have seen a number of similarity/distance metrics

Similarity: Why Bother?

- We have seen a number of similarity/distance metrics
- Each metric may be suitable for specific task(s) in a particular domain

Similarity: Why Bother?

- We have seen a number of similarity/distance metrics
- Each metric may be suitable for specific task(s) in a particular domain
- **Clustering** is one of these tasks!

CLUSTERING

What is Clustering?

- A procedure to group a set of objects into classes of **similar** objects

What is Clustering?

- A procedure to group a set of objects into classes of **similar** objects
- A standard problem in many (big) data applications:

What is Clustering?

- A procedure to group a set of objects into classes of **similar** objects
- A standard problem in many (big) data applications:
 - Categorizing documents by their topics
 - Grouping customers by their behaviors
 - ...

What is Clustering?

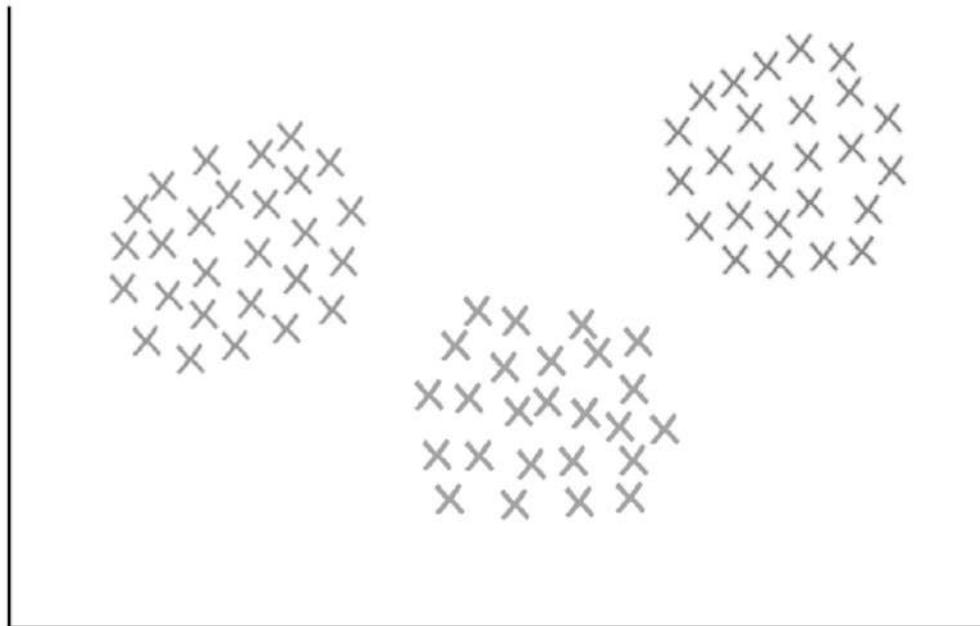
- A typical example of **unsupervised learning** technique

What is Clustering?

- A typical example of **unsupervised learning** technique
- A method of **data exploration**, i.e., a way of looking for patterns of interest in data

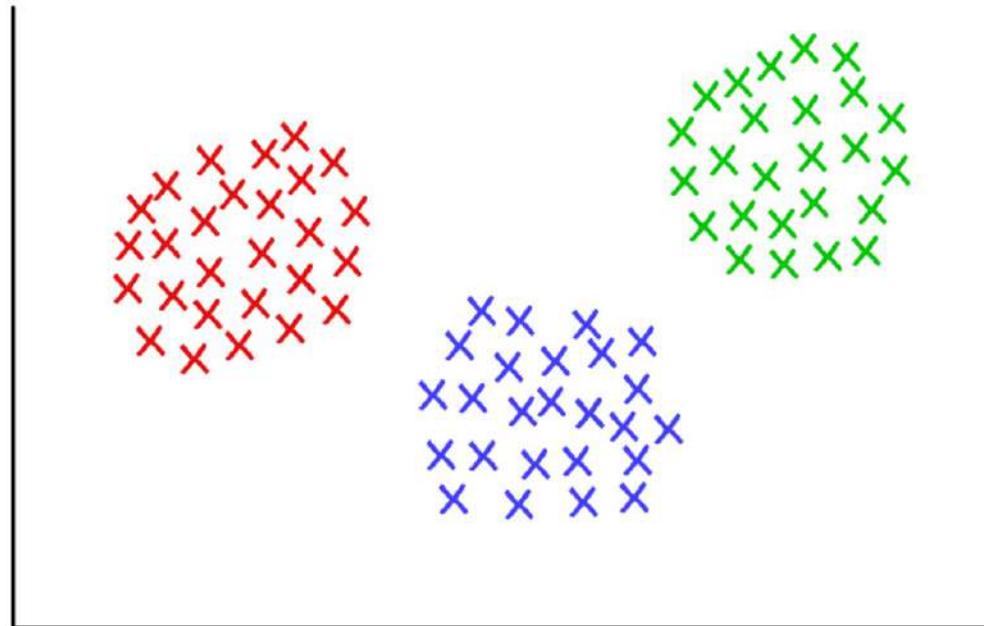
Clustering: Intuition

Given a set of 2-dimensional data points



Clustering: Intuition

We'd like to understand their "structure" to find groups of data points



Clustering: Formal Definition

- Given a set of data points and a notion of **distance** between those

Clustering: Formal Definition

- Given a set of data points and a notion of **distance** between those
- Group the data points into some number of clusters so that:

Clustering: Formal Definition

- Given a set of data points and a notion of **distance** between those
- Group the data points into some number of clusters so that:
 - Members of a cluster are close/similar to each other (i.e., **high intra-cluster similarity**)

Clustering: Formal Definition

- Given a set of data points and a notion of **distance** between those
- Group the data points into some number of clusters so that:
 - Members of a cluster are close/similar to each other (i.e., **high intra-cluster similarity**)
 - Members of different clusters are dissimilar (i.e., **low inter-cluster similarity**)

Clustering: Practical Issues

- Object **representation**
 - Data points may be in very high-dimensional spaces

Clustering: Practical Issues

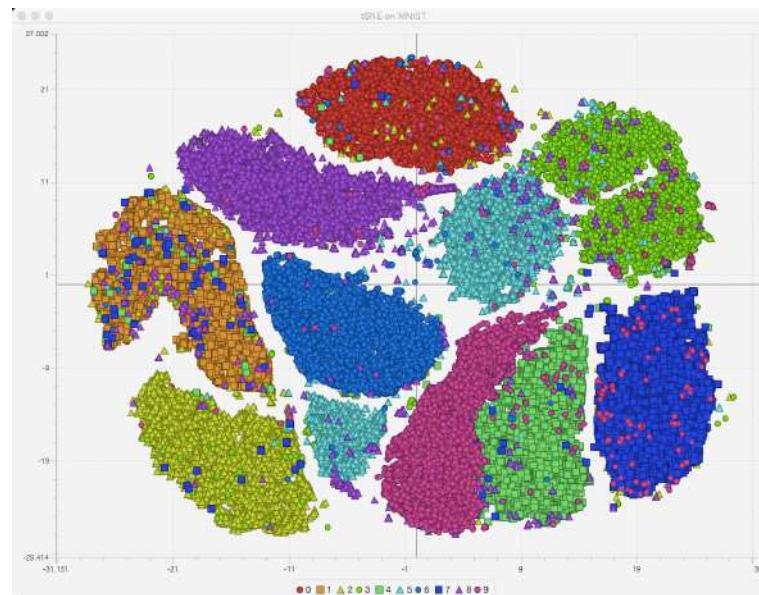
- Object **representation**
 - Data points may be in very high-dimensional spaces
- Notion of **similarity** between objects using a distance measure
 - Euclidean distance, Cosine similarity, Jaccard coefficient, etc.

Clustering: Practical Issues

- Object **representation**
 - Data points may be in very high-dimensional spaces
- Notion of **similarity** between objects using a distance measure
 - Euclidean distance, Cosine similarity, Jaccard coefficient, etc.
- Number of **output clusters**
 - Fixed apriori? Data-driven?

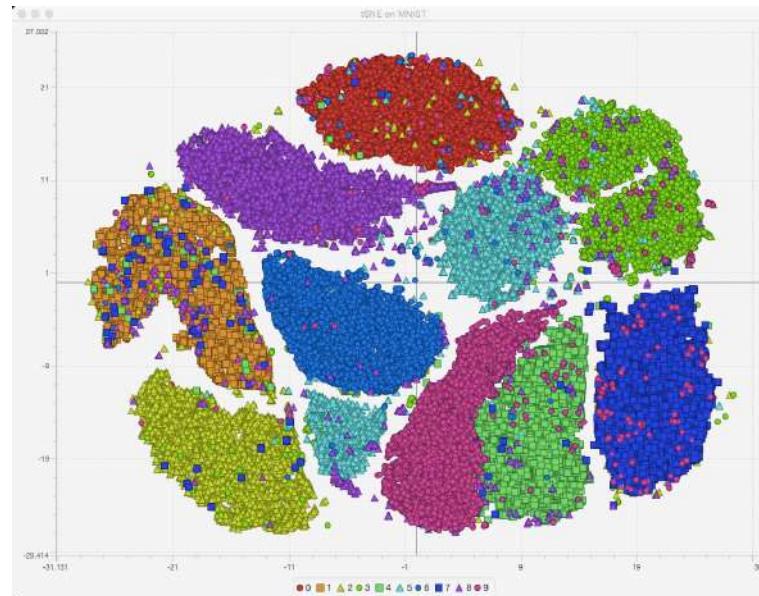
Clustering: A Hard Problem

Data points are not always easily and clearly **separable**



Clustering: A Hard Problem

Data points are not always easily and clearly **separable**



Finding a clear boundary between clusters may be **hard** in the real world

Clustering: A Hard Problem

- Clustering in 2 dimensions looks easy

Clustering: A Hard Problem

- Clustering in 2 dimensions looks easy
- So does clustering of a small number of data points

Clustering: A Hard Problem

- Clustering in 2 dimensions looks easy
- So does clustering of a small number of data points
- **What does make things hard, then?**

Clustering: A Hard Problem

- Clustering in 2 dimensions looks easy
- So does clustering of a small number of data points
- **What does make things hard, then?**
 - Many real-world applications involve 10s, 100s, or 1,000s of dimensions

Clustering: A Hard Problem

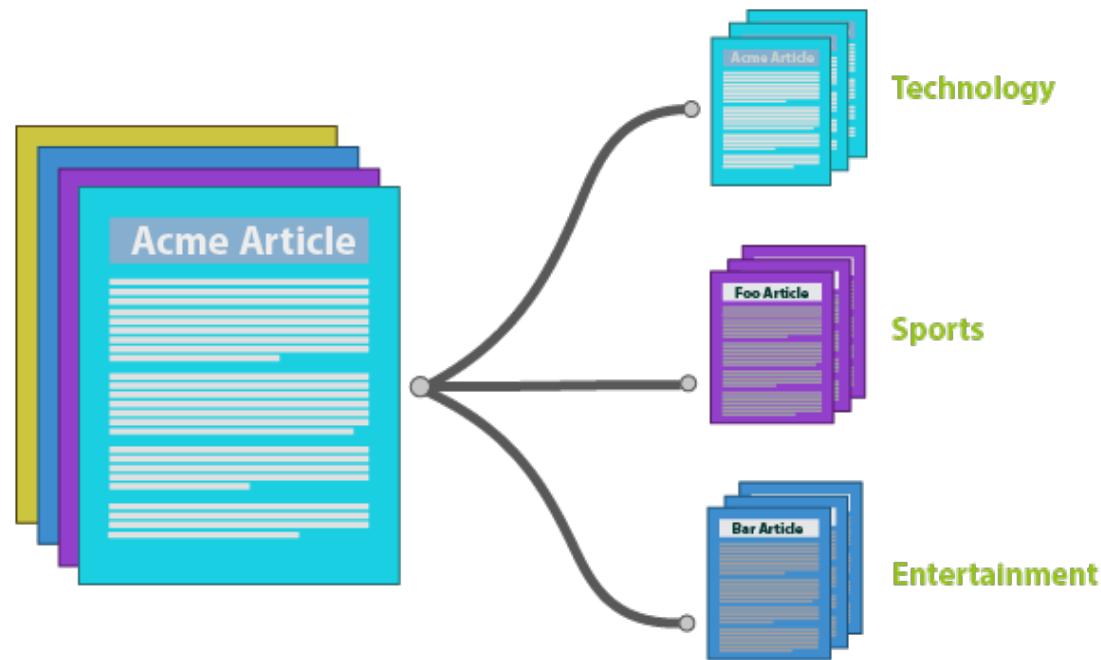
- Clustering in 2 dimensions looks easy
- So does clustering of a small number of data points
- **What does make things hard, then?**

Many real-world applications involve 10s, 100s, or 1,000s of dimensions



In high-dimensional spaces almost all pairs of points are at the same (large) distance

Example: Text Document Clustering



source: <https://towardsdatascience.com/applying-machine-learning-to-classify-an-unsupervised-text-document-e7bb6265f52>

Example: Text Document Clustering

- **Problem:** Group together documents on the same **topic**

Example: Text Document Clustering

- **Problem:** Group together documents on the same **topic**
- Documents with similar sets of words may be about the same topic

Example: Text Document Clustering

- **Problem:** Group together documents on the same **topic**
- Documents with similar sets of words may be about the same topic
- **Key Issues:**
 - Representing documents (in the space of words)

Example: Text Document Clustering

- **Problem:** Group together documents on the same **topic**
- Documents with similar sets of words may be about the same topic
- **Key Issues:**
 - Representing documents (in the space of words)
 - Measuring document similarity (in the space of words)

Document Representation

- Different ways of representing documents (in the space of words)

Document Representation

- Different ways of representing documents (in the space of words)
 - As a **set of words** (disregarding the order and multiplicity)

Document Representation

- Different ways of representing documents (in the space of words)
 - As a **set of words** (disregarding the order and multiplicity)
 - As a **bag-of-words** (i.e., a multiset disregarding the order yet keeping multiplicity)

Document Representation

- Different ways of representing documents (in the space of words)
 - As a **set of words** (disregarding the order and multiplicity)
 - As a **bag-of-words** (i.e., a multiset disregarding the order yet keeping multiplicity)
 - As a **bag-of- n -grams** (i.e., the more general case of bag-of-words)

Document Representation

- Different ways of representing documents (in the space of words)
 - As a **set of words** (disregarding the order and multiplicity)
 - As a **bag-of-words** (i.e., a multiset disregarding the order yet keeping multiplicity)
 - As a **bag-of- n -grams** (i.e., the more general case of bag-of-words)
 - More advanced representations derived from (Large) Neural Language Models (e.g., word2vec, BERT, Transformers)

Document Representation

- Different ways of representing documents (in the space of words)
 - As a **set of words** (disregarding the order and multiplicity)
 - As a **bag-of-words** (i.e., a multiset disregarding the order yet keeping multiplicity)
 - As a **bag-of- n -grams** (i.e., the more general case of bag-of-words)
 - More advanced representations derived from (Large) Neural Language Models (e.g., word2vec, BERT, Transformers)
- The choice of document representation affects the similarity measure

Document Representation: Set of Words

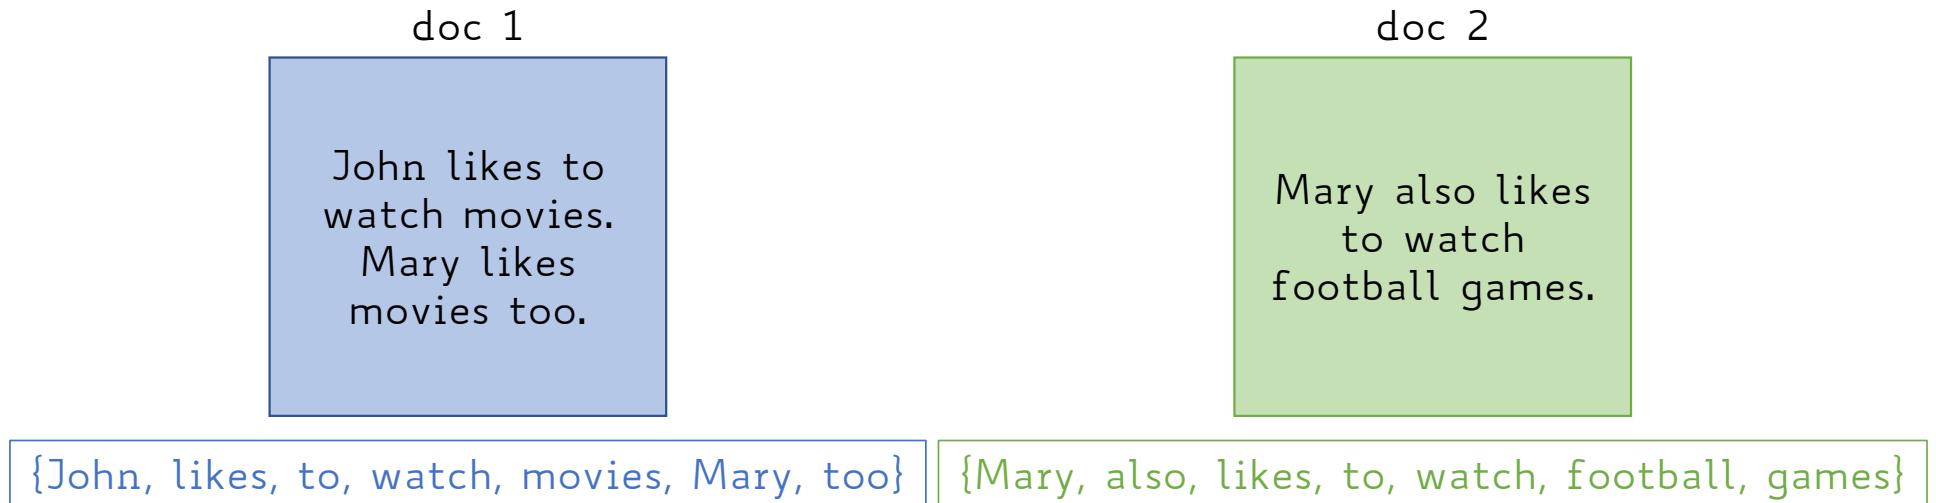
doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

Document Representation: Set of Words



Document Representation: Bag-of-Words

We keep **multiplicity**

doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

Document Representation: Bag-of-Words

We keep **multiplicity**

doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

```
{  
  John:1, likes:2, to:1,  
  watch:1,  
  movies:2, Mary:1, too:1  
}
```

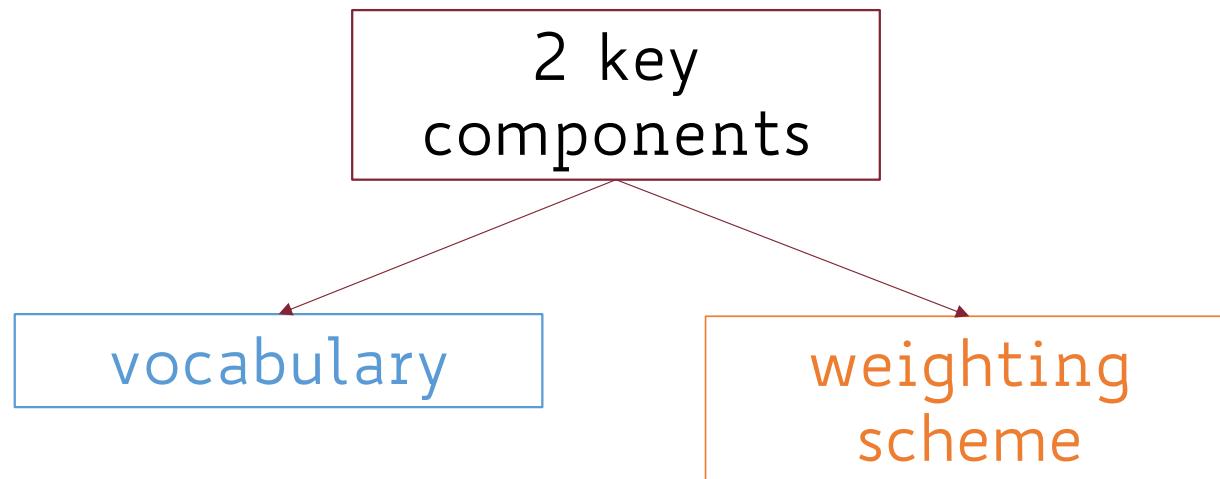
```
{  
  Mary:1, also:1, likes:1, to:1,  
  watch:1, football:1, games:1  
}
```

Document Representation: Bag-of-Words

Bag-of-Words (BoW) model is just a preliminary step for more complex document representations

Document Representation: Bag-of-Words

Bag-of-Words (BoW) model is just a preliminary step for more complex document representations



Bag-of-Words:Vocabulary

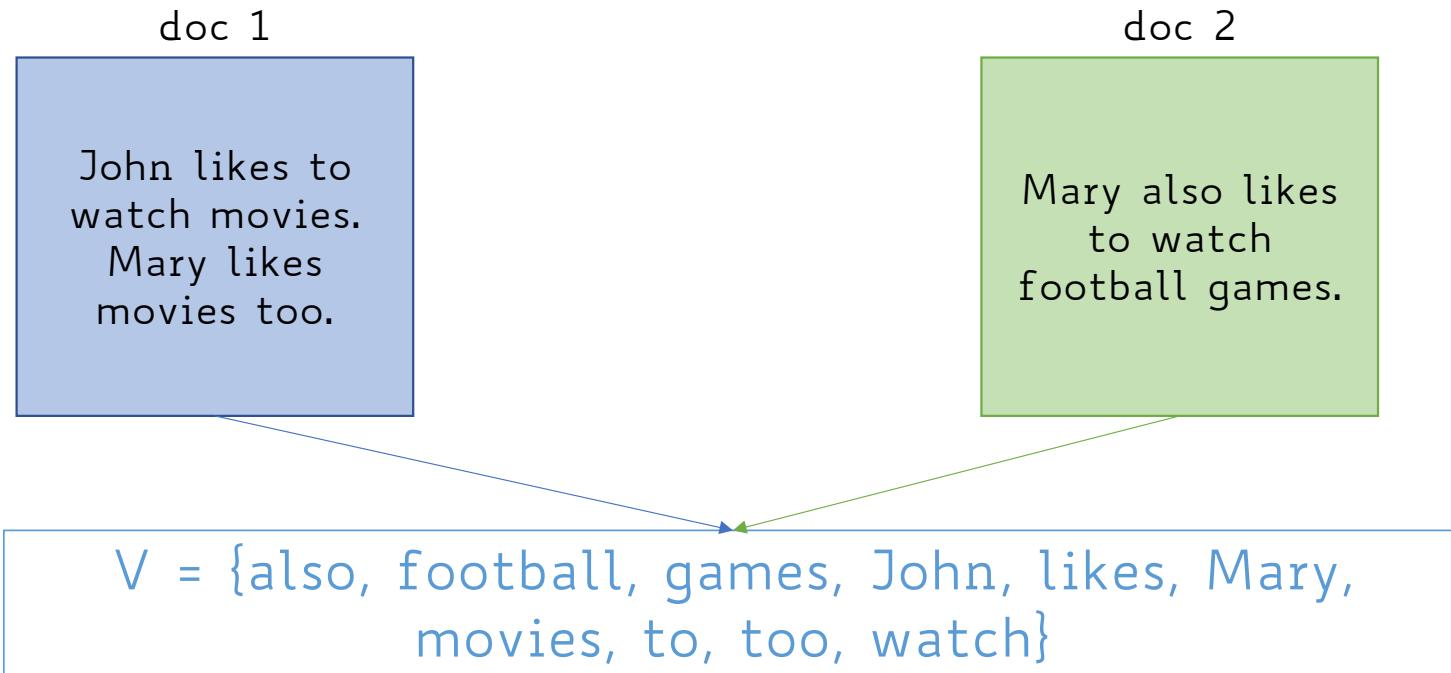
doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

Bag-of-Words:Vocabulary



Bag-of-Words: Weighting Scheme

doc 1

John likes to
watch movies.
Mary likes
movies too.

A $|V|$ -dimensional vector, where the i -th component indicates the multiplicity of the i -th word of the vocabulary

$V = \{\text{also, football, games, John, likes, Mary, movies, to, too, watch}\}$

Bag-of-Words: Weighting Scheme

doc 1

```
John likes to  
watch movies.  
Mary likes  
movies too.
```

A $|V|$ -dimensional vector, where the i -th component indicates the multiplicity of the i -th word of the vocabulary

$(0, 0, 0, 1, 2, 1, 2, 1, 1, 1)$

$V = \{\text{also, football, games, John, likes, Mary, movies, to, too, watch}\}$

Bag-of-Words: Weighting Scheme

doc 2

Mary also likes
to watch
football games.

A $|V|$ -dimensional vector, where the i -th component indicates the multiplicity of the i -th word of the vocabulary

$$V = \{\text{also, football, games, John, likes, Mary, movies, to, too, watch}\}$$

Bag-of-Words: Weighting Scheme

doc 2

Mary also likes
to watch
football games.

A $|V|$ -dimensional vector, where the i -th component indicates the multiplicity of the i -th word of the vocabulary

$(1, 1, 1, 0, 1, 1, 0, 1, 0, 1)$

$V = \{\text{also, football, games, John, likes, Mary, movies, to, too, watch}\}$

Bag-of-Words: A Formal Perspective

$D = \{d_1, \dots, d_N\}$ = collection of N documents

$V = \{w_1, \dots, w_{|V|}\}$ = **vocabulary** of $|V|$ words extracted from D

$\mathbf{d}_i = (f(w_1, i), \dots, f(w_{|V|}, i))$ = $|V|$ -dimensional vector representing d_i

$f : V \times D \mapsto \mathbb{R}$ is a function that maps each word of a document to a real value (**weighting scheme**)

Bag-of-Words: A Formal Perspective

One-Hot (binary) weighting scheme

$$f(w_j, i) = \begin{cases} 1 & \text{if } w_j \text{ appears in } d_i \\ 0 & \text{otherwise} \end{cases}$$

Bag-of-Words: A Formal Perspective

Term-Frequency weighting scheme

$$f(w_j, i) = tf(w_j, i)$$

tf computes the number of times word w_j occurs in document d_i

Bag-of-Words: A Formal Perspective

TF-IDF weighting scheme

$$f(w_j, i) = tf(w_j, i) * idf(w_j)$$

$$idf(w_j) = \log \left(\frac{N}{n_j} \right)$$

n_j is the number of documents in D containing the word w_j

Bag-of-Words: A Formal Perspective

TF-IDF weighting scheme

$$f(w_j, i) = tf(w_j, i) * idf(w_j)$$

$$idf(w_j) = \log \left(\frac{N + 1}{n_j + 1} \right)$$

Any idea
why?

n_j is the number of documents in D containing the word w_j

BoW: Limitations and Improvements

- **2 main limitations** of BoW model:
 - High dimensionality → sparseness
 - No sequential information nor semantics included → unigram model

BoW: Limitations and Improvements

- **2 main limitations** of BoW model:
 - High dimensionality → sparseness
 - No sequential information nor semantics included → unigram model
- Possible improvements:
 - Use n -grams rather than unigrams to capture sequentiality between consecutive words (i.e., context)
 - Even better, use so-called Neural Language Models like word2vec, BERT, and, more recently, Transformers (LLMs)

Bag-of- n -grams

Example: bigrams ($n=2$)

doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

Bag-of- n -grams

Example: bigrams ($n=2$)

doc 1

John likes to
watch movies.
Mary likes
movies too.

doc 2

Mary also likes
to watch
football games.

```
{"John likes", "likes to", "to  
watch",  
"watch movies", "Mary likes",  
"likes movies", "movies too"}
```

```
{"Mary also", "also likes", "likes to",  
"to watch", "watch football", "football games"}
```

Document Similarity

- We have examined a number of possible document representations

Document Similarity

- We have examined a number of possible document representations
- Depending on those, several similarity measures can be used

Document Similarity

- We have examined a number of possible document representations
- Depending on those, several similarity measures can be used
- For example, if documents are represented as:
 - set of words → **Jaccard coefficient**
 - one-hot bag-of-words → **Euclidean distance**
 - tf or tf-idf bag-of-words → **Cosine similarity**

High-Dimensional Spaces

- It's easy to get to very high-dimensional spaces

High-Dimensional Spaces

- It's easy to get to very high-dimensional spaces
- In the word space, the size of the vocabulary can be very high!

High-Dimensional Spaces

- It's easy to get to very high-dimensional spaces
- In the word space, the size of the vocabulary can be very high!
- Moreover, only few dimensions are non-zero

High-Dimensional Spaces

- It's easy to get to very high-dimensional spaces
- In the word space, the size of the vocabulary can be very high!
- Moreover, only few dimensions are non-zero
- Other domains like images, audio, etc. suffer from the same issue

High-Dimensional Spaces

- Data in a high-dimensional space tends to be **sparser** than in lower dimensions

High-Dimensional Spaces

- Data in a high-dimensional space tends to be **sparser** than in lower dimensions
- Data points are **more dissimilar** to each other

High-Dimensional Spaces

- In Euclidean space, the distance between two points is large as long as they are far apart along **at least one dimension**

High-Dimensional Spaces

- In Euclidean space, the distance between two points is large as long as they are far apart along **at least one dimension**
- The higher the number of dimensions the higher the chance this happens

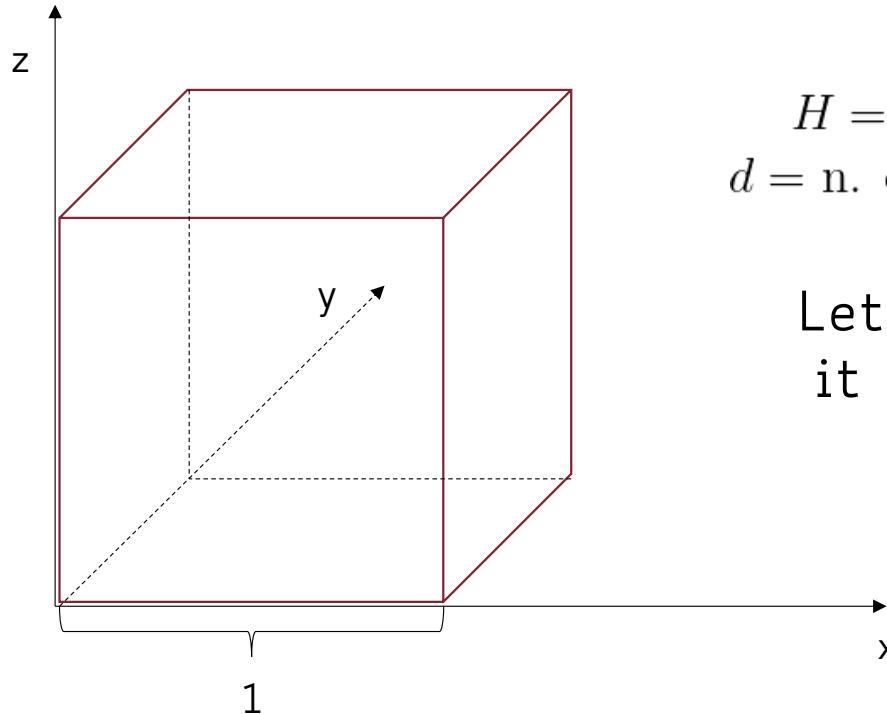
High-Dimensional Spaces

- In Euclidean space, the distance between two points is large as long as they are far apart along **at least one dimension**
- The higher the number of dimensions the higher the chance this happens



The Curse of Dimensionality

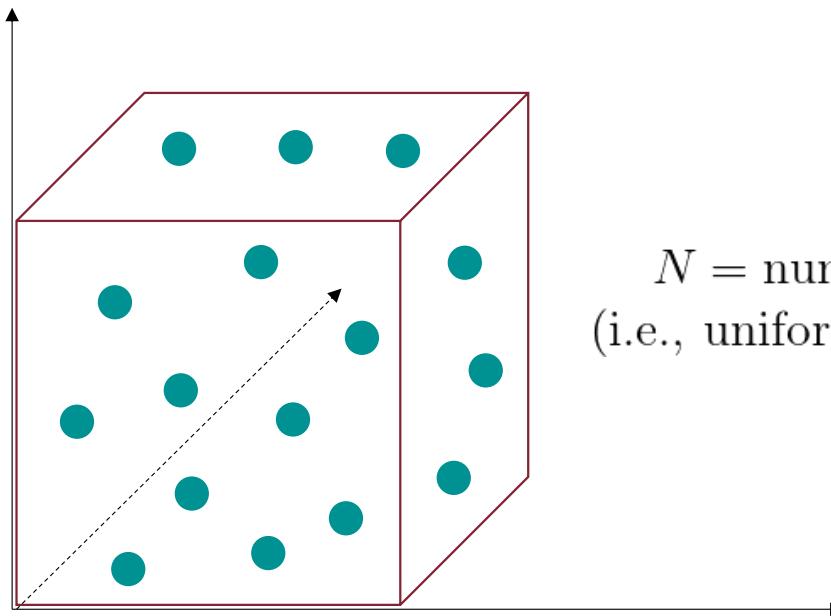
The Curse of Dimensionality



H = unit-length hypercube in \mathbb{R}^d
 d = n. of space dimensions

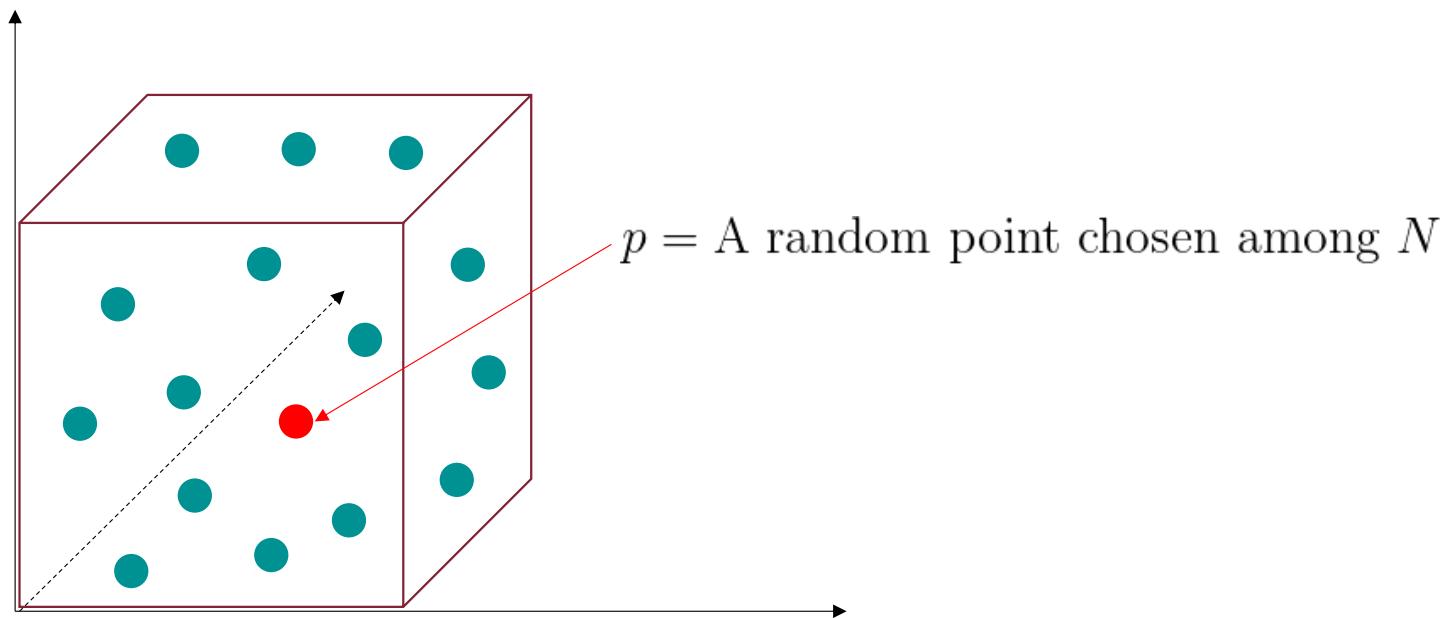
Let $d=3$ as beyond that
it is hard to visualize
the space

The Curse of Dimensionality

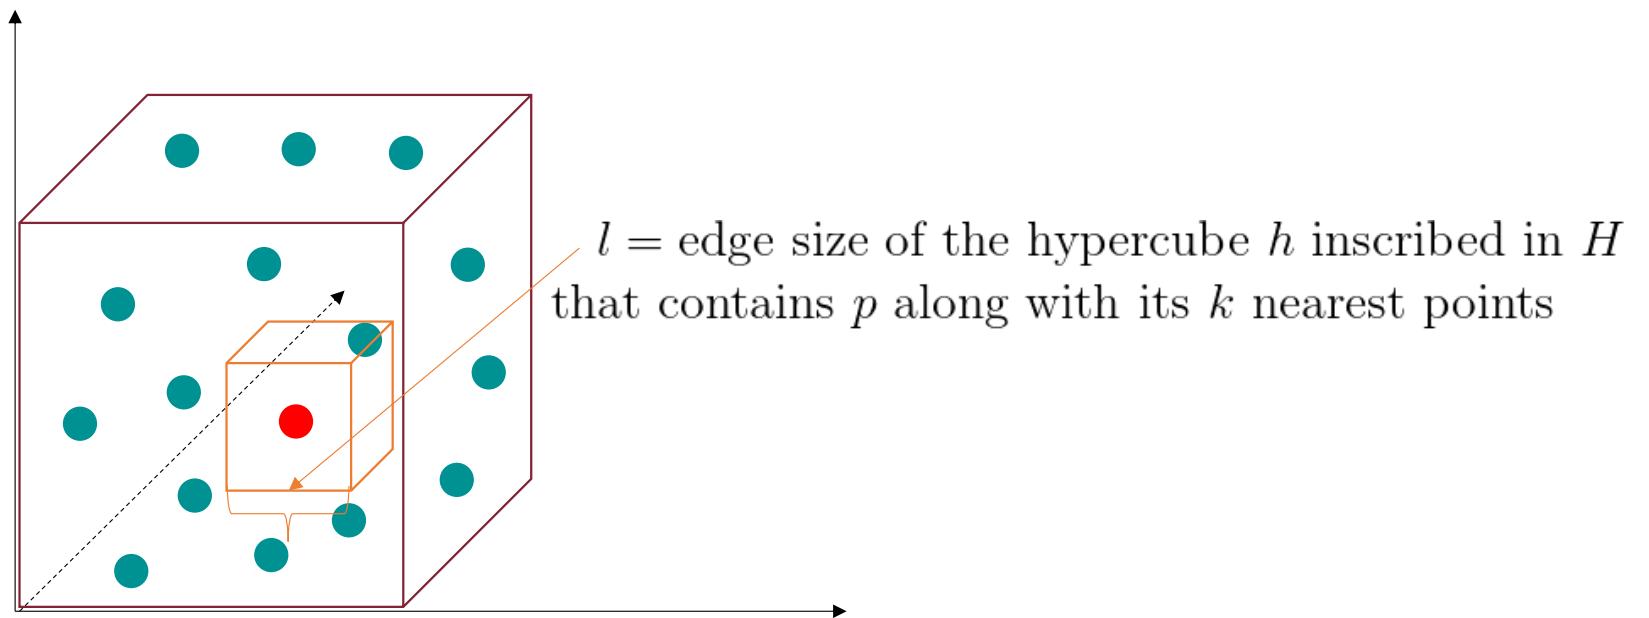


N = number of data points randomly
(i.e., uniformly) distributed in H

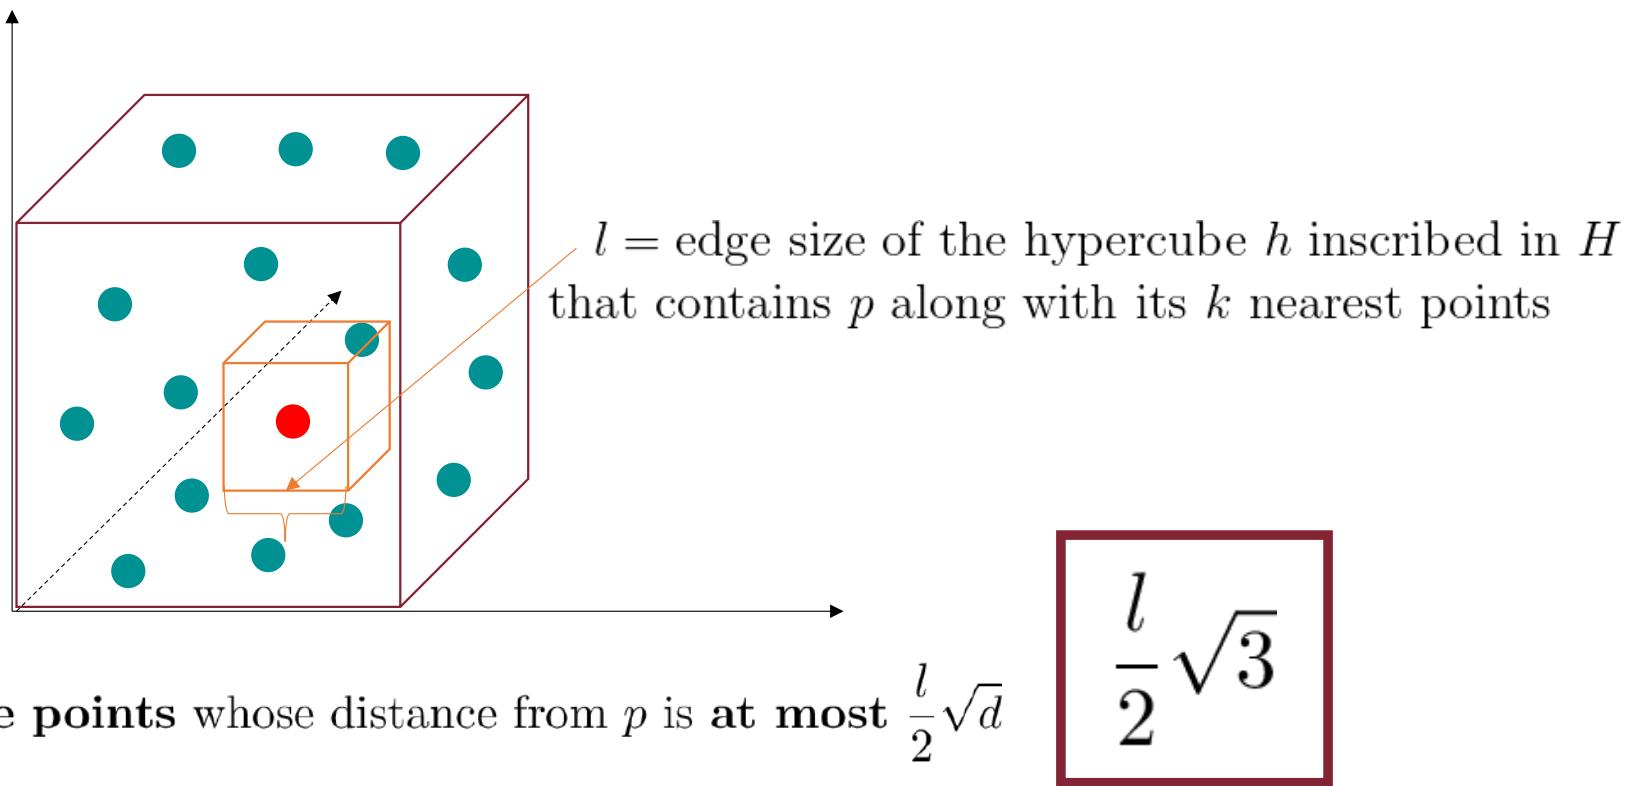
The Curse of Dimensionality



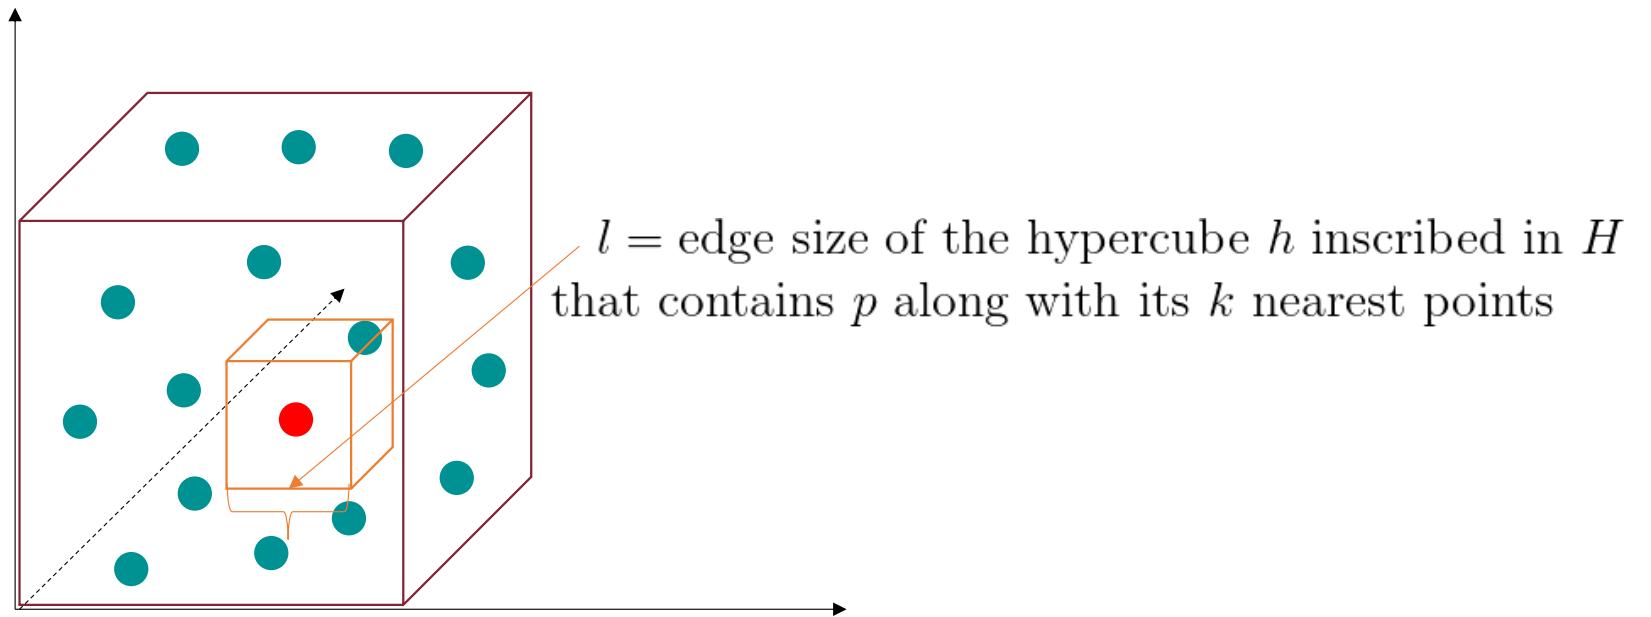
The Curse of Dimensionality



The Curse of Dimensionality

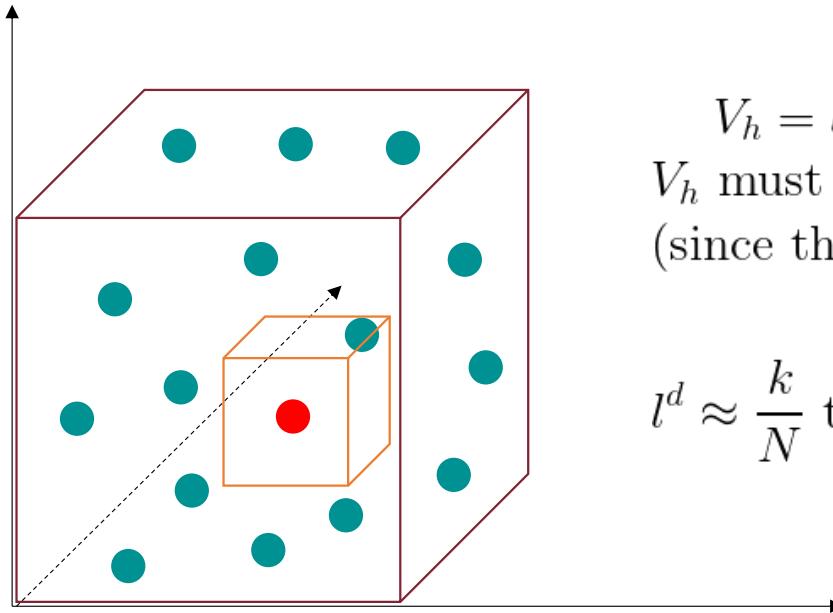


The Curse of Dimensionality



The same question can be formulated in terms of the radius l of an inscribed hypersphere

The Curse of Dimensionality



$V_h = l^d$ volume of the hypercube h
 V_h must roughly contain k/N points
(since those are randomly distributed)

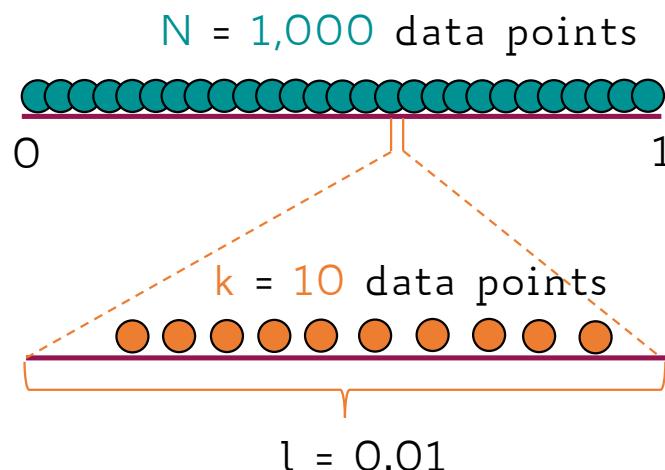
$$l^d \approx \frac{k}{N} \text{ therefore } l \approx \left(\frac{k}{N}\right)^{1/d}$$

The Curse of Dimensionality

A few numbers...

$$N = 1,000; k = 10 \quad l \approx \left(\frac{10}{1000} \right)^{1/d} = \left(\frac{1}{100} \right)^{1/d}$$

d	l
1	0.01

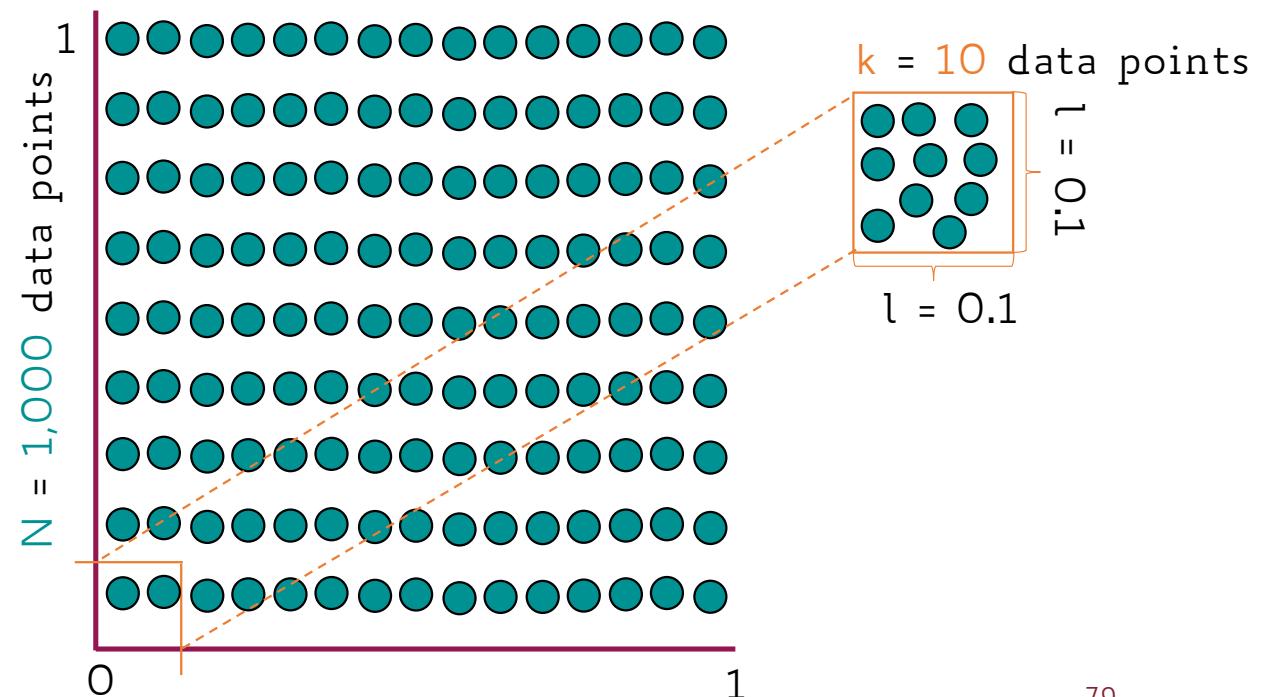


The Curse of Dimensionality

A few numbers...

d	l
1	0.01
2	0.1

$$N = 1,000; k = 10 \quad l \approx \left(\frac{10}{1000} \right)^{1/d} = \left(\frac{1}{100} \right)^{1/d}$$



The Curse of Dimensionality

A few numbers...

$$N = 1,000; k = 10 \quad l \approx \left(\frac{10}{1000} \right)^{1/d} = \left(\frac{1}{100} \right)^{1/d}$$

d	l
1	0.01
2	0.1
3	0.215
...	...
10	0.631

When d is equal 10 the length of the edge of the inscribed hypercube is already about 63% of the largest hypercube

The Curse of Dimensionality

A few numbers...

$$N = 1,000; k = 10 \quad l \approx \left(\frac{10}{1000} \right)^{1/d} = \left(\frac{1}{100} \right)^{1/d}$$

d	l
1	0.01
2	0.1
3	0.215
...	...
10	0.631
...	...
1000	0.995

When d is equal 1,000 there is basically no difference between the two hypercubes!

The Curse of Dimensionality: Why Bother?

- Points are more likely to be located at the edges of the region

The Curse of Dimensionality: Why Bother?

- Points are more likely to be located at the edges of the region
- Nearest points are not close at all!

The Curse of Dimensionality: Why Bother?

- Points are more likely to be located at the edges of the region
- Nearest points are not close at all!
- Distance between points indistinguishable (**distance concentration**)
 - Hard to separate between nearest and furthest data points
 - Hard to find clusters among so many pairs that are all at approximately the same distance

The Curse of Dimensionality: The Edge

Let ε define the **edge** (i.e., border) of our space

The Curse of Dimensionality: The Edge

Let ε define the **edge** (i.e., border) of our space

See how the probability of picking a data point that is **not** located at the edge changes as the number of dimensions grow

The Curse of Dimensionality: The Edge

Let ε define the **edge** (i.e., border) of our space

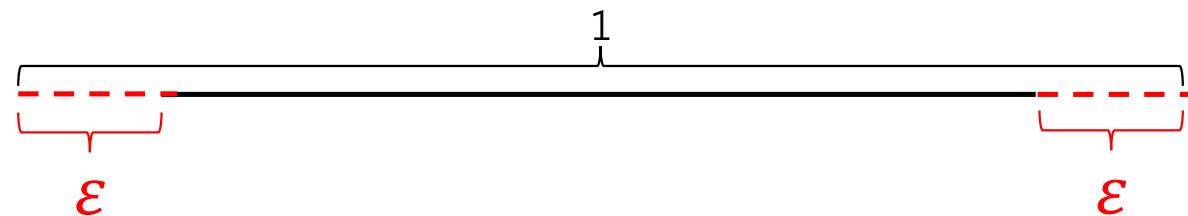
See how the probability of picking a data point that is **not** located at the edge changes as the number of dimensions grow

Remember:

We assume data points are **uniformly distributed at random** on the space

The Curse of Dimensionality: The Edge

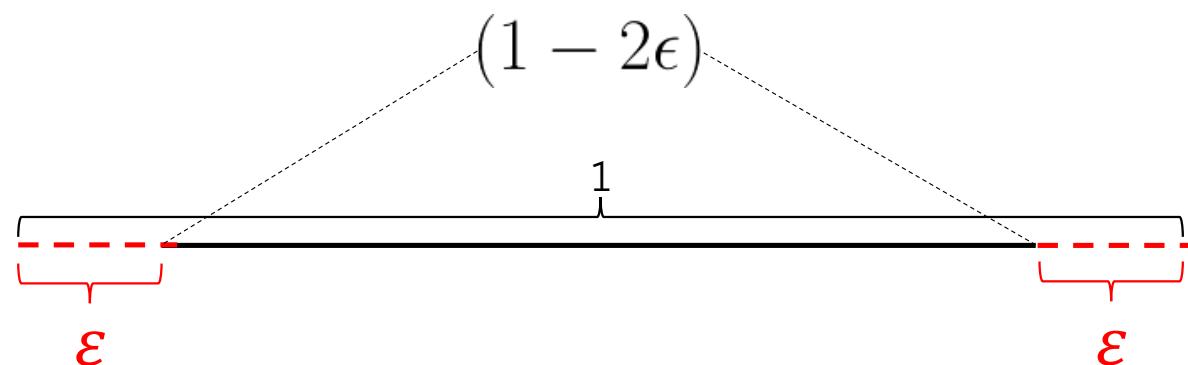
$$d = 1$$



The Curse of Dimensionality: The Edge

$$d = 1$$

The probability of being **not** at the edge is just



The Curse of Dimensionality: The Edge

$$d > 1$$

The probability of being **not** at the edge is
the probability of being not at the edge on
every single dimension

The Curse of Dimensionality: The Edge

$$d > 1$$

The probability of being **not** at the edge is
the probability of being not at the edge on
every single dimension

$$(1 - 2\epsilon)^d$$

assuming each dimension is independent from each other

The Curse of Dimensionality: The Edge

$$d > 1$$

The probability of being **not** at the edge is
the probability of being not at the edge on
every single dimension

$$(1 - 2\epsilon)^d$$

assuming each dimension is independent from each other

$$\lim_{d \rightarrow \infty} (1 - 2\epsilon)^d = 0$$

The Curse of Dimensionality

A Notebook where the Curse of Dimensionality is (visually) explained is available at the following link:

https://github.com/gtolomei/big-data-computing/blob/master/notebooks/The_Curse_of_Dimensionality.ipynb

So What Can We Do?

- If data are really uniformly distributed in a high-dimensional space... nothing!

So What Can We Do?

- If data are really uniformly distributed in a high-dimensional space... nothing!
- Luckily, though, real-world (interesting) data have patterns underneath (i.e., they are **not random!**)

So What Can We Do?

- If data are really uniformly distributed in a high-dimensional space... nothing!
- Luckily, though, real-world (interesting) data have patterns underneath (i.e., they are **not random!**)
- Lower intrinsic dimensionality

So What Can We Do?

The Manifold Hypothesis

- High dimensional data (e.g., images) lie on low-dimensional manifolds (i.e., sub-space) embedded in the high-dimensional space

So What Can We Do?

The Manifold Hypothesis

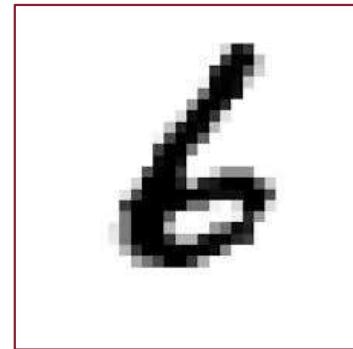
- High dimensional data (e.g., images) lie on low-dimensional manifolds (i.e., sub-space) embedded in the high-dimensional space
- Dimensionality reduction techniques (more on this later...)

Modeled vs. True Dimensionality

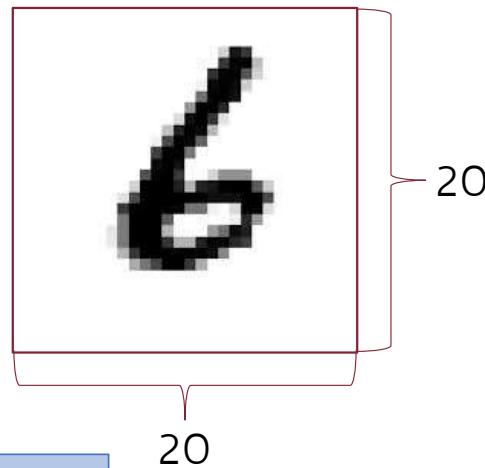
Example
Handwritten digit recognition

0	4	1	9	2	1	3	1	4	3
5	3	6	1	7	2	8	6	9	4
0	9	1	1	2	4	3	2	7	3
8	6	9	0	5	6	0	7	6	1
8	7	9	3	9	8	5	9	3	3
0	7	4	9	8	0	9	4	1	4
4	6	0	4	5	6	1	0	0	1
7	1	6	3	0	2	1	1	7	9
0	2	6	7	8	3	9	0	4	6
7	4	6	8	0	7	8	3	1	5

Modeled vs. True Dimensionality



Modeled vs. True Dimensionality

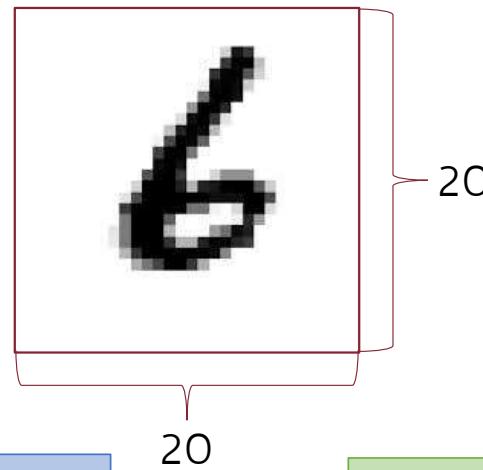


Modeled dimensionality

Each digit represented by
 20×20 bitmap

400-dimensional binary vector

Modeled vs. True Dimensionality



Modeled dimensionality

Each digit represented by
20x20 bitmap

400-dimensional binary vector

True dimensionality

Actual digits just cover a tiny fraction of all this huge space

Small variations of the pen-stroke

Modeled vs. True Dimensionality

Random samples
from 400-d space

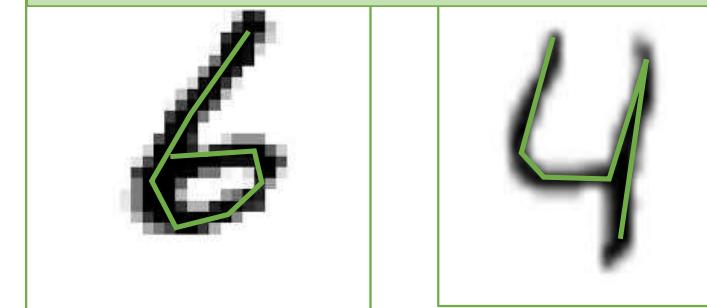


Modeled vs. True Dimensionality

Random samples
from 400-d space

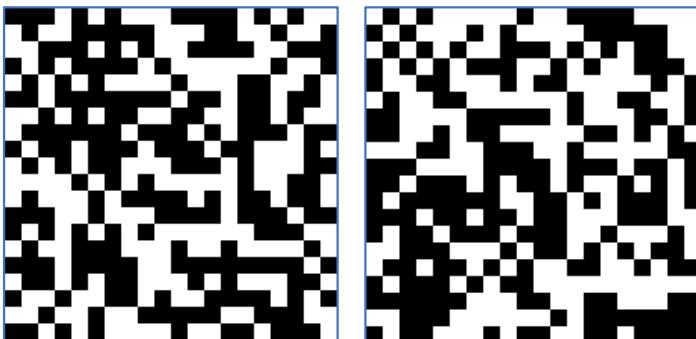


True digits living
in a
400-d space

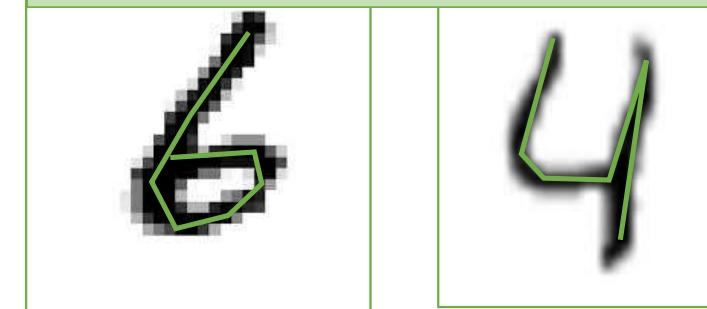


Modeled vs. True Dimensionality

Random samples
from 400-d space



True digits living
in a
400-d space



We model data (i.e., digits) as very high dimensional...
... In fact, they are not so

Take-Home Message of Today

- Many big data tasks are based on finding (hidden) commonalities between input data points

Take-Home Message of Today

- Many big data tasks are based on finding (hidden) commonalities between input data points
- For example, clustering is an unsupervised learning technique to group "similar" objects together

Take-Home Message of Today

- Many big data tasks are based on finding (hidden) commonalities between input data points
- For example, clustering is an unsupervised learning technique to group "similar" objects together
- Depends on:
 - **object representation**
 - **similarity measure**

Take-Home Message of Today

- In the Euclidean space, when data dimensionality gets large, similarity/distance becomes meaningless!

Take-Home Message of Today

- In the Euclidean space, when data dimensionality gets large, similarity/distance becomes meaningless!
- Any set of random data points are far away from each other → **curse of dimensionality**

Take-Home Message of Today

- In the Euclidean space, when data dimensionality gets large, similarity/distance becomes meaningless!
- Any set of random data points are far away from each other → **curse of dimensionality**
- Luckily, real-data may live in lower-dimensional spaces

